

Pull-Request Workflow with Agents

Last modified: 2026-10-03 11:30:33 (PDT)

The practices below apply whether you are driving an agent’s work or doing the work yourself. They keep parallel sessions from colliding and keep a pull request moving toward a clean, mergeable state.

1 File an Issue Before Starting

When starting a **new** piece of work, go **issue-first**: before branching, editing, or opening a PR, make sure a tracking issue exists. Search the tracker first with a qualifying all-state search, not an open-only listing; if no existing issue covers the task, **file one** (`gh issue create` / `glab issue create`), then proceed. Never jump straight into a PR without a tracking issue behind it.

On GitHub that search is `gh issue list --state all --search`; on GitLab it is `glab issue list --all --search` (glab has no `--state`). Not `--state open`: a closed issue for the same bug is the duplicate an open-only search cannot see, per [check-open-prs-before-duplicating](#). `hooks/warn-pr-create-without-dupe-check.py` warns (never blocks) when a `gh issue create` / `glab issue create` runs with no such query earlier in the session. It prompts the search; it does not judge the terms ([grep-is-not-coverage](#)).

The issue is the durable record of intent, scope, and “done” criteria — it gives reviewers context, lets the PR auto-close it via `Closes #N`, and keeps the work discoverable even if the PR stalls. Skip only when the task is already tracked by an open issue. A closed match is not a skip: surface it and confirm before re-doing the work.

This rule settles *whether* something is tracked, not *where* it goes. An item whose deliverable is a decision rather than a diff belongs on the discussion board instead, per [choose-issue-or-discussion](#) — so read “file one” here as “file one in the right venue”, which for actionable work is the tracker.

When the issue is a **bug report**, include a minimal reproducible example (a rerex — <https://reprex.tidyverse.org/>) whenever you can. A rerex is what a maintainer needs to

confirm and fix the bug, and it's what they'll ask for anyway, so providing it up front saves a round trip. The **reprexes** skill helps reduce the problem to a minimal, self-contained example.

When filing an issue that contains a list of independent subissues, file each subissue as a child issue linked under the parent (GitHub sub-issues feature: `mcp__github__sub_issue_write` in remote sessions, or `gh api` with the sub-issues endpoint in local sessions).

That splitting rule has teeth, and they are worth stating: a PR's Closes #N closes the whole issue, including every item in it the PR never addressed. Read as tidiness, the rule is easy to skip when the second item feels like a footnote. The actual consequence is that GitHub cannot partially close an issue, so the residual items are not deferred and not reopened — they are silently gone, and nothing in the merge, the PR, or the closed issue reports that anything was dropped.

It is worse than an ordinary lost to-do, because a closed issue is *evidence that the work was handled*. A later reader searching the tracker finds it closed and reasonably concludes every item in it was dealt with, so the loss is not merely silent but actively misleading.

So before writing **Closes #N**, re-read #N and confirm the diff covers all of it. When it doesn't, either split the remainder into its own issue first, or reference the parent with **Refs #N**, which links without closing.

- **Do:** split at filing time, or at the latest before the closing PR merges.
- **Do:** use **Refs #N** when a PR advances an issue without completing it.
- **Don't:** let **Closes #N** ride on an issue whose scope is wider than the diff.

(Morrison-Lab/ai-config#847, 2026-07-29: an issue was filed carrying a primary bug and a secondary note, and the PR fixing the first said **Closes #847**. The second item survived only because the maintainer asked about it before the merge, which is not a mechanism; it was split into #852 and shipped as #853, and both PRs merged within the following half hour. The splitting rule directly above already existed and was simply not applied when #847 was filed, which is the argument for stating its consequence rather than only its instruction.)

1.1 A closing keyword plus #N closes #N even when the sentence negates it

GitHub's parser matches **KEYWORD #N** as a substring. It does not read the rest of the sentence. A line that says the keyword is not being used still closes the issue when the keyword sits next to the number. The squash commit of #1718 closed #1717 that way, and the hook that commit shipped stayed unregistered until #2275 / #2294.

- **Do:** keep the number off the keyword (**Refs #N**, or “the closing keyword was not used for #N”).
- **Don't:** write a sentence that places a closing keyword next to #N in order to say you are not using it.

See [ardi.cases.md](#), “A negated closing-keyword sentence still closes the issue”, and [github-closing-keywords.md](#).

1.2 Deferring a request out of the current change is allowed, and the tracking issue is what allows it

The rule at the top governs work you are about to start. Its mirror governs work you are declining to start now: a request that arrives while a change is already in flight, and that would grow that change past what it set out to do. Such a request may be deferred, on your own judgment, **provided the deferred item is filed as an issue in the same reply**. The permission and the condition are one rule rather than two. An untracked deferral is not a deferral, it is dropping the request in the vocabulary of scope discipline.

The requests this covers come from the user, which is what makes it worth stating. A reviewer’s finding already has a *Defer* disposition, per [ardi](#)’s *ARD* step, and a request the user explicitly defers already routes to [defer-issue](#). Neither reaches the commonest case, where the user asks for something adjacent mid-review and the standing instinct treats any direct request as automatically in scope for whatever happens to be open. A request can be genuinely wanted and genuinely out of scope for the current PR at once, and saying so is a service rather than a refusal.

It is a grant of latitude and not an instruction to defer. The default is unchanged: do what was asked. What the grant removes is the bind a mid-flight request creates, where responsiveness and scope discipline pull opposite ways and doing everything asked is the only move that reads as cooperative.

File the issue so it stands alone, by this fragment’s own standard. The conversation that produced the request will not survive it, so an issue reading “do the thing we discussed” defers nothing and only moves the loss somewhere harder to notice.

Say which parts you deferred and why, in the same reply. This is the near-miss, and it reads as compliance from the inside: three things were asked, two were done, the reply describes the two, and nothing states that a third existed. A silent partial delivery is indistinguishable from having done the whole thing, so the user finds out what was dropped only by rereading their own request. Name the deferred item, give the reason, and link the issue.

1.2.1 The boundary with technical debt

[dont-incur-technical-debt](#) says a filed issue records debt rather than paying it, and that a defect you have already diagnosed inside your own diff is yours to fix now. Nothing here softens that, and the two rules read as contradictory until the boundary is drawn.

That fragment supplies the discriminator, so use its question:

Does the diff I am about to push contain the thing I just diagnosed as wrong?

When it does, the request is not out of scope, it is the scope, and no issue number buys it out. This rule covers work **adjacent to** the diff instead: pre-existing prose the change never authored, a broader sweep the change happens to touch one instance of, a follow-on improvement that would be welcome later.

- **Do:** defer an out-of-scope request on your own judgment, and file the tracking issue in the same reply that declines it.
- **Do:** name each deferred item, its reason, and its issue, so a partial delivery is visible as partial.
- **Do:** ask the technical-debt question first, and fix rather than defer whatever the current diff itself introduced.
- **Don't:** treat a request as in scope merely because the user made it directly.
- **Don't:** defer without filing — an untracked deferral is a dropped request, and it reads as scope discipline while being the opposite.
- **Don't:** read this as a reason to defer; the default is still to do what was asked.

(Directive from the user, 2026-08-09: “cai: it’s ok to defer out-of-scope requests from me; just make sure to track them in issues”. It came mid-review on UCD-SERG/serocalculator#654, a Quarto methodology-vignette formalization, where adjacent requests kept arriving in quick succession — convert propositions to theorems, sweep the chapter for overclaims, reformat multi-equality display equations — several of them touching prose the PR had never authored.)

2 Claim a PR or Issue Before Working on It

Before starting a work session on a GitHub PR or issue — i.e. before fetching the branch, making edits, posting a review, or invoking an automated review cycle — post a brief comment on the PR/issue so other people and any automated review bots know not to start a conflicting parallel session.

Use:

```
gh pr comment <N> --body "Working on this --- please hold off on pushing to this branch until
```

```
_Posted by Claude Code (AI agent) --- not written by a human._"
```

```
gh issue comment <N> --body "Working on this --- please hold off until I'm done.
```

```
_Posted by Claude Code (AI agent) --- not written by a human._"
```

Both halves of that body are load-bearing, and they answer different questions. The first line says the thread is claimed. The trailing line says **who claimed it**, and it is required on every

comment an agent posts to a forge — not only on a claim. See [disclose-agent-authorship](#), which carries the rule, the exact marker, and why the marker deliberately avoids the robot emoji.

Then proceed with the work. After the session ends (PR merged, issue closed, or work otherwise paused), follow up with a closing comment so the PR/issue is unclaimed for the next person.

Skip the claim step if the most recent comment already says you are working on it **and that claim is still live under the expiration rule below**. This applies to any task that will push commits to a PR branch, run iterative review loops, or post a review. A posted review races HEAD the same way a write session does: other sessions push while the reviewer is still reading, and the posted comment then stamps a SHA that is already stale. Post the claim **before** the review starts. When the session is **review-only** (it will not push), unclaim when the SHA-stamped review comment lands so the author can address findings. When the session is also driving the branch (implementing, ARDI), keep the write claim until that work ends — posting a review mid-loop does not release it. A persistent watch is not a standing claim — re-claim only when a new review round starts. A review-only claim still expires under the 2-hour rule below; reassert it if the pass is still running and the thread has been idle that long. It does **not** apply to read-only inspection that will not post (showing a PR, checking status, explaining a diff).

- **Do:** post a **hold off** claim before starting a posted review of a PR.
- **Do:** unclaim a review-only pass when that review comment is posted.
- **Do:** keep a still-driving write claim after posting a review in the same session.
- **Don't:** skip the claim because the session is “only reviewing” and not pushing — the collision is on HEAD, not on the working tree.
- **Don't:** leave a review-only claim standing after the verdict lands, or claim every open PR at the start of a sweep.

This includes a PR **you opened yourself**: in repos with an active `@claude` agent (`claude.yml`), the agent can push commits to your branch on PR activity — e.g. merging `main` in — and collide with your in-flight push, so claim early to flag the branch as actively worked. (See [memories/claude-bot-workflows.md](#), “`@claude` CI action”, for the collision-recovery steps.)

When starting work from an issue, follow the claim comment with an immediate draft PR — see [pr-on-claim](#) for the mechanics. An open PR is a stronger “in-flight” signal than a comment alone.

A claim expires 2 hours after the most recent push or comment on the PR/issue — reassert it rather than resuming under a stale one. A claim comment with no expiry binds the thread indefinitely: a crashed or abandoned session leaves its “hold off” standing forever, and a second session has no rule for when the claim stops blocking. So the convention is time-boxed and keyed to observable activity: a claim is **live for 2 hours from the most recent push or comment** on the PR/issue, and **expired** past that.

The rule cuts both ways.

- **As the claimant:** resuming work after more than 2 idle hours — no push and no comment in that window — starts with a fresh claim comment, not with an edit. The skip-if-already-claimed shortcut above covers only a live claim; an expired claim of your own no longer covers you, because a parallel session is entitled to treat it as lapsed.
- **As a would-be second session:** another claimant’s claim whose PR/issue shows no push or comment in over 2 hours no longer blocks you. Take over by posting your own claim comment, never by starting silently — the fresh claim is what flips the thread’s state, and it is what tells the stale claimant they were superseded if they return. A claim’s age is evidence about the *claim*, not proof the branch is quiet, so the mid-task checks below — the “already done” cross-check against the PR’s actual commit list, and the rejected-push tree comparison — still apply before your first push, as does the branch-head re-fetch in the [claim-pr](#) skill’s Notes.

Staleness is decidable by one read rather than by judgment, per [algorithmatize-checks](#):

```
gh pr view <N> --json updatedAt --jq .updatedAt      # VIEW_PR
gh issue view <N> --json updatedAt --jq .updatedAt    # VIEW_ISSUE
```

`updatedAt` moves on more events than pushes and comments (labels, reviews, body edits), so it only ever **over-approximates** freshness: a stale verdict from it is definitive, and a borderline-fresh one defaults to respecting the claim — the safe direction, since over-respecting a dead claim costs a wait while under-respecting a live one costs a collision.

- **Do:** post a fresh claim before resuming work when more than 2 hours have passed since the most recent push or comment on the PR/issue.
- **Do:** treat another session’s claim as expired on the same 2-hour reading, and post your own claim before touching anything.
- **Don’t:** read “the most recent comment already says I’m working on it” as a standing skip — that shortcut covers only a claim under 2 hours old.
- **Don’t:** start work under an expired claim, your own or anyone else’s, without a fresh claim comment — a silent resumption and a silent takeover collide identically.

(Directive from the user, 2026-08-15: “let’s set a convention that pr and issue claims last 2 hours from the most recent push or comment; if it’s been longer than that, reassert your claim.”)

Every detector of a claim matches the OLD wording as well as the new one, and dropping the old alternation is the one edit that fails silently. There were **two** retired invariants, not one, and enumerating them from the file in front of you is how the second was missed for a whole review round. Most emitters carried the **paws off** invariant — `claim-pr`, `gi`, `st`, `gip`, `pr-on-claim`, `post-merge`, `handoff` (as “still claimed, paws off.”) and the orchestrator (as “paws off until done”). [ardi](#) did not. It said “back off until done”, and had done since 2026-06-17 — and it is the corpus’s highest-traffic claim emitter, run on every PR in every repo. Both now read “please hold off ...”.

Derive that set from history rather than from the current tree, which no longer contains any of them. Match the **wordings**, not the `--body` flag. Keying on the flag looks tighter and reaches neither `gip`, which states its claim as quoted prose inside a worker brief, nor the orchestrator, which builds the body as a Python f-string on a different line from the flag — the two emitters least like the others, and so the two a flag-shaped search is least able to find. Widen the pathspec past `skills/` for the same reason.

```
git log -p --all -- 'skills/**' 'shared/**' 'commands/**' 'scripts/**' \  
| grep -oiE '^\+.*(paws off|back off|hold off)[^"]*' | sort -u
```

That returns the matcher and prose lines too, so it wants a skim rather than a count — the deliberate trade for a search that cannot miss an emitter because of how it happened to spell the call.

Claims posted before that are still sitting on open PRs and issues, and a claim stays live on activity rather than on age — so a thread claimed under the old wording and pushed to this morning is live right now.

A detector narrowed to the new phrase alone still returns cleanly on such a thread. It returns **nothing**, which is indistinguishable from an unclaimed thread, and that reading licenses exactly the parallel session this whole convention exists to prevent. Nothing in the output announces the miss: a claim search that finds no claim looks the same whether the thread is free or the matcher went blind.

Match the two-word invariant, never a whole sentence. The claim body varies by target — a PR claim says “please hold off on pushing to this branch until I’m done” and an issue claim says “please hold off until I’m done” — so neither sentence contains the other, and a detector keyed on either one is blind to half the claims. Under the old single-string wording that distinction did not exist, which is exactly why it is easy to carry a whole-sentence matcher across the rename without noticing it has narrowed. `hold off` is the invariant. `paws off` and `back off` are its two predecessors, and a matcher naming only the first is the failure this very section describes, committed by the section itself.

So match the alternation, case-insensitively, everywhere a claim is read:

```
gh pr view <N> --json comments \  
-q '.comments[] | select(.body | test("hold off|paws off|back off"; "i"))' # READ_PR_COM
```

Keep both old alternatives until no claim under a retired wording can plausibly still be live — which, given the 2-hour rule keys on activity and not on the comment’s own age, means until every PR and issue open on 2026-08-24 has closed. Removing it is a deliberate later edit, not tidying to do in passing.

Then check the same comment for a release term, because one release marker contains a claim invariant. The retired release wording is ... done --- paws off released.,

which matches `paws off` — so the invariant that fixes the whole-sentence bug introduces a second one, and this one fails the quiet way round: a *released* PR reads as claimed, the reader backs off, and nothing reports why. The sentence matcher this replaced did not collide, so the collision arrived with the fix. Treat a comment as a release rather than a claim when it also matches `unclaim|released|PR is free|now mergeable`, and derive that list rather than copying it: `grep -rn "unclaim\|released\|PR is free\|now mergeable" skills/commands/`.

- **Do:** match `hold off|paws off|back off` case-insensitively wherever a claim is read, then exclude the comment if it also carries a release term.
- **Do:** treat both old alternatives as load-bearing until the threads carrying them have closed.
- **Don't:** read an empty claim search as an unclaimed thread without first confirming the matcher covers both wordings — the two results are identical.
- **Don't:** drop a back-compat alternative as part of an unrelated change.
- **Don't:** enumerate the retired wordings from the files you happen to be editing — `back off` was invisible to exactly that method for a full review round, because the one file that posted it was not one of the seven that agreed with each other.

Verify a mid-task “already done” claim against real PR state before trusting or redoing it. A PR you claimed and are actively driving can still gain commits from a **second, independently-running session** under the same account. A `<github-webhook-activity>` `review-comment-reply` event can describe work this session never did, in the form “Addressed, pushed in `<sha>`”. Don’t assume it’s fabricated or injected, and don’t reflexively redo the same fix: cross-check the PR’s actual commit list (`gh pr view --json commits / pull_request_read get_commits`) and review threads before either (a) trusting the claim, or (b) starting the same fix yourself. If a commit with that SHA genuinely exists, authored close to when the event arrived, treat it as confirmation a live parallel session owns this PR right now — stop pushing further speculative fixes yourself, and, if genuinely in doubt, ask whether to keep driving or step back, rather than racing the other session’s pushes. This gap is distinct from the initial claim check above: it’s not about claiming a PR before starting, but about **re-verifying you’re still the sole active driver** once work has been under way for a while — especially when you picked up the PR mid-session (e.g. by answering a diagnostic question about it) rather than through the normal claim-then-branch flow, so no fresh claim check ever ran right before you started pushing.

(Morrison-Lab/gha#286, 2026-07-24: a webhook event delivered a review-comment reply attributed to the **repository owner**, reading exactly like a Claude-authored reply and claiming a fix this session hadn’t made, worded “Addressed, pushed in `3fb8c5b`”. It was verified real via `get_commits` before proceeding — a second live session, not injection.)

The git-level variant of that check: a rejected push whose remote commit is byte-for-byte what you were about to push. The section above covers a *comment* claiming work was done. Here the parallel session makes no claim at all. Your `git push` is simply

rejected because it pushed first, and what it pushed is the same merge you just made. The reflex on a rejected push is to merge again, which would stack a redundant merge commit on top of an identical one.

Four reads settle it before you touch anything:

```
git rev-parse HEAD^{tree}          # your merge's tree
git rev-parse origin/<branch>^{tree} # theirs
git show -s --format=%P HEAD       # your merge's parents
git show -s --format=%P origin/<branch> # its parents
```

An identical tree plus identical parents means the two merges are the same merge, so the right action is `git reset --hard origin/<branch>`.

- **Do:** compare trees and parents before deciding what a rejected push means.
- **Do:** discard your local merge with `git reset --hard origin/<branch>` once both match.
- **Don't:** re-merge reflexively on a rejected push — that is what produces the redundant merge commit.
- **Don't:** force-push over the other session's commit.

(Morrison-Lab/ai-config#965, 2026-07-31: `main` moved one commit, a local `git merge origin/main` was made, and the push was rejected. The remote carried `b8d2273`, a merge of the same two parents, with tree `1bda1bc`, identical to the local merge's.)

Matching parents with a differing tree means the two sessions resolved the same merge differently — merge the two commits together, don't reset onto either one. Identical parents but a different tree is not the “same merge” case above: a concurrent session (a human, or an @claude-style bot reacting to PR activity) merged the same `main` commit into the same branch, but resolved a real conflict (or made an additional fix) differently than you did. `git reset --hard origin/<branch>` here silently discards whatever your version got right that theirs didn't — e.g. a version-parity bump their merge didn't carry, or a merge-conflict resolution theirs got wrong. Instead, fetch and merge the remote branch into your local one (an ordinary three-way merge, since the two commits share both parents as a common ancestor between them); resolve any conflict on its own merits, the same as any other merge, then push the result.

- **Do:** compare parents first, then trees; matching parents with a differing tree calls for a merge of the two commits, not a reset.
- **Do:** merge the remote branch in normally, resolving whatever differs on its own merits.
- **Don't:** `git reset --hard` onto a same-parents-different-tree remote commit — that discards your own resolution outright, on the assumption the two were interchangeable.

(UCD-SERG/serocalculator#654, 2026-08-08: `main` had absorbed a same-version dev bump from an unrelated PR, so this session merged `main` in and bumped the version past it while separately resolving a real conflict in `inst/WORDLIST`. The `@claude` review bot's own `main-sync` had pushed a merge with the same two parents in the meantime, but it left the version at parity — still failing `version-check` — and had never seen the `WORDLIST` conflict at all, since its sync predated that conflict existing. Merging the two commits, rather than resetting onto either, kept both fixes.)

Second occurrence, 2026-08-24, with DIFFERENT parents — the same lesson in the case the section above does not describe. Both branches of the rule above turn on the parents matching, which happens when two sessions resolve the same `main`-merge. The commoner collision has matching parents nowhere in sight: two sessions independently fix the **same review round's findings** on one branch, five minutes apart, so the two commits share only their base and the divergence is ordinary.

What carries over is the part that is not about parents. “They addressed the same findings” is exactly the belief that makes a reset feel safe, and it is the same interchangeability assumption the `Don't` above rejects. Measured on [ai-config#2185](#), 2026-08-24 Pacific (3b8d04e6 at 18:23 and cf195e46 at 18:28, merged as 5c577a54 at 18:35): each side had a fix the other lacked — 3b8d04e6 an attached `-F` body (`-Fbody.md`), cf195e46 a flag-boundary lookbehind without which a **compliant** comment warned — so either reset would have shipped a regression that no check could see, since both sides were green. A third apparent difference was inert: cf195e46 dropped the `--comment=` alternative the base already carried, and `--comment\b` matches `--comment=x` anyway, since the boundary sits between `t` and `=`. That one is the reason to **execute** both patterns rather than read both diffs.

- **Do:** merge two independent fixes for one review round, whatever their parents, and **run** each side's version against the inputs at issue — a diff shows what each side changed, and only executing it shows which of those changes did anything.
- **Don't:** reset onto the other session's commit because the two rounds answered the same findings — addressing the same list is not producing the same fix.

Third occurrence, 2026-08-30, with a posted claim comment already standing — and this time the reconciliation itself was orphaned by the merge outracing it. The two occurrences above both kept content from each side by pushing a merge that combined them. This one shows a third way to lose the race: the reconciling merge can be **correct and still never ship**, because the PR merges before it lands.

On [ai-config#2668](#), this session posted a claim comment (“Claude Code CLI ... is taking over this PR ... please hold off until I'm done”, 17:48:33Z) and kept driving the branch. The claim did not stop a second driver: at 20:55:26 this session committed c2153168, continuing its own paragraph-window design for the re-raise veto, and at 21:00:55 a second identity's commit 1bea0e0b (“scan entire containing section for reraise veto”) landed as a structurally different design for the identical gap. This session then prepared a local merge, b5794bc4, reconciling the two — its own message states plainly that “the `@claude` bot and this session fixed the same

three findings concurrently” and gives the reason for taking the bot’s design whole: “on the central one its design is strictly better: scanning the whole containing section ... is wider than a window of the containing paragraph plus one on each side, so it vetoes more and keeps more citations – the safe direction.” That reconciliation never entered history. `5dfcb6e7`, the commit that actually merged as the PR’s squash, is byte-identical in its full tree to `1bea0e0b` — the bot’s commit taken as-is, with none of `b5794bc4`’s additions — and it merged at `21:09:30`, six seconds **before** `b5794bc4` was authored at `21:09:36`. The reconciliation was correct (its rationale for preferring the wider scan holds, and the shipped code is the safe design either way), but it was already too late: the PR had merged while it was still being written, so its extra rationale comments were simply never pushed and are absent from `main` today.

- **Do:** when two independent fixes turn out to be competing designs for the same gap rather than complementary ones, compare them **on the merits**, but re-check the PR’s live state before trusting that comparison as the resolution — a merge racing ahead of your own reconciliation is a distinct failure from the collision the reconciliation was meant to fix.
- **Do:** treat a posted claim comment as a signal, not a lock; keep checking `git ls-remote` before every push exactly as [check-before-pushing](#) already prescribes, claim or no claim.
- **Don’t:** assume a live claim comment stops a bot (or a second session) from pushing to the branch it names — this is the same collision [check-before-pushing.md](#)’s “Ownership is what suppresses the check” section already lists, observed again with a claim already standing.
- **Don’t:** treat a local reconciliation commit as shipped just because it exists and is correct — check the PR’s actual merged head (`gh pr view <N> --json mergeCommit,mergedAt`) before crediting it with anything, since a merge that completes first ships the branch as it stood, not your unpushed fix-up on top of it.

Handing off mid-task to another agent, on user request (“finish what you’re doing, then relinquish holds; I’ll put another agent on them”): don’t just stop — leave the next agent a clean starting point. On each claimed PR/issue: (1) post a status comment on the PR itself distinguishing what’s **done** from what’s genuinely **not done** (the actual point of the issue, not just the side-fixes found along the way) and any blocker still open, so the next agent doesn’t have to re-derive it from the diff; (2) post the closing/unclaim comment on the issue per the pattern above; (3) `unsubscribe_pr_activity` (or stop babysitting locally) so you don’t keep auto-fixing a PR you no longer own; (4) stop any background watch/poll task tied to that work (e.g. a `ScheduleWakeup` or a `Monitor/background-Bash wait`) so it doesn’t fire into a session that’s moved on. A merge-conflict-free `git status` and a pushed branch are not enough on their own — the status comment is what makes the handoff legible. (ucdavis/bcs `gia` session, 2026-07-06: handed off PRs #310 and #311 mid-implementation this way, each blocked on the same slow `renv::restore()`.)

3 Keep Your Branch Synced with Main

Whenever `main` has moved ahead of a PR branch you're working on, **merge `main` into the PR branch** before the next push or review trigger. Don't wait for a conflict to surface or for someone to ask.

Worked-example case records for the rules below live in [sync-with-main.cases.md](#), moved out of the auto-loaded context.

This fragment covers the `single-branch-vs-main` case. When orchestrating a multi-agent `ultracode` session, merges can happen at more points than that — see [ultracode-merge-conflicts](#) for the broader check (worktree-isolated agent branches, concurrent `parallel()` results) and the note on GitHub's mergeable indicator not evaluating custom `.gitattributes` merge drivers.

Always check for merge conflicts with `main` before pushing results to remote. Run this before every push, not just before triggering a review:

```
git fetch origin main
git log --oneline ..origin/main | head    # any commits? main is ahead --- merge it in
git merge origin/main
```

If the push is rejected because `main` has moved (! [rejected] with (fetch first) or (non-fast-forward)), fetch and merge before retrying — don't force-push.

Always do this before triggering a fresh review too, so the reviewer evaluates the PR against current `main` rather than a stale snapshot.

Don't rebase or squash-rewrite a published PR branch unless explicitly asked — a merge commit is the right move because it matches GitHub's "Update branch" button and preserves the PR history.

If the merge has conflicts, resolve them, run the project's standard pre-commit checks (render / lint / spell / tests), commit, then push. Don't push a half-resolved merge.

A sync-only push invalidates the previous commit's review verdict — never arm auto-merge after syncing. Merging `origin/main` in and pushing creates a new HEAD commit ref that is unreviewed until fresh reviews land. Arming `gh pr merge --auto` after a sync push risks merging an unreviewed head as soon as CI passes ([fully-clean](#), [Pattern 12](#)). Always wait for fresh reviews and CI on the new head, re-verify with `check-pr-fully-clean.py`, and merge directly.

After merging `main`, re-check version parity. In R packages with a `version-check` CI job, the branch's `DESCRIPTION Version:` must *exceed* `main`'s. A conflict-free merge can silently put them at parity — `main` advanced (e.g. another PR merged between when you last bumped and now). After every merge of `main`, compare versions:

```
git fetch origin main
git show origin/main:DESCRIPTION | grep ^Version
grep ^Version DESCRIPTION
```

If they match, bump the branch’s `Version:` by one patch level before pushing.

Re-check main again right before the final push, not just at the start of a merge. Resolving a conflict (rerunning generators, fixing prose, updating a CHANGELOG entry) can take long enough for `main` to advance a second time. A `git fetch origin main` immediately before `git push` — after conflict resolution is done, not only before it started — catches that case; an earlier CI failure on a commit you thought was current is a symptom of skipping this second check.

A conflict-free merge does not mean derived artifacts are in sync. If your branch regenerates a generated tree (e.g. `codex-skills/`, a lockfile, rendered docs) and `main` added a new *source* input the generator consumes (a new skill, a new dependency), `git` merges both cleanly — but the generator never ran against the new input on your branch, so its output is missing or stale and the sync check fails on `main` after both land. After merging `main`, re-run the generator and commit the result whenever `main` touched the generator’s inputs — don’t trust the absence of conflicts. (Concretely: merge the PR that adds the new skill *first*, then sync the wrapper-regenerating branch and rerun `scripts/sync-codex-skill-wrappers.py` before merging it.)

A CI failure on a brand-new PR’s very first commit (e.g. the empty claim-commit from `pr-on-claim`) is a signal to check main’s position before debugging the failure itself. A local checkout that sat around since before the session started can already be many commits behind `main` — the failure (a stale generated-tree check, a check `main` has since added or dropped) often isn’t a real problem with your change at all, just `main` having moved. `git fetch origin main && git log --oneline ..origin/main` first; if `main` is ahead, merge it in and re-run the checks before treating the failure as something to fix in the diff.

The same staleness trap has a silent variant with no CI failure to flag it: a `worktree/branch` named after a PR’s followup can still be based on a `main` from before that PR actually merged. A `worktree` directory or branch name suggesting “after PR #N” (e.g. `pr-N-followup-...`) is not proof the branch’s actual base commit postdates #N’s merge — it can have been created earlier and simply named for its intended purpose. Trusting that naming, then reasoning from `git show <hash>` for a commit found via `git log --all` (which lists every reachable commit across all refs, not just your branch’s ancestry) can make content look present when it isn’t actually in your branch yet. Verify with `git log --oneline HEAD..origin/main` or by reading the actual blob your branch would produce (`git show HEAD:<path>`, or the working tree itself before assuming what it contains), not a commit hash pulled from `--all`. If `main` has moved, merge it in before building further edits on the assumption the missing content exists.

A real conflict inside a file whose logic is also copied elsewhere (an extracted script, a doc example) needs the copy re-synced too, not just the conflicted file resolved. When a PR extracts inline logic (e.g. a workflow step's shell block) into a standalone script for testability, and `main` independently changes that same inline logic while the PR is open, resolving the merge conflict in the workflow file is not enough — the extracted script must be updated to match `main`'s new logic exactly, or the PR silently reverts `main`'s fix the moment it merges. Diff the extracted copy against `main`'s current inline version line-for-line (strip indentation, `diff`) to confirm an exact match, not just “looks about right.” If the PR carries tests against the extracted copy (fixtures, unit tests), add regression coverage for whatever `main`'s change fixed — the merge is the natural moment to catch a gap the original PR's tests didn't anticipate, and to prove the new fixtures actually catch the regression (temporarily revert the fix, confirm the test fails, then restore).

When `main` DELETES a file your branch references, resolving the marked conflict is not enough — `grep` the whole tree for the deleted path. Git only conflicts where both sides edited the same lines, so a merge that brings in a deletion flags the file that *used* the thing, and nothing else. Any other reference to the deleted path — a docstring citing it as precedent, a comment, a doc cross-reference — merges cleanly and silently becomes a dangling reference, because those files were never in the conflict's scope. After resolving any merge that removed a file, run `grep -rn "<deleted-path>"` across the repo and re-point or reword each hit; then distinguish live references (must be fixed) from historical citations of the removal itself (correct as-is, leave them). This is the deletion counterpart to the extracted-copy case above: there the logic moved and a copy went stale, here it vanished and the pointers went dead.

A textual conflict in a skill file can be the symptom of a conceptual duplicate, not just competing edits to the same line. When merging `main` into a branch that's authoring a new skill, if the conflict lands in a `## Relationship to other skills` section (or `main` added an entirely new skill in the same territory), that's a signal to re-run `skill-builder`'s Step 0 judgment — not just resolve the diff mechanically. Compare the new skill against whatever landed on `main`: are they the same concern (fold into one, redirect), or genuinely distinct (cross-link both directions so neither reads as an unexplained near-duplicate)? `skill-builder`'s in-flight-work scan only runs once, at the start; `main` can grow a colliding skill in the time a PR is open, so the check has to be repeated at merge time too.

The same collision can land before you write a line, and then it produces no conflict at all — just duplicated work nobody flags. The bullet above catches a duplicate at merge time, via a conflict. When `main` gains the colliding content while your change is still *planned* rather than written, there is nothing to conflict with: you write the duplicate, push it, and the review has to argue you out of content that was already redundant on arrival. So re-run the dupe check after any fetch that brings in new commits, not only at merge — a plan researched an hour ago was researched against a different `main`.

The cheap version is to read what actually arrived rather than only the count: `git log --oneline <old>..origin/main` plus `git diff --stat` over the same range, then ask whether

any of it covers something still on your list. In a session that loads skills or plugins from the repo, a new one appearing in the session’s own skill listing is the same signal arriving for free.

This is a *timing* gap, and it composes with the *scope* gap rather than replacing it. `check-open-prs-before-duplicating` covers work that is still in flight, unmerged, and therefore invisible to any check against `main`; run that one too, since a duplicate is just as wasted whether the collision has landed yet or not. Both checks share the same weakness — each runs once, at the start, and answers for the moment it ran.

Dropping the planned work is the cheap outcome, so record why in the issue and the PR body rather than deleting it silently — otherwise the next person re-proposes it.

A routine merge from main can create the duplicate inside your own diff. The collision above lands before you write, so the duplicate is redundant on arrival. A later `main` merge is quieter: both branches were non-duplicative when they were written, and the duplicate appears only when you bring the other branch’s text into yours. Git reports a clean merge because the two copies sit in different files. Diff-scoped added-line checks do not help either, because the duplicated lines already existed on one side or the other. So after merging `main` into a prose branch, run the duplicate check against the branch’s full current diff and the neighbouring corpus, not only against lines added by the merge commit.

- **Do:** after a `main` merge, re-run a cross-file duplication check over the merged branch’s whole prose diff.
- **Do:** treat a reviewer finding on such duplication as correct even when each copy was independently right before the merge.
- **Don’t:** assume a conflict-free `main` merge preserved DRY, or that the duplicate would have appeared in an added-lines-only scan.
- **Don’t:** answer by asking which branch “introduced” the duplication; the merge introduced the state that made both copies coexist.

Two PRs that each append a new terminal numbered subsection to the same file (e.g. `### 5. ...` in a `CLAUDE.md` review-guidelines list) will conflict on merge even when neither side’s content actually disagrees. This isn’t an editorial clash — it’s two authors both writing to “the next number” at the same insertion point. Resolve by keeping **both** additions and renumbering sequentially from the collision point on, not by dropping either side; then grep the file for any other place that names the old numbering (a cross-reference, an index). This is also a reason `fully-clean`’s CI-green-and-review-clean verdict is a snapshot, not a mergeability guarantee — `main` can pick up its own append in the same spot after your last review round, so a PR can go from “reviewed clean” to “needs a merge conflict resolved” with no defect in its own diff. Before reporting a PR ready to merge, re-check with `git fetch origin main` plus the `git merge-tree` command from `resolve-conflicts`, not just a cached mergeable flag or an earlier green CI run.

After merging a PR that extracts an inline block into a reusable unit (a composite action, a shared script/function), check other open PRs that still edit that same

inline block — your merge just broke their textual diff, even though their intended change is usually trivial to re-apply to the new location. This is the mirror image of the case above: there, you’re the one resyncing after `main` moved a copy of your logic; here, *you* are the one who moved the logic, so the burden of noticing and fixing the resulting conflict falls on you, not on the sibling PR’s author waiting to hit it. Don’t wait for that PR’s own merge/CI to surface the conflict — check every open PR touching the same file right after your extraction merges: `git merge-tree "$(git merge-base origin/main origin/<sibling-branch>)" origin/main origin/<sibling-branch>` (or `gh pr diff <N>` against the new `main`) shows whether it still applies cleanly. Re-apply the sibling PR’s actual semantic change (not a mechanical `--theirs`) to the new location, verify with a direct diff that the extracted unit now differs from `main` by exactly that PR’s intended change and nothing else, then push to their branch and flag what you did in a PR comment.

See [sync-with-main.cases.md](#), “Check other open PRs after merging an extraction”.

That “push to their branch” is scoped by standing, not only by cause. `gha#201/#202` were CI workflow files in a repo the author drove, where a push saves the sibling’s author a round and risks nothing they were relying on. The same push onto a branch you do not own — a colleague’s active work, and most sharply a release branch carrying an out-of-band process — can disrupt something a comment would not. There, name the extraction, the deletion, or the rename and where the content went in a PR comment, and leave the push to whoever owns the branch. Causing the conflict obliges you to *surface* it. It does not by itself license editing someone else’s branch. See [batch-merge-and-resolve](#), “A conflict your sweep found is not a conflict your merge caused”, for the attribution step that says which conflicts are yours in the first place.

An add/add conflict on a *shared config file* usually means two PRs independently fixed the same root cause — reconcile the reasoning, don’t just pick a side. This generalizes the skill-file case above beyond skills: a repo-wide CI/lint/build config fix (a new tool config file, a workflow tweak) is exactly the kind of change multiple sessions or bots are likely to attempt in parallel once a check starts failing on `main` for everyone. When the conflict is a whole-file add/add (not just competing edits to an existing file), read both sides’ reasoning — code comments, commit messages, the PR discussion — before resolving; usually one side’s explanation is more complete (covers a case the other missed, cites the tool’s actual constraint) and should win outright rather than mechanically merging fragments of both. Re-diff the PR against `origin/main` after resolving to confirm the PR’s remaining changes are its own original scope, not a reintroduction of what the other, now-merged PR already added.

A dirty mergeable_state on a bot-opened PR can mean a sibling PR already closed the same issue, not just that main drifted. An issue-triggered `@claude` workflow can fire twice on the same issue in quick succession (a duplicate dispatch, or two people independently routing the same request), producing two independent PRs that both fully resolve it — including adding the identical new file. The second PR’s merge conflict is an add/add on that new file, and it looks like ordinary main-drift, but treating it that way and mechanically resolving in favor of “ours” silently reintroduces a duplicate the other PR’s merge already published.

Before resolving, check the PR’s linked issue for **other** cross-referenced PRs/closing events — if one already merged and closed it, diff the conflicting file against **main**: if it’s the sibling PR’s already-published version, keep **main**’s content and keep only this PR’s genuinely distinct remainder (a piece the sibling PR never did), rather than re-adding a second copy.

The same parallel resolution can be a whole-file split, and then files can vanish from your diff with no deletion hunk to read. The add/add and duplicate-issue cases above both say to keep **main** when a sibling PR already published the same new file. The split case adds a second check, because resolving the one conflict can also make other files disappear from your PR’s diff entirely. Those files look harmlessly gone, and there is no deleted line for **ardi**’s pre-push deletion sweep to inspect. Two causes are indistinguishable from the final diff alone: **main** absorbed your cross-reference edit, or the merge dropped your work. So verify each vanished file against the pre-merge head before calling the collapse correct. For each file that left the diff, compare the original head against the merge-base to recover what your branch intended, then confirm current **main** now carries that same change. Search **main**’s whole corpus for that change rather than the path it used to live at: a sibling PR that relocated the content — a companion-file split, a rename, a section moved between files — leaves it present but elsewhere, so a path-scoped confirmation reports it missing and invites you to re-add a copy **main** already has. **ardia**’s **Superseded** terminal state carries the measurement. Only after that per-file check is it safe to treat the smaller diff as a successful conflict resolution rather than as lost work.

- **Do:** save or read the original pre-merge head, list the files that left the PR diff after the merge, and verify each one’s intended change is already on **main**.
- **Do:** keep **main**’s version for the overlapping split file when the sibling PR has already published the same refactor, then carry forward only this PR’s distinct remainder.
- **Don’t:** infer that a vanished file was safely absorbed merely because the final diff got smaller.
- **Don’t:** rely on the deleted-lines sweep for this case; content that left the diff has no deletion hunk for that sweep to show.

The same silent reversion happens one line at a time, and there the file never vanishes from the diff at all. The bidirectional check above triggers on a *file* disappearing from the PR’s diff. A merge can revert a single sentence inside a file that stays in the diff for other reasons, and then nothing about the diff’s shape changes to prompt a look. **main** resolves the conflicting region toward its own side, the branch’s fix is gone, and the restored text is byte-identical to **main**’s copy — so `git diff origin/main...HEAD` for that region shows **zero lines**, the same as if the branch had never touched it.

This is the near-miss worth naming: reviewing the PR’s diff against **main** feels like reviewing what the merge did, and it is not. The PR diff answers “how does the branch currently differ from **main**”, and a reverted line that now matches **main** again answers that question with nothing, by construction. The only comparison that can see the revert is the branch’s own pre-merge tip against the merge result — a different diff than any ordinary review runs.

```
git diff <pre-merge-branch-tip> <merge-commit> -- <file>
```

A line that appears there as a **deletion**, with no corresponding line re-added elsewhere in that hunk, is a fix the merge discarded. “No conflict marker” says nothing about this case either: the region may never have raised a <<<<<< at all, since a merge is free to resolve a three-way diff silently in favor of one side when git’s own heuristics call it unambiguous.

The general form is **batch-merge-and-resolve**’s survival check, generalized past whole files: for every line the branch added between the merge-base and its pre-merge tip, confirm that line — not just the file it lived in — survives unchanged in the merge result. A whole-file diff-stat delta cannot substitute, since the file’s total line count can be unchanged or even larger while one specific sentence inside it was swapped back to **main**’s wording.

- **Do:** diff the branch’s pre-merge tip against the merge commit, per file, and treat any line that appears only as a deletion there as a discarded fix.
- **Do:** run this check even when the file in question is still present and still shows other changes in the PR’s diff against **main**.
- **Don’t:** treat `git diff origin/main...HEAD` as evidence about what a merge did — a reverted line that now matches **main** produces no diff there, which is exactly what makes the revert invisible.
- **Don’t:** read “no conflict marker in this region” as evidence nothing was discarded; a silent three-way resolution needs no marker to still pick the wrong side.

(Morrison-Lab/ai-config#2243, 2026-08-25/26: commit 6a537734 fixed a sentence in **CLAUDE.md** (“This webhook-driven loop never formally invokes...” to “This subscription never formally invokes...”). The later merge of **main**, commit 061bffcc, resolved that region toward **main**’s side and restored the pre-fix wording verbatim, byte-identical to **main**’s copy. An adversarial re-review caught it only via `git diff b42b344d 061bffcc -- CLAUDE.md`, the pre-merge-tip-to-merge-result comparison above, and fixed it in c0c658d2. Tracked as [ai-config#2374](#).)

When the whole PR is superseded, not just one file, the conflict is telling you to close it rather than resolve it. The two cases above keep **main**’s version of a file a sibling PR already published, and carry forward the current PR’s distinct remainder. The remainder can be empty. When a **main**-merge conflict pits *every* added line of an idle PR against a better-formatted copy already on **main** — a sibling PR having landed the same content — resolving toward **main** leaves nothing, and the PR’s own prior review findings are moot. Confirm by grepping `origin/main` for the PR’s distinctive added phrases before resolving anything: all present means superseded. The right action is then to recommend closing the PR, since its content is preserved on **main**, not to push an empty diff to a clean verdict. For an ARDIA sweep this is a terminal state of its own — see [ardia](#)’s **Superseded**, which also gives the up-front check that catches it before rounds are spent.

A merge into a growing numbered list (e.g. [gha](#)’s **CLAUDE.md “Code review guidelines” section) can produce zero blank lines between two adjacent headings even with no**

textual conflict — lint catches it, git doesn't. When a section is a hotspot several PRs independently append items to (each PR adding its own `### N.` block at the end), a clean three-way merge can still splice one PR's closing line directly against the next PR's heading with no blank line between them — this doesn't produce a `<<<<<<` conflict marker (git resolves it as a straightforward insertion), so it's easy to push without noticing. `markdownlint`'s MD022 (blanks-around-headings) is what actually catches it, as a CI failure with no proximate code change to explain it. Re-run the repo's markdown lint (or at minimum re-read the diff around every `### N.` boundary you didn't personally write) after any merge that touches a shared growing list, not just after a merge with conflicts.

The same splice happens to LIST ITEMS, and there `markdownlint` most likely does NOT catch it — so nothing turns red at all. The case above is a heading spliced against preceding text, which MD022 decides. The changelog case is a *bullet* spliced onto the previous item's continuation line:

```
`data-raw/precompute-true-effects-chunk.R` (#429).  
* The `docs` workflow's "Build site" step no longer times out intermittently.
```

That is a valid **tight** list item, so a list in which every other entry is blank-line separated silently starts mixing tight and loose items and renders inconsistently. `markdownlint`'s blanks-around-lists rule governs the boundaries *of* a list, not the gaps *between* its items, and no default rule enforces consistent looseness within one list — so unlike the heading case, CI stays green and only a human reading the merged section notices.

Check it mechanically instead of by eye; one line decides it:

```
awk 'prev !~ /^[[:space:]]*$/ && /^[*+\\-] / {print FILENAME":"NR": "$0" {prev=$0}}' NEWS.md
```

Use `[[[:space:]]*` rather than a bare `/^$/` — a whitespace-only preceding line is not a violation and produces false positives. The pattern `[*+\\-]` covers all three common Markdown unordered-list markers; `^\\*` alone would miss `-` and `+` bullets.

Two consequences. Run this after any merge into a changelog or other growing bulleted list, alongside the heading check above. And note that a `merge=union` driver on such a file (see [configure-gitattributes](#)) *increases* the rate of this defect, since union resolves an append collision by keeping both sides with no conflict to review — so confirm a detector is wired into CI before enabling one, not after.

Run that check as a whole-file count, and compare it before and after — scoping it to the lines you added cannot see this defect at all. The check above is the right instrument; the natural way to apply it is the wrong one. Having found the file's spliced bullets, the obvious next question is which of them are yours, and the obvious way to answer is to intersect them with the lines the branch added. That question is unanswerable, because the defect is a **deleted blank line** before a bullet that was already there. The bullet is *context* in

the diff, never an addition, so the intersection is empty by construction and the check reports a confident zero.

Note how this differs from the scope failures elsewhere in this corpus, where a check’s **inputs** were too narrow — a glob, a missing flag, a two-dot range. Here the inputs were right and the **question** was wrong, which no widening fixes. And it fails in the direction that reads as an all-clear, on the one file a reviewer will not re-derive.

The sound form is a count delta over the whole file, which needs no judgment about ownership and no diff at all:

```
git show origin/main:NEWS.md | awk '...' | wc -l # before
awk '...' NEWS.md           | wc -l             # after
```

A merge must not increase the count. That is an [algorithmatize-checks](#) instrument in the strict sense — two integers decide it — and it holds whoever authored the surrounding lines.

Generalize past changelogs, because the property is about the defect rather than the file: **when a defect can be introduced by deleting a line, any instrument keyed on added lines is unsound**. Ask instead whether a whole-file measurement got worse. The version-parity rule above is the same shape — a conflict-free merge leaves the branch at parity with `main`, `version-check` goes red, and there is nothing in the diff to point at — which is why that rule is a direct comparison of two `DESCRIPTION` versions rather than a diff inspection.

The shared trigger is the practical part: **a conflict-free merge is exactly when nothing prompts anyone to look**. Both defects arrive through one, both are invisible to diff-scoped checking, and the merge reports success.

- **Do:** measure the whole file before and after a merge, and treat any increase as the merge’s fault regardless of who wrote the lines.
- **Do:** ask, of every diff-scoped check, whether the defect it targets could be caused by a deletion — and replace it with a count if so.
- **Don’t:** intersect a whole-file finding with the branch’s added lines to decide ownership; for a deletion-caused defect that always returns zero.
- **Don’t:** read a conflict-free merge as a merge that changed nothing beyond what the diff shows.

A commit claiming “I’ve pulled main and resolved the merge conflicts” can be lying — verify it actually merged before trusting the claim. A genuine conflict-resolution commit is a merge commit (two parents); a commit that just hand-edits files to *look* resolved, without running a real `git merge`, is an ordinary single-parent commit — and it never actually incorporates whatever new state of `main` prompted the “resolve conflicts” request in the first place. This is easy to miss because GitHub’s own `mergeable/mergeStateStatus` fields don’t distinguish the two: both look identical from the PR page until you check the commit graph. Verify with `git show -s --format="%P" <commit>` — one hash means no real merge

happened, regardless of what the commit message says. If a branch needed conflict resolution more than once and each attempt claimed success but the PR still shows **CONFLICTING**, check every “resolved conflicts” commit in its history this way before trying yet another resolution attempt on top of a foundation that was never actually re-merged.

4 Driving a Pull Request to Clean

Whenever you are working a PR/MR, run the full **ARDI** loop by default, without being asked: **A**ddress every flagged item, **R**ebut findings that are wrong, **D**efer out-of-scope items to tracked issues, then **I**terate with a fresh review — repeating until every reviewer’s latest verdict is **fully clean**. Don’t stop at “review-clean, just needs approval” and hand triage back; keep the cycle going until it’s genuinely clean.

Posting a review as comments, with no request to edit the branch, is not working the PR. Do not start ARDI, do not push fixes, and do not merge. Leave the findings and stop unless asked to iterate. A later request to iterate is a driving request. “Watch and ARDI every PR you touch” applies when you are driving the branch, not when you were asked only to read it. (UCD-SERG/shigella#31, 2026-08-25.)

Extended rationale — the mechanism, evidence, and argument behind each rule below — lives in [ardi.rationale.md](#), moved out of the auto-loaded context. Each rule here keeps its statement and its Do/Don’t pair; read the companion when the reasoning or the evidence is the question.

Worked-example case records for the rules below live in [ardi.cases.md](#), moved out of the auto-loaded context.

Continuously monitor every PR/MR you are actively working until it reaches that terminal state.

That wait is conditional on a run having been scheduled, and on some repos a push schedules nothing.

- **Do:** read the review workflow’s **on:** block before the first push to a PR in an unfamiliar repo, and dispatch explicitly after the round’s last push when it carries no push-based trigger.
- **Do:** treat a non-zero `check-pr-fully-clean.py` on a dispatch-only repo as a prompt to dispatch.
- **Don’t:** read green CI at the current head as evidence a review is in flight — on a dispatch-only repo that is the steady state, not a transient one.
- **Don’t:** let a verdict from an earlier head stand because the repo’s trigger class was already known. Knowing it is not the same as acting on it each round.

Dispatching needs no permission, so do not ask about the spend.

The rule above is about the mechanism and says nothing about authorization, which leaves a session free to know it and still stall — by reasoning that a review round costs money and the spend is the maintainer’s call. That sounds like restraint and is indistinguishable from it from the inside.

The asymmetry runs the other way. A green, unreviewed PR is **parked, not clean**, so declining to dispatch does not save a round; it holds the PR in a state that reads as finished and is not. The stall also spends the user’s attention every time, which is the thing the review loop exists to conserve.

So dispatch when the round is ready, and put the run in the status report rather than the question.

- **Do:** dispatch the review yourself once the round’s last push has landed, on any repo whose reviewer is dispatch-only.
- **Do:** name the run you are waiting on, so the report carries a fact rather than a request.
- **Don’t:** write “spending a round is the maintainer’s call” into a status report, or hold a ready PR pending a spend question.
- **Don’t:** read this as a general spending grant — it covers scheduling a review, and merging is still [mwc](#)’s to govern.

(Directive from the user, 2026-08-16: “always dispatch”. Three PRs reached green CI on a `workflow_dispatch`-only repo in one session, and each time the session asked before dispatching, citing rounds that had billed \$12.14, \$10.37 and \$12.44 against a monthly limit already reached. Both earlier dispatches came back **Needs more work**, one with a blocking correctness bug, so the round was not a formality. Tracked as [ai-config#1571](#).)

Dispatch once, after the round’s LAST push — a per-push rhythm cancels its own reviews.

Dispatch with `--ref <PR-branch>`, or the resulting failure is invisible on the PR.

- **Do:** finish pushing, then dispatch once, and name in the status report which run you are waiting on.
- **Do:** pass `--ref <PR-branch>` on every dispatch, so the review’s check runs attach to the PR head.
- **Do:** diagnose a missing verdict by reading the run, since a cancelled dispatched run leaves no trace on the PR’s check-run list.
- **Don’t:** dispatch per push — each one cancels the last, and the round spends review time producing nothing.
- **Don’t:** re-dispatch reflexively when a verdict is missing. If one is in flight, the retry cancels it.
- **Don’t:** read a green, nothing-pending PR as reviewed on such a repo — that is also what an invisible cancelled gate looks like.

A cancelled run is invisible to a session `READING` the PR and loud to one `SUBSCRIBED` to it, and the second is the dangerous direction.

The bullet above is scoped to the check-run list deliberately. GitHub lists check runs for the **head commit**, so a run cancelled after the head moved leaves a **require-review** failure hanging off the superseded SHA, where a session reading the PR will not find it.

The webhook stream is a different surface and it does not filter that way. The cancel fires a **check_run.completed** with **conclusion: failure**, so a session subscribed to PR activity is woken by a red **required** check on its own PR.

That inverts the risk the bullet above describes. An invisible failure costs you a verdict you thought you had, and you find out by looking. A visible failure on a superseded commit costs more, because the drive-to-green posture says not to end a CI-failure wake without pushing a fix or replying with a blocker — so the reflex is to fix, against a commit that is no longer in the PR’s timeline. At best that is wasted work. At worst you change the current head on the authority of a red check that was never about it.

One field decides it, and it is the field **fully-clean** already names for the neighbouring problem: compare the event’s own **head_sha** against the PR’s current head. Equal means act. Unequal means confirm a run is live at the real head, and leave the diff alone.

- **Do:** compare a CI-failure event’s **head_sha** against the PR’s current head before diagnosing anything.
- **Do:** reply naming the superseded SHA rather than staying silent, so the wake is visibly dispositioned rather than dropped.
- **Don’t:** read “leaves no trace on the PR” as covering the webhook stream — it describes the check-run list, which is filtered by head commit.
- **Don’t:** push a fix in response to a red check whose **head_sha** is not the head; the check is not about the code you would be changing.

See [ardi.cases.md](#), “A cancelled dispatch that fired a failure webhook against the superseded SHA”.

See [ardi.cases.md](#), “A per-push dispatch cancels its own review, invisibly”.

A dispatch you make while idle can still be cancelled — by a push you did not make. The per-push rule above governs your own pushing rhythm colliding with your own dispatch. The concurrency group is keyed on the PR number alone, not on who triggered the run or how, so a **third party’s** push into the same PR branch cancels your dispatch exactly as a push of your own would — even when you have pushed nothing since dispatching. On a repo with an active **@claude** agent, the bot itself is such a third party: reacting to review activity, it can push a **main-sync** merge into the PR branch (see [claim-pr](#)’s note on this), which is a **synchronize** event and, on a repo that reviews on push, schedules a fresh review run in the same group — cancelling yours.

The tell: your dispatched run reads **cancelled**, its **headBranch** is whatever ref you dispatched against — the default branch if you omitted **--ref**, the PR branch if you passed it, so this conjunct alone proves nothing — and a newer **pull_request**-event run exists for the same PR, at a newer head. The right response is to do nothing — the cancellation is benign, since the newer run supersedes at a better head, and re-dispatching would cancel *that* one instead.

- **Do:** before treating a cancelled dispatch as a lost review, check for a newer **pull_request**-triggered run on the same PR. Its existence and its head are the whole explanation.
- **Do:** re-fetch the PR branch head before concluding a dispatch failed — a head that moved without you pushing explains a cancellation the per-push rule above cannot.
- **Don't:** re-dispatch to “fix” a cancelled run that a newer, better-placed run already superseded — that cancels the survivor instead.
- **Don't:** read **require-review: failure** on the cancelled run as a live problem. It is a side effect of the superseded run, and the newer run supplies the real verdict.

See [ardi.cases.md](#), “A third-party push cancels an idle-dispatched review”.

Always execute `python3 scripts/check-pr-fully-clean.py --quorum <number-of-reachable-providers> <pr>` synchronously in the foreground turn to evaluate clean verdicts. Whenever ending a turn while waiting for an AI review or CI completion on an active PR after pushing code, launch a **schedule** timer (e.g. 120s) to check back. When the timer fires: - Check if a review for the HEAD SHA has arrived. - If no review has posted yet, verify whether review workflow runs are still in progress (`gh run list / gh pr view --json statusCheckRollup`). - If review workflows are still running: schedule another timer to check back. - If the reviewer failed, was canceled, skipped with no replacement (e.g. quota limit), or produced a stub review with no stated verdict: invoke self-review fallback per [self-review-fallback.md](#) rather than stalling the loop. - Otherwise, fix any underlying workflow or dispatch issues discovered along the way and schedule another timer to maintain continuous monitoring until a review lands, self-review fallback triggers, or CI completes. This applies transitively to PR-driving workflows such as **gi**, **gii**, and **ardia**; only monitor PRs the session owns or has explicitly claimed, so the rule does not authorize changing someone else's work.

The loop's terminal action is to **report the PR ready, not to merge it**. Merging is human-gated — it happens only on an explicit human “merge it” (the **merge-it** skill), never as a step ARDI takes on its own. So when you carry a PR across a **ScheduleWakeup** or **/loop** wait, **never** bake a self-merge directive like “if clean and CI green, merge it” into the wakeup/loop prompt: a scheduled prompt fires back as a user-role turn, so a self-authored “merge it” only *looks* like human approval (and Claude Code's auto-mode classifier will rightly deny it as a self-authored merge). Drive to fully clean, report ready, and leave the merge — and any other destructive one-off, e.g. a `gh workflow run` that force-pushes — for explicit human authorization.

Because the loop ends there, **the clean verdict remains a UMS checkpoint** — don't hold that pass for the merge, which is on the human's clock rather than this session's and may land after a **/clear** or not at all. **Reading the review is an earlier checkpoint than the**

verdict. Run UMS when the review arrives (Rebut and Defer included), not only once every finding is Addressed or the round comes back clean. See `CLAUDE.md`’s “Run UMS proactively, as learnings accumulate”; the merge-time pass in `post-merge` then only has to cover what the merge itself taught.

The one exception: if the human has explicitly granted the `mwc` (merge-when-confident) session permission, that grant is a live human instruction, not a self-authored one, so baking a self-merge step into a wakeup/loop prompt is fine for the rest of that session. See `mwc` for the grant’s scope and limits. The Scope Limit still binds: a disagreement among reviews — one all-clear and another with findings, nits included — is not fully clean, so `mwc` does not authorize merging it. ARD every item from every review, then request fresh reviews (`ai-config#2274`).

A scheduled check-in can outlive the PR it names, and its stale premise arrives as a user instruction. Nothing reviews a wakeup prompt between authoring and firing — a PR body gets read by a reviewer, a changelog entry sits in a diff, but a scheduled prompt is written at T, stored, and delivered at T+N with no intervening reader. And it arrives as a user-role turn, the most authoritative framing a turn can receive, which makes a stale premise persuasive: the turn opens by telling you what the situation is, in the voice of an instruction. If the PR merged between arming and firing, acting on “drive #N’s findings” means working a branch GitHub has already auto-deleted — the orphaned-branch recovery `use-existing-pr-branch.md` documents, arrived at from the other direction.

- **Do:** word a wakeup against a re-derivable set (“re-check every PR this session opened that has not merged or closed”) rather than a fixed identifier, so it survives any number of merges in the gap.
- **Do:** treat a named PR’s `state/merged` as a claim from a past self, and verify it (`gh pr view <N> --json state,mergedAt / pull_request_read get`) before acting on anything the prompt asserts — `merged: true` means the correct action is `post-merge`, not another ARD round.
- **Don’t:** hard-code a PR number, or a head SHA, into a wakeup that may outlive either.
- **Don’t:** treat a wakeup’s premise as current because it arrives in the user role — it is a message from a past self, and the state it describes is as old as the prompt.

(`ai-config#902`, 2026-07-30: a `send_later` check-in named PR #873, which merged between arming and firing; the wakeup arrived asserting a live PR needing an ARD round when the actually-live PR was #892, its own UMS follow-up.)

A head SHA is the same defect on a far shorter clock, and it fails more quietly.

A stale PR number still names something that exists, so verifying it returns an answer that disagrees with the prompt. A stale SHA names a commit that is simply no longer the head, so a check-in that verifies the verdict against it confirms a review of superseded content and reports success. The clock is the ARD round rather than the merge queue: on `ai-config#2623` the two commits either side of one review round were 5m38s apart, and the check-in naming the earlier one was armed 1m23s before the later push landed — so it was obsolete before it

ever fired. Word the check-in to *derive* the head (“fetch the current head SHA”) rather than to carry one.

In the **clear-all family** (*ardia*, *gia*, *gii*, *gip*), “report ready, don’t merge” gates only the merge — it does **not** pause the sweep. A clean-but-unmerged PR is not a stop; move to the next item, and stack it when it isn’t naturally independent of that PR. See [stack-dont-pause](#).

The same gate does not pause the loop *within* a single PR either, and that is the harder half to see.

The tell is lexical, and it sits in your own outgoing message: a RECOMMENDATION or question whose proposed action is ordinary ARDI work.

- **Do:** resolve conflicts, sync, push fixes, and re-dispatch reviews on a PR whose merge you are correctly withholding, and report those in the past tense.
- **Do:** name the one action that is gated, so “blocked” stays a claim about a step rather than about the PR.
- **Don’t:** generalize a withheld merge into withholding the rest of the loop as though one authorization covered both.
- **Don’t:** write a recommendation proposing work ARDI already mandates — a request to do the required thing is the error, not a courtesy.

See [ardi.cases.md](#), “A merge gate is not a work gate”.

Self-review against the project’s own stated conventions before every push, not just the first — and don’t just re-read the criteria, actually run the applicable review skills against your own diff and iterate on what they find, the same ARD cycle you’d run against an external reviewer’s findings. Don’t treat the review bot as the mechanism that discovers a project’s documented conventions — self-apply them first.

See [ardi.cases.md](#), “A review round surfacing five findings your own conventions already covered”.

4.0.1 Pre-push checklist

Pause point: after committing, before `git push`.

- **A separate adversarial-reviewer subagent reviewed this diff and returned a clean verdict** — dispatched in the foreground against `git diff origin/<default-branch>...HEAD`, briefed with the standards rather than with your rationale for the change, with every finding Addressed, Rebutted, or Deferred to a tracked issue, and re-dispatched after the last commit so its `Reviewed-Commit:` fingerprint names the commits the push would ship ([adversarial-self-review](#)). An inline pass under a reviewer framing does not satisfy this, and reads identically in the output — the test is whether an `Agent` call was made.

- **The whole test suite ran**, not the files you predicted the change touches, and the tests/failed/skipped triple was read — a non-trivial skip count means re-running with the gating flags set (`NOT_CRAN=true`, and whatever else un-gates a conditional skip).
- **Generated trees were regenerated** if the diff (or a `main` merge) touched a generator’s inputs, and the PR body states how many changed files are generated.
- **Added lines were scanned** for banned punctuation and multi-sentence lines, run *after* committing, *after* every pass that edited the diff (your own reflow included), and with the three-dot range (`origin/main...HEAD`) — a pre-commit run reports on the wrong tree, a later edit retires the lines an earlier run scanned, and a two-dot range re-attributes whatever `main` deleted to you. The mirror direction fires at merge-decision time rather than at line-scanning time, and reads as an emergency: a two-dot `git diff origin/main` shows whatever `main` **added** as *your* deletions, so a behind branch appears to be reverting a sibling PR that just landed. It is not. A merge is three-way, so a file this branch never touched keeps `main`’s version, and a deletion count in an untouched file means the branch is behind rather than dangerous. Settle it with `git merge-tree --write-tree origin/main <head>` and read the resulting tree, rather than acting on the two-dot diff.
- **The changelog entry and EVERY PR description this round touched were re-read** against the new behavior, not just the code — none is in the diff, so no reviewer and no `grep` will catch a stale one. Read “every” literally: a round that corrects a claim appearing in two PRs’ bodies discharges the *feeling* of having synced bodies as soon as one of them is done, and the one most likely to be skipped is your own, because fixing the other repo’s copy is the part that felt like the work. This fires on a **prose** diff too: a body that explains the claim the round just walked back is stale in the way that matters most, and “reconciling prose” does not feel like changing what the PR does. Every **number** in the body was *re-derived* by command rather than re-read, run *at this push* rather than carried from the last one, with the command pasted beside it — a wrong count reads exactly as plausible as a right one, so reading is no instrument for it, and a base figure owes its own derivation rather than riding on the delta’s. “At this push” includes a push that only answers a self-review finding, and it includes the figures a “Corrections to this body” entry already refreshed — that entry is a claim about the previous head, so the current push is what expires it. A figure whose deriving command carries a **precondition** owes one more step, because deriving it freshly discharges “don’t recall it” and says nothing about whether the command was right for this diff: cross-check it against a quantity computed by something else (`git diff --shortstat` against a hand-run added-lines count), since re-reading a correct-looking pipeline confirms it (**fail-fast**).
- **The diff’s deleted lines were read** (`git diff origin/main...HEAD | grep '^-'`), and each one was a decision rather than collateral from an edit’s blast radius — a reviewer reads every deletion as deliberate and will rationalize an accidental one.
- **main was merged in** if it moved, with version parity re-checked afterward, so the round costs one review run rather than two — and any whole-file count a merge can worsen (spliced changelog bullets) compared before against after, since a defect caused by a *deleted* line is invisible to every added-lines check (**sync-with-main**).

- **Killer item: the push landed.** `git rev-parse HEAD origin/<branch>` agree before any reply asserting a fix. This one is marked because its failure is not an omission but a **false claim about state**, which a reviewer has no reason to doubt: CI reports green because it correctly validated the older head, and the session's own recollection agrees with the reply. It answers whether the **branch** moved, and nothing about whether the **PR** is still open — a closed PR keeps accepting pushes and stops tracking its branch, so both SHAs agree while the PR's own head stays frozen. Read the PR's **state** as a second check, per [use-existing-pr-branch](#), rather than letting this item stand for both.

Review a round's fixes as one diff, not as N independent fixes: two of them, each correctly addressing its own finding, can compose into a defect neither introduces alone.

- **Do:** re-read the round's full diff as a unit once every finding is addressed, before running the pre-push checklist.
- **Do:** treat a multi-item commit message as the prompt to check the items against each other.
- **Don't:** conclude the round is sound because each finding's fix is sound; that is a claim about the parts.
- **Don't:** rely on the next review round to catch it — it may, but it is then spending a round on something the previous round created.

See [ardi.cases.md](#), “Two correct fixes composing into a defect neither introduces alone”.

A clean verdict does not certify that your diff contains only what you meant, because a reviewer cannot tell an accident from a decision.

- **Do:** read your diff's deleted lines before pushing, and confirm each one was a decision rather than a casualty of an edit's blast radius.
- **Do:** say plainly, in the thread, when a review has blessed something unintended — the reviewer cannot know, and its verdict will otherwise stand as the record.
- **Don't:** treat a clean verdict as evidence about intent; it is evidence about correctness only.
- **Don't:** keep an unintended change because the reasoning offered for it turned out to be good.

When the edit is a regex or string patch rather than the Edit tool, two mechanisms turn that displacement into a silent over-deletion, and a self-check can wave both through.

- **Do:** anchor a string or regex patch on text unique to the intended site, and prefer the Edit tool's exact-string matching over a broad `.*?` DOTALL span.
- **Do:** make a patch self-check assert that neighbouring structures survive — the sibling function or loop still present, untouched element counts unchanged — not merely that the changed element's count is as expected.

- **Don't:** trust a `re.sub(..., flags=DOTALL, count=1)` whose non-greedy `.*?` start anchor is non-unique; it binds to the first occurrence.
- **Don't:** read “the assertions passed” as “the patch is correct”; a count-based check can pass by coincidental balance, and the `git diff` deletion-review is the real gate.

A clean verdict does not discharge the self-review against project conventions either, and the reviewer's own “not a finding” is where that shows up.

- **Do:** re-run the project-conventions check against your own diff after a clean verdict, not only before the push — dispatched to the [adversarial-reviewer](#) subagent like any other self-review ([adversarial-self-review](#)), rather than performed inline.
- **Do:** read a reviewer's “observations” and “not a finding” items as candidate violations, and `grep CLAUDE.md` for whatever they discuss.
- **Don't:** let a reasoned “belt-and-suspenders is fine” settle a question the repo already answered in writing.
- **Don't:** treat a non-blocking label as deciding whether an item gets checked at all.

Proactively self-correct a technical claim you already told a reviewer, the moment further testing shows it was wrong — don't wait for the reviewer to catch it. If you stated a rationale (an approach is safe, a risk doesn't apply, a backstop exists) and then discover through your own follow-up verification that it's false, post the correction with the actual evidence immediately, rather than leaving the stale claim standing until a review round re-raises it. This keeps the review loop converging instead of churning on a claim you already know is wrong.

See [ardi.cases.md](#), “Self-correcting a rationale before the reviewer re-raises it”.

A fix is not “pushed” until it is on the PR's head commit — verify with a SHA comparison before telling a reviewer you pushed it. From inside a session, an edited working tree and a pushed commit feel identical, so a round that edits the files, writes the reply, and never runs `git push` produces a reply asserting a fix that does not exist on the branch. Nothing contradicts it: CI reports green, because it correctly validated the older head; the next review round reviews code without the fix; and the session's own recollection of having made the change agrees with the reply. That makes it worse than an ordinary wrong claim — it is a false statement about *state*, which a reviewer has no reason to doubt and no cheap way to check.

A SHA you put in a PR body or a reply must be read, never recalled — and the PR body is where an invented one survives longest.

Knowing the prefix genuinely does not discharge this for the full SHA a link wants.

- **Do:** read every SHA you cite out of `git rev-parse` or `git log`, and confirm it resolves before pasting it.
- **Do:** correct a published wrong SHA with a visible note naming the real one.

- **Don't:** write a short SHA from recollection because it looks like the commit you just made.
- **Don't:** extend a genuinely-read short prefix into a full SHA by hand — the prefix discharges nothing for the 33 characters it does not contain.
- **Don't:** expect review to catch it — a reviewer has no reason to suspect a citation, and the body is not in the diff they are reading.

See [ardi.cases.md](#), “A genuinely-read prefix, extended into a fabricated link”.

The same rule governs a merge or squash commit message, which is worse than a PR body on both counts the bullet above names.

The trigger is a PR with no closing issue, which is why “verify identifiers” does not reach it.

An invented number here can close someone else's live work, because issues and pull requests share one number space.

- **Do:** read any Closes/Fixes/Refs number in a merge message out of the PR body it came from, and confirm the target is what you think it is.
- **Do:** say plainly that a PR closes nothing when it has no tracking issue, rather than leaving the slot empty for a number to fill later.
- **Do:** correct a published wrong reference visibly, in a comment, since the message itself cannot be amended once it is on the default branch.
- **Don't:** treat a closing keyword as inert against a pull request — it closes one, and the number space is shared with issues.
- **Don't:** infer from “nothing changed state” that a wrong reference was harmless; check whether the target was open.

See [ardi.cases.md](#), “An invented Closes in a merge commit message” and “A negated closing-keyword sentence still closes the issue”.

A SHA's provenance is the question its source command answers, not merely that a command produced it.

- **Do:** name the claim, then run the one command that answers it, and paste that command's output.
- **Don't:** lift a SHA out of nearby command output because it is genuine and close at hand — `git stash list` and `git reflog` answer about their own subjects, not about the tip.

See [ardi.cases.md](#), “A read SHA can answer a different question”.

A verification table you write in the PR body is the same defect one artifact over, and re-reading it cannot catch a wrong number.

It goes stale rather than being wrong on arrival.

It sits in the PR body, which nothing re-reads.

- **Do:** re-derive every count in the PR body with a command at push time, and publish the command next to the count.
- **Do:** treat any round that changes the diff as expiring every figure the body already states, not only the figure that round was about.
- **Don't:** substitute re-reading for re-deriving — re-reading is the right instrument for a stale description and no instrument at all for a stale number.
- **Don't:** report a delta without deriving its base; a base carried from recollection is unfalsifiable by any later check of the delta.

See [ardi.cases.md](#), “A verification table in the PR body going stale as rounds change the diff”.

A reviewer’s round-one confirmation of that table does not expire when the diff moves, and the confirmation is what makes the stale figure dangerous.

The rule above says the figures go stale. This is about the one artifact in the review record that argues they have not.

Round 1 verifies the table, in detail, because a body full of derived counts is exactly what a first review checks, and it says so, naming each figure it matched. Round 2 does not re-verify it, because round 2 is not about the table. Round 2 is about whether round 1’s findings were addressed, so the body sits outside what that round set out to read.

The confirmation is therefore a claim about one head, and nothing retires it. An unverified table at least invites suspicion. A table a reviewer explicitly confirmed reads as settled by someone other than its author, and that reading survives every push that falsifies it.

Sharper still, and this is the part worth pinning: round 2 can derive the correct new figure and use it in its own prose while the body carries the old one, and flag nothing. The reviewer is not diffing its numbers against the body’s. The reviewer derives fresh ones for its own purposes, so the two figures sit one round apart in a single comment thread, contradicting each other, with nobody comparing them.

- **Do:** re-derive every figure in the body at each push, whatever an earlier round confirmed, and record the SHA the new figures were derived at.
- **Do:** compare any figure a later review states in its own prose against the figure the body states, and read a mismatch as the body being stale.
- **Don't:** carry a round-one confirmation forward to a later head — it verified the diff that existed when it ran.
- **Don't:** read a later round’s clean verdict as evidence the body is still accurate; that round checked the findings, not the table.

See [ardi.cases.md](#), “A round-one confirmation laundering a body the next round contradicts”.

A “Corrections to this body” entry is itself a figure in the body, so the next push expires it too — and it reads as more settled than the figure it corrected.

- **Do:** re-derive every figure a corrections entry vouches for at each push, and record the SHA the new figures were derived at alongside them.
- **Do:** append a further numbered entry when a later push moves the figures again, rather than editing the previous one, so the round that expired them stays visible.
- **Don't:** read a corrections entry as discharging the figures it names — it is a claim about one commit, and the next push is what falsifies it.
- **Don't:** treat having written the correction as having done the check; the note is that check's output, never a substitute for re-running it.

See [ardi.cases.md](#), “A corrections entry expires with the next push”.

Verifying that a stale figure is gone needs a SECTION-scoped search, because the corrections entry legitimately quotes it.

The two rules above compose into a check that cannot discriminate. The table must stop claiming the superseded figure, and the corrections entry must quote that same figure in order to say what changed — so the string is still in the body after a fully correct fix, and a whole-body search for it reports that fix as having failed.

That is a check whose pass path and failure path look alike, which [fail-fast](#) says is not yet a check. It also fails in the direction that invites damage: the natural response to a “still present” hit is to delete the quotation, which is the one part of the entry carrying the record.

Scope the search to the section that makes the claim, and assert the corrections entry in the opposite direction:

```
ver = body[body.find("## Verification"):body.find("### Corrections")]
corr = body[body.find("### Corrections):]
assert "484 added" not in ver      # the table no longer claims it
assert "484 added" in corr        # the entry still records what changed
```

- **Do:** scope a staleness check to the section that makes the claim, and assert separately that the corrections entry still quotes the old figure.
- **Do:** write the two assertions in opposite directions, so a deleted quotation fails as loudly as an uncorrected table.
- **Don't:** search the whole body for the superseded figure — a correct fix leaves it present, so that search reports every correct outcome as a failure.
- **Don't:** answer a “still present” hit by removing the quotation from the corrections entry; that quotation is the record the entry exists to carry.

See [ardi.cases.md](#), “A whole-body staleness check that reported a correct fix as failed”.

The one case where a figure does **not** expire is a push that leaves the tree unchanged — a revert-and-restore returns the tree to an object it already had, and a measurement is a function of the tree rather than the commit. [dont-incur-technical-debt](#)'s “The one exception” section carries that mechanic, and the deferral it licenses.

The read side of that comparison can lag a push by a few seconds, so test the two *local* refs against each other before concluding anything failed.

- **Do:** compare HEAD against origin/<branch> first when the PR API disagrees, and re-read rather than re-push when those two agree.
- **Don't:** amend, force-push, or re-commit on the strength of an API SHA alone.

A brand-new branch can read back at the wrong commit, so the local two-ref comparison above is not sufficient there.

The gap is in the trigger rather than in the remedy.

The likeliest explanation is a local one, and it reproduces offline.

See [ardi.cases.md](#), “A brand-new branch reading back at main’s tip, reproduced offline”.

- **Do:** run `git ls-remote origin <branch>` after the first push to a new branch, and compare its SHA against `git rev-parse HEAD`.
- **Do:** run `git rev-parse HEAD <branch>` first when those two disagree, since a branch ref left behind accounts for the whole signature (and note that `--short` rejects a second revision, so pass neither).
- **Do:** re-run plain `git ls-remote` as well, so a ref that self-corrects stays distinguishable from one a re-push repaired.
- **Do:** re-push with `git push origin HEAD:refs/heads/<branch>` when the mismatch persists, and read the SHA range it prints as the confirmation.
- **Don't:** treat a `git push` that exited 0 and printed `* [new branch]` as evidence the commit reached the remote.
- **Don't:** assume `git push -u origin <branch>` sent the commit you just made – it sends the branch ref, which HEAD may have moved past.
- **Don't:** credit a corrective re-push with having repaired a remote-side fault when neither of those two controls was run.
- **Don't:** answer a `No commits between main and <branch>` error by re-checking the base branch argument before checking where the head ref actually points.

The same false claim arrives as *incoming* state when you pick a PR up mid-flight, and there the SHA comparison usually has nothing to compare.

- **Do:** run `gh pr diff <N> --name-only` against any inherited “already fixed” claim before deciding a finding is closed.
- **Do:** state plainly in your own summary that the prior claim did not hold, and name the head it was false at.
- **Don't:** treat green CI as evidence that a claimed fix landed.
- **Don't:** infer that a finding is stale because a comment says it was addressed.

Your own fix-round commit makes the same claim one step earlier, and a file-count coincidence is what lets it through. A commit message that enumerates a review round’s findings asserts that the commit’s diff touches every file those findings name, and nothing checks that assertion by default. Measured 2026-08-27 on [ai-config#2229](#): a seven-finding fix commit changed seven files, and the match read as confirmation — but the findings named eight distinct files, because two findings each spanned two files while one file was named by two findings. The commit covered seven of the eight, so the one unfixed finding hid behind the equal counts, and the next review round re-raised the finding against a commit message that claimed to have fixed it. A matching count is not a matching set. Derive the union of files the round’s findings name, compare it against the fix commit’s own changed-file list (`git show --name-only`, which prints full one-per-line paths where `--stat` may ellipsize them) before pushing, and account for every member missing from either side.

- **Do:** compare the fix commit’s changed-file list against the files the findings name, member by member, before pushing a round.
- **Don’t:** read “N findings, N files changed” as the round being covered — the count coincidence is exactly what masked the miss.

Run that same command before *any* readiness claim, not only against an inherited one — a PR whose branch carries no implementation is green on every check.

- **Do:** run `gh pr diff <N> --name-only` before reporting a PR ready, and read the returned paths against what the PR says it does.
- **Do:** treat an empty return, or a return holding only a `main` merge’s incidental paths, as the PR carrying no implementation.
- **Don’t:** count the claim commit or a `main` merge as work — neither is implementation, and both give the branch a plausible history.
- **Don’t:** read all-green CI plus a finding-free review as evidence a PR contains anything; on an empty diff that is the expected result.

When the change affects downstream consumers, validate it against a real consumer repo before reporting the PR ready — a package’s own test fixtures are built to exercise its code, not to resemble the packages that will actually use it.

- **Input shapes no fixture happens to contain.** A real package carries metadata the fixtures never needed — an entry of a different kind, an extra tag, an unusual name — so a branch written for it has never actually run on real input.
- **Message formatting under real counts.** Fixtures usually trip the plural path; a real repo hitting the same code with exactly one item exercises the singular wording, which no test asserted.
- **The migration/upgrade path, as opposed to the fresh-install path.** This is the one fixtures can never reach: a fixture is created new by the test, so it always gets the current templates. An existing consumer has the *old* config, and whether the feature

reaches it at all is a different question from whether it works. Verify the claim in the changelog by running the documented migration step, rather than describing it.

See [ardi.cases.md](#), “Validating against a real consumer repo covers what fixtures cannot”.

Verify a blocker you assert in a PR body or a reply, with the same rigor you apply to a reviewer’s claims — a stated blocker becomes a premise other people build on.

A hold SOMEONE ELSE stated is a premise you inherit, and nothing re-tests it. The rule directly above governs a blocker *you* assert, and it is about your own claims. So is the restated-blocker rule this file carries elsewhere. The mirror is the case nobody owns: a hold written by another agent or another session — in a PR body, an issue comment, a handoff note — with a reason attached, whose reason later stops being true.

It is worse than a stale blocker of your own, for a reason that has nothing to do with diligence. The natural response to someone else’s stated hold is deference, and deference costs nothing at the moment it is chosen: honouring it is always defensible, and re-testing it looks like second-guessing a peer. So the hold is read as a fact about the artifact rather than as a claim with a checkable condition inside it, and it survives exactly as long as nobody feels entitled to check.

The tell is a hold with a reason in it. “Do not merge” alone is a standing instruction and stays binding. “Do not merge, because X may not have happened” is a *conditional*, and X is checkable — usually in one query. The two look identical at a glance, because the reason reads as an explanation of the instruction rather than as its condition.

A reviewer honouring the hold is not evidence the hold is warranted. A reviewer honours instructions; it does not audit them. It can dissolve the premise itself, in the same comment, and still defer — which is correct behaviour and is precisely why it cannot serve as confirmation.

Testing the premise is yours. Lifting the hold is not, when it is another party’s, so the deliverable is a verified report to a human rather than the action the hold was blocking.

- **Do:** treat a hold carrying a stated reason as a conditional, and test the condition before reporting the item blocked.
- **Do:** report what you found to whoever can lift it, naming the condition and the evidence, rather than lifting another party’s hold yourself.
- **Don’t:** repeat “held” across status reports without having tested the premise once — the repetition is what converts a claim into a fact.
- **Don’t:** read a reviewer’s deference to the hold as an independent endorsement of it.

(Morrison-Lab/gha#678, 2026-08-26 to 2026-08-28. Its body held the PR because the dogfood review might `self_mod-skip` and “a skip is not a verdict”. Two skips did occur; then gha#674 reached v2, a genuine review ran at the exact head on 2026-08-27 returning **Ready for merge**,

and deferred to the hold anyway. The PR sat fully clean for two days, and one session reported it as “held by its own do-not-merge” in three separate status sweeps before testing the premise even once.)

Attempting the base form of a command is not attempting its variants — a refusal describes the invocation you ran, never the flag you did not try.

- **Do:** run the flag variant the docs or the error itself name, before generalizing a refusal into an impossibility.
- **Do:** scope the published claim to the invocation actually run, naming the exact command and what it exited with.
- **Don’t:** read an unconditional-sounding error as covering flags you never passed.
- **Don’t:** count an attempt at the base form as discharging the rule above for a variant of it.

Name the specific gate when you report a blocker, not a category word that happens to be one of several.

- **Do:** quote the clause that distinguishes the failure, and name the gate it belongs to.
- **Do:** re-read a blocker you have restated several times, since a paraphrase repeated across status reports hardens into the record.
- **Don’t:** use one mechanism’s own name as a generic word for its category.
- **Don’t:** treat having verified *that* something is blocked as having verified *why*.

When the blocker is a hang, inspect the process rather than re-guessing what it is waiting on.

- **Do:** read `ps -o stat=, lsof -d 0`, and the process tree before describing what a hung command is waiting for.
- **Do:** say which read produced the answer, so the gate is checkable rather than asserted.
- **Don’t:** substitute one guessed mechanism for another because a probe produced no output.
- **Don’t:** report a timeout signal as evidence about *why* something blocked; it is evidence only that it had not finished.

A blocker that was true when you published it can stop being true while the PR is open, and withdrawing it is your job, not the reviewer’s.

- **Do:** after every `main` merge, scan the PR’s touched files for merge-status hedges with whitespace-normalizing search, then re-read each hit against the new base.
- **Don’t:** assume a hedge survived because the file that contained it merged without conflicts, or because literal grep missed a phrase split across semantic lines.

Landing a fix falsifies whatever prose documented the defect, and that prose is never in your diff — so grep for it rather than expecting to be reminded.

- **Prose staled by the fix.** It was accurate when written, so nothing about it reads as a defect, and a workaround it prescribes becomes active misdirection the moment the thing it worked around is gone. Keep the entry where the old behaviour explains something — most of a corpus is written against it — but mark plainly that it is history and name the change that ended it.
- **Prose asserting conformance to a reference.** A docstring saying the code “follows” some reference implementation is a claim about two artifacts, and your own divergence falsifies it. This one is not staleness at all: it was false before you arrived, and it is load-bearing, because a reader checking the code against the reference stops at the sentence saying someone already did.
- **Do:** grep the repository for the defect, the workaround, and the behaviour you changed, before calling a fix complete.
- **Do:** mark a superseded entry as history and name the change that ended it, rather than deleting it, when the old behaviour still explains other text.
- **Don’t:** treat a clean grep over the diff as coverage — the stale prose is outside it by construction.
- **Don’t:** leave a doc asserting conformance to a reference standing when the code diverges; correct the claim in the same change that establishes the divergence.

An instruction’s own suggested code is not exempt from the project-conventions self-review above.

See [ardi.cases.md](#), “An instruction’s own suggested code breaking a project convention”.

When the code path under test has a staging or transform step between input and output, a passing unit suite is not evidence it works — exercise the real path once.

See [ardi.cases.md](#), “A staging step the unit fixtures could not reach”.

When new code branches on a third-party tool’s behavior, read that tool’s own config or docs for the specific behavior — don’t infer it from what the tool broadly does.

A regression test written alongside a fix can lock the bug in rather than catch it — assert the two paths that diverge, not the one you just touched.

A systematic audit done by skimming is worse than the one-at-a-time version it replaces.

Adding an explanation supersedes whatever the file already said about the same thing, so re-read the older passage — your own diff is the likeliest source of a contradiction nobody flags.

The same rule applies within a single diff, and there nothing prompts the check at all.

And when the explanation you add is a *mechanism* claim, test the class it distinguishes, not just the sample in front of you.

See [ardi.cases.md](#), “A mechanism claim whose population held no true positive”.

A symptom that stops reproducing is a fix having landed, until you have checked otherwise — reaching for nondeterminism is the attractive wrong answer.

- **Do:** look for a merged fix, and date it, before attributing a vanished symptom to anything.
- **Do:** report the before/after with its timestamps, so the negative control is visible rather than asserted.
- **Don’t:** explain a symptom’s disappearance as nondeterminism on the strength of one clean run.
- **Don’t:** carry such a claim into an issue or a decision doc, where it argues against the very fix that produced the silence.

The mirror runs the other way, and it is the one that discards a good fix: a symptom that KEEPS reproducing after a fix landed is not evidence the fix failed.

The bullet above governs a symptom that vanished, where the attractive wrong answer is nondeterminism. Here the symptom is still there, and the attractive wrong answer is that the diagnosis was wrong — which sends you back to re-litigate a fix that is working, and leaves the real remaining cause unread.

The mechanism is ordinary and worth naming, because it makes the persistence expected rather than surprising. A failure can have causes in series, and only the first one is observable while it stands. Removing it does not change the outcome; it changes which cause produces the outcome. So the job’s colour is the same before and after, and the outcome is the one thing everybody checks.

The discriminator is the error, not the outcome. Both runs failed, so comparing conclusions establishes nothing. Comparing the error text is decidable in one read, and a changed error means the first cause is gone and a second was behind it. Where the fix is upstream, pin the comparison to the dependency version each run actually resolved, since a run predating the fix is not evidence about it — [dont-reinvent-wheel](#)’s “mirror direction” section owns that lookup.

Note the asymmetry that makes this worth a rule. Reading the new error costs one glance and usually names its own remedy. Re-litigating the first fix costs a round, and it argues for reverting something correct — the same shape the bullet above warns about, where a claim ends up arguing against the fix that produced the change.

- **Do:** diff the error text across the fix, not the pass/fail outcome, before concluding anything about whether the fix worked.

- **Do:** resolve which dependency version each run used, when the fix landed upstream, so a pre-fix run is not read as evidence against it.
- **Do:** report a changed error as a second cause found, and file it, rather than as the first fix having failed.
- **Don't:** re-open a landed fix because the symptom persists — that is a claim about the outcome, and the outcome is what a serial second cause preserves.
- **Don't:** read the earlier bullet as covering this; it fires on a symptom that stopped, and this one fires on a symptom that did not.

See [ardi.cases.md](#), “A trust-gate fix that revealed a tool-name mismatch behind it”.

A third cell: a symptom that stops reproducing AFTER you applied a remedy is not evidence the remedy worked.

The two bullets above cover the case where you changed nothing and the symptom vanished, and the case where you changed something and the symptom stayed. This is the remaining one, and it is the only one of the three that ends in a green check — which is why it is the one that gets reported.

A pass following your change is consistent with the change having fixed it. It is equally consistent with an intermittent failure that did not fire this time, with something else having moved in the base, and with the failing condition never being reached on this run. Applying a plausible remedy and observing success is not a controlled test, because no run of the unchanged head exists to compare against. [review-verdict-pitfalls](#) states the complementary half, and its wording is the discriminator: a retry is a genuine negative control “because nothing changed between the two runs”. Change something and you have spent that control.

Three things make this harder to catch than either neighbour.

The remedy is usually correct to apply. Porting a fix that already exists on `main` is what [sync-with-main](#) asks for, so nothing about the action is a mistake. The causal claim then inherits the action’s correctness and is never examined on its own, which is the near-miss: right action, unchecked story, and no moment at which the two come apart.

Green terminates inquiry. A red check invites investigation and a green one closes the question, so the wrong story survives in exactly the place least likely to be revisited.

The refuting datum is usually free. A timeout failure carries its own duration, which is in the job metadata whether or not anyone reads it.

The test is one question, asked before the claim is written: **name the observation that would differ if your change were the cause, and go and look for it.** For a raised timeout that is a duration falling between the old limit and the new one, so a passing run that finishes inside the OLD limit never reached the raise. For a widened pattern it is an input matching the new alternative and not the old. For an added dependency it is the code path that requires it. Where the passing run does not show that observation, you have a green check

and no evidence, and the honest report says the check passed rather than that your change fixed it.

- **Do:** name what the passing run would have to show if your change were the cause, and read that field before reporting causation.
- **Do:** report the pass as “the check passed” until that observation is in hand, and say separately that the remedy is worth keeping on its own merits.
- **Do:** compare a timeout’s passing duration against the OLD limit — a run finishing inside it proves the raise was never exercised.
- **Don’t:** read a correctly-applied remedy as licensing the claim that it worked; the action being right is what carries the claim past review unexamined.
- **Don’t:** treat a single pass as a controlled test — changing something is precisely what removes the negative control a bare retry would have given.

See [ardi.cases.md](#), “A green check credited to a timeout raise that never ran”.

Verify a command, path, or flag *you* write into a doc, with the same rigor [address-every-comment](#) demands for one a reviewer suggests.

- **Do:** confirm every literal you invent against the tool’s own source or help output before it lands in a doc.
- **Do:** cite the file or command you checked, so the claim stays falsifiable.
- **Don’t:** infer a subcommand from a family that has its siblings (`gh label list/create/edit` does not imply `gh label view`).
- **Don’t:** treat a literal as exempt because the prose around it is well-sourced — the literal is the part a reader executes.

Run that check over your own fix, too — the remedy for an unverified literal is where the next unverified literal goes.

- **Do:** re-run the rule you are applying against the text of your own fix, before committing it.
- **Do:** say in the thread when a fix’s own draft tripped the same rule, since that is the only place the near-miss is visible.
- **Don’t:** treat the effort of writing a correction as evidence the correction is verified.

The same rule reaches past a literal, to the defect CLASS a code fix just closed.

A consolidation commit is the highest-risk host for it, and the likeliest to be trusted.

- **Do:** ask whether the fix’s own new code instantiates the class it closed, before committing it.
- **Do:** treat a commit that consolidates one duplicated concept as owing a check that it forked none, since its framing argues the other way.

- **Do:** compose an existing shared anchor or helper into a new site rather than hand-rolling an equivalent, so the site inherits later fixes too.
- **Don't:** read a diff that removes a duplicate as evidence that it added none.
- **Don't:** treat the next round's finding at a new address as a fresh gap without first checking whether your own previous fix created that address.

When regenerating a generated tree makes it most of the diff, say so in the PR body — otherwise a reviewer reads it as pollution and blocks.

- **Do:** grep a file for a generated-by header before editing it, and change the source instead.
- **Do:** state in the PR body how many of the changed files are generated, and name the hand-written ones.
- **Don't:** revert generated output because a reviewer calls it noise — check first whether the sync check requires it.
- **Don't:** assume a reviewer sees the source files; on a large diff they frequently do not.

See [ardi.cases.md](#), “Editing generated output, then being read as pollution once regenerated”.

Run the whole test suite before pushing, not the files you predict the change touches — and check that the ones you ran were not silently skipped.

Matching the tool's VERSION is not matching its ENVIRONMENT, and when the tool GENERATES a file you are about to commit, the gap ships.

Read the generator's own diagnostics first, because a good one says so outright and names the cause.

The file list is the backstop, for a generator that degrades with no diagnostic, or one whose diagnostics scroll past in a long run.

- **Do:** read the generator's own warnings before its output — roxygen2 names the missing package and the tag that needed it.
- **Do:** compare your generator's changed-file list against the CI log's, and treat any extra file as an environment mismatch until explained.
- **Do:** install the optional/dev dependency set as well as the tool, when a generator loads the package to do its work.
- **Don't:** read “I installed the same version CI installs” as having matched CI — version is one input to the output, and rarely the one that differs.
- **Don't:** commit generated output whose diff is wider than the job you are trying to satisfy reported.

See [ardi.cases.md](#), “A generator's environment, not its version, changed the committed artifact”.

- **Do:** run the full suite before pushing, and state the tests/failed/ skipped triple rather than “tests pass”.
- **Do:** set the flags that un-gate conditional skips, and re-run if the skip count is non-trivial.
- **Don’t:** scope a local run to the files you edited — the test asserting the old behaviour is usually somewhere else.
- **Don’t:** read a green subset as a green suite, or a skip as a pass.

Running a script is not running its tests, and an “advisory” check can have a hard-gating twin.

- **Do:** run every check the CI job runs, its test files included, before pushing.
- **Do:** grep the job definition for other steps touching the same property before saying anything about whether it gates.
- **Don’t:** substitute a production script’s exit code for its test file.
- **Don’t:** infer a job’s behaviour from one step’s label — “(advisory)” describes that step, not the job.

A third failure mode of the whole-suite rule above: the suite holds no case that could have failed.

- **Do:** construct the input class the change is supposed to handle and diff its behaviour against the pre-change code, before calling a guard verified.
- **Do:** name which cases could have exercised the defect class, rather than quoting a suite total — the total is a fact about the suite, not the diff.
- **Don’t:** offer a pre-existing suite’s green as verification of a change it holds no case for; those cases predate the defect and cannot speak to it.
- **Don’t:** read the tests/failed/skipped triple above as covering this — it makes the report more precise without making it any more relevant.

A fourth failure mode: the case exists, and which branch it reaches is decided by the host.

- **Do:** name in the test which host-derived value selects which branch, and add a case that pins each branch regardless of that value.
- **Do:** run a host-dependent suite in both environments before believing its coverage, and say which branch each run took.
- **Don’t:** read green in CI as covering a branch whose selection depends on an input CI happens to supply one way.
- **Don’t:** reach for the skip count here — nothing is skipped, so that component is identical on both machines even when the failed counts diverge.

See [ardi.cases.md](#), “A suite whose branch coverage varies by host”.

5 What “Fully Clean” Means

“Fully clean” is the terminal state the ARDI review loop drives toward. A GitHub PR is **fully clean** when **both** of these hold, verified via `python3 scripts/check-pr-fully-clean.py --quorum <number-of-reachable-providers> <pr-number>`. For a GitLab MR, establish the same criteria from GitLab’s current-head pipeline, review, and discussion APIs; `check-pr-fully-clean.py` queries GitHub and cannot verify a GitLab MR.

Active monitoring and polling are required after every push until fully clean.

Reaching fully clean requires an active polling loop or scheduled wake mechanism after each push. Do not pause passively or assume check runs and reviews will complete without polling. Actively query the current head’s CI/pipeline runs and review verdicts (`gh` for GitHub, `glab` for GitLab) until each round reaches a terminal state, re-arming the poll while work remains.

- **Do:** actively poll and re-arm monitoring after every push until CI and review reach a terminal state at the current head.
- **Don’t:** stop polling while CI or reviews are in flight, or assume automated pipelines completed without querying the forge.

A forge’s mergeable result is an integration-state signal, not a review verdict. It can be true while a reviewer has left resolvable findings open. Do not report a PR/MR fully clean, ready to merge, or merge it until the review thread sweep is also clear.

- **Do:** for a GitLab MR, page through `projects/<project>/merge_requests/<iid>/notes` and confirm that every resolvable, actionable `DiffNote` is resolved before reporting it fully clean. The notes endpoint is authoritative because it can expose unresolved diff notes absent from a discussion-level sweep; use `discussions` to resolve the thread after finding it.
- **Do:** for a GitLab MR, obtain the current head SHA and page through every pipeline on that SHA; confirm each has completed successfully or was skipped.
- **Do:** accept a GitLab review verdict as current-head evidence only when its body names that SHA or its diff discussion has `position.head_sha` equal to it; then read every review note’s full body for a clean verdict.
- **Do:** re-read the GitLab MR head SHA immediately before reporting fully clean; when it changed during the sweep, restart the pipeline, review, and discussion checks against the new head.
- **Do:** state an MR is *mergeable but not fully clean* when CI and merge status are green but an actionable review thread remains open.
- **Don’t:** treat GitLab’s `detailed_merge_status: mergeable` as a substitute for a current clean review verdict and resolved discussion threads.
- **Don’t:** treat an unresolved clean final verdict as an actionable finding; GitLab can mark that all-clear note resolvable too.
- **Don’t:** treat a resolved GitLab discussion sweep as a clean current-head review; findings can appear in an overall review note without a resolvable discussion.

- **Don't:** let an earlier review's green pipeline or later code push erase an unresolved finding without a reviewer-confirmed clean round.

In a remote/web session the instrument still runs, and hand-checking the axes in its place is not acceptable (user directive, 2026-08-29, ai-config#2441). Those sessions have no `gh` CLI, and an MCP tool cannot be called from inside a Python subprocess, so the split is that the **agent** retrieves and the **script** judges: gather the PR's state via MCP, write it to a file, and pass `--from-json <file>`.

Payload key	Gather with	Notes
<code>repo</code>	the OWNER/REPO under check	Or pass <code>-R</code> instead.
<code>pr</code>	<code>pull_request_read</code> (get, <code>get_reviews</code> , <code>get_comments</code>)	See the field list below.
<code>check_runs</code>	<code>pull_request_read</code> (<code>get_check_runs</code>)	Bare list or the REST <code>{"check_runs": [...]}</code> envelope.
<code>actions_runs</code>	<code>actions_get</code> (<code>get_workflow_run</code>), keyed by run id	Omitting it changes verdicts — see below.

`pr` needs `headRefOid`, `headRefName`, `state`, `reviewDecision`, and `commits[].committedDate`, plus two nested shapes the scan reads directly: each entry of `reviews[]` needs `state`, `submittedAt`, `body`, `commit.oid`, `author.login`, and `authorAssociation`, and each entry of `comments[]` needs `body`, `author.login`, `createdAt`, and `authorAssociation`. `commit.oid` is the exact-SHA gate, and `submittedAt/createdAt` order the latest-verdict selection, so a payload omitting them is accepted and scored on weaker evidence.

The field names are `gh pr view --json`'s rather than the MCP tool's, so a small mapping is needed — `head.sha` becomes `headRefOid`, and an author becomes `{"author": {"login": ...}}`. That mapping is the agent's job precisely because it is the part that differs between session kinds.

Page through `pull_request_read`. It caps at 100 per page, and a truncated `comments[]` is indistinguishable from a complete one — so dropping a later standing not-clean yields a **false clean**, which is the one error this instrument exists to prevent.

`actions_runs` is optional but not inert. Omitting it disables the ai-config#2277 suppression of a cancelled run superseded by a later success, and disables run-based attribution of a verdict whose comment cites its run URL rather than a SHA. Both omissions push toward a false **not-clean**, which is the safe direction — but if a PR reads not-clean for a reason you cannot see in its checks or reviews, gather this key before concluding anything. Never add entries until the verdict flips: that is tuning the instrument to agree with you.

Unusable payload data exits 2, never 1. Exit 1 is the script’s real not-clean verdict, so a data problem reported as exit 1 would be indistinguishable from a finding about the PR. That covers a missing or wrong-typed top-level key, a `pr` missing `headRef0id`, malformed JSON, and — while `--from-json` is active — any other exception, since none of them can be a statement about the PR.

Note what the rationale is *not*: an absent `check_runs` read as `[]` does not score clean. `check_ci_runs` already reports “No check runs found” and returns not-clean. Refusing absent data is still right, because substituting an empty value would manufacture a *finding bullet* out of missing data, which is worse than a crash rather than better.

Extended rationale — the mechanism, evidence, and argument behind each rule below — lives in [fully-clean.rationale.md](#), moved out of the auto-loaded context. Each rule here keeps its statement and its Do/Don’t pair; read the companion when the reasoning or the evidence is the question.

Worked-example case records for the rules below live in [fully-clean.cases.md](#), moved out of the auto-loaded context.

1. **All CI workflows and check runs are green AND completed.** Every workflow and check run passes — not just the required checks and not just the review job.

status itself can be stale, so never infer a job’s *duration* from it.

- **Do:** read elapsed time from log timestamps whenever the length of a run is the thing being judged.
- **Don’t:** conclude a job is still running, or has passed some duration threshold, from `in_progress` plus the wall clock.

When you are waiting for a job rather than timing one, poll its step list instead of its status — the steps are not subject to the same lag.

- **Do:** poll `actions/jobs/<id>’s steps[]` when waiting on a specific job, and treat its terminal step completing as the signal.
- **Do:** report which step the job is on, so a stalled job is distinguishable from a slow one.
- **Don’t:** poll a check run’s `status` in a loop and read repeated `in_progress` as evidence the job is still working.

See [fully-clean.cases.md](#), “Poll a job’s step list, not its check-run status”.

A `BlobNotFound` / HTTP 404 on the job-log fetch means the job has not completed, not that it has hung.

- **Do:** read a 404 / `BlobNotFound` on the job-log endpoint as “the job has not finished”, and wait for completion (or read the live UI log) before judging its outcome.

- **Do:** take a job’s real state from its `status/conclusion`, since the same 404 covers a still-running job and a completed-with-no-logs one.
- **Don’t:** read a 404 on the log fetch as positive evidence of a hang or a stall — it is the opposite, evidence the job is still running.
- **Don’t:** file an issue reporting a review job as hung or “no verdict produced” while its log fetch still 404s and its status is `in_progress`.
- **Don’t:** run the rule backwards: a log URL being **served** is not evidence the job completed. The blob can exist mid-run, so a successful log fetch and a still-running job coexist. Completion comes from `status/conclusion` alone, in both directions.

A hung check is a distinct state from a crashed one, and the step list is what tells them apart. `status/conclusion` look identical for both — empty and `in_progress` either way — so the remedies above for a *finished* check that misled you do not apply here: there is no failure yet to `rerun_failed_jobs`, and a fast-crash signature has nothing to match against a job that has not crashed. Read the step list instead: which step is `in_progress`, and for how long relative to that step’s usual duration. A retry step stuck for tens of minutes where its sibling attempt finished in under a minute is a hang, not a slow-but-normal run.

- **Do:** read the job’s step list when a check runs unusually long, and name which step is hung and for how long.
- **Do:** state a hung check as a tooling status when another reviewer has already given a verdict at the same head, rather than either declaring clean or chasing it as a finding.
- **Don’t:** try to re-run an `in_progress` job, or apply a fast-crash remedy to a hang — neither condition has fired yet.
- **Don’t:** infer from `status: in_progress` alone that the check is merely slow; read the step list before concluding either way.

See [fully-clean.cases.md](#), “A hung retry step, not a crashed one”.

gh pr checks is not a complete enumeration of a head’s check runs, so read the commit check-runs endpoint before deciding that everything has finished. GraphQL `statusCheckRollup` is the same kind of short surface for a *progress* report — not enough for a terminal “fully clean” / “ready to merge” claim. That claim needs `scripts/check-pr-fully-clean.py` (ai-config#2277, 2026-08-26).

--paginate is load-bearing, not tidiness.

The endpoint covers check runs only, so a repo that still uses legacy commit statuses needs a second query.

Why the two surfaces disagree is unexplained, so do not assert a mechanism for it.

- **Do:** take the check-run half of criterion 1 from the paginated check-runs endpoint, and add `commits/<sha>/status` where the repo uses commit statuses, rather than treating either query as sufficient alone.
- **Do:** report both counts when the endpoint and the rollup disagree, so the gap stays visible to whoever reads the status next.
- **Do:** re-derive check state from that endpoint on the PR's current head when a completion notification wakes you, having first compared the wake's own `head_sha` against that head.
- **Don't:** read 0 `pending` from `gh pr checks` as evidence that nothing is still running.
- **Don't:** drop `--paginate` — an unfinished run on page 2 returns the same empty result as a finished head.
- **Don't:** offer a reason for the omission — none was established.
- **Don't:** read a wake reporting check suites finished as an all-clear — `ardi`'s superseded-head case is a **red** wake inviting a needless fix, and this is its **green**-sounding mirror, inviting a needless merge.

A paginated sweep with an inconsistent page size silently skips items, and every response still reads as complete coverage. `--paginate` above is the CLI answer; the same failure reaches a manual REST/GraphQL sweep (an MCP tool taking explicit `page/perPage` arguments, say) when the page size changes partway through — `page: N` is relative to whatever size was last requested, so changing it mid-sweep silently renumbers what “page 2” means, and a page fetched at the wrong offset that way returns real, green data with nothing in the response flagging the gap. Both API surfaces hand you the check for free (`total_count` on REST, `pageInfo.hasNextPage/endCursor` on GraphQL), so this is a count to assert, not a judgment call.

- **Do:** hold the page size constant across a paginated sweep, and assert items seen equals `total_count` before drawing a conclusion from it.
- **Do:** prefer cursor pagination (`after/endCursor`) where offered, since a cursor cannot be invalidated by a page-size change.
- **Don't:** change the page size between pages of the same sweep.
- **Don't:** treat “every item I looked at was green” as “every item is green” without the count check — the same mistake made while sweeping review threads (below) is worse, since a skipped range there means an unresolved thread reads as a clean PR.

See [fully-clean.cases.md](#), “A paginated sweep with inconsistent page size skipped 8 of 30 items and read as complete”.

See [fully-clean.cases.md](#), “A `check_suite.completed` wake at a superseded head”.

A polling loop needs the same negative control a sweep does, because an EMPTY check list satisfies “nothing is pending” exactly as well as a finished one. The rules above that concern the check list each found one that came back **short**; this is the case where it comes back **empty**, which none of them reaches. “Not yet

started” and “finished successfully” produce an identical reading, which is `fail-fast`’s pass-path-equals-failure-path shape failing in the dangerous direction — it reports a PR ready. `batch-merge-and-resolve` states the governing rule for a **sweep**, and a sweep and a poll do not resemble each other from the inside — one feels like a measurement, the other like waiting — so that rule loads and matches nothing here.

A non-empty population is necessary and **not sufficient**, because the population grows while the poll runs. Measured on the corrected run: the total went 13, then 16, 17, 18 across two minutes, as later workflows registered their checks. So a threshold only rules out the empty case, and a poller that happened to observe zero pending at total 13 would have exited before five further checks existed. Require the terminal reading to repeat — zero pending **and** an unchanged total across two consecutive polls — and print the total each tick, so growth is visible rather than inferred.

- **Do:** require a non-empty population before reading zero-pending as done, and report how many check runs were examined.
- **Do:** confirm the total is unchanged since the previous poll, since the population grows as workflows register.
- **Don’t:** arm a poller assuming the push already created the checks — a draft-to-ready transition creates them on a separate event, so a poller armed at push time can run entirely inside a zero-check window.
- **Don’t:** treat a single zero-pending reading as terminal, however large the population was when you took it.

See [fully-clean.cases.md](#), “A poller exited on an empty check list”.

The population also grows LATE, from a job PASSING rather than from registration lag, so a total that has been stable for many polls is not evidence the set is final. The paragraph above measured growth in the two minutes after a push, as workflows registered — 13 to 18 across $t=150s$ to $t=270s$ — and that shape invites reading the growth window as bounded by the push. It is not. A job gated on **needs:** creates its checks only when its dependency **completes**, so a green job is a cause of new checks rather than one fewer thing to wait for.

That inverts what a stable total means. Under registration lag a stable total is at least weak evidence the set has settled, because registration is contiguous with the push and then stops. Under a gated successor the total is stable *because* the spawning job has not finished, so the stretch that looks most settled is the one immediately preceding the growth.

The two-poll rule above held in the case below, and for a reason that does not generalize: the spawning job sat inside the polled set and was itself pending, so no poll could have read zero. A spawner outside that set — a job in another workflow, a `workflow_run` trigger — leaves the guard nothing to see. So the reading that licenses a merge is one

taken **at the merge decision**, not a count carried forward, however many polls agreed on it.

- **Do:** re-read the check list at the moment of the merge decision, and compare its total against the one the earlier polls agreed on.
- **Do:** read a job completing as a reason to expect new checks.
- **Don't:** read a long-stable total as evidence the set is final — the stable stretch is where a gated successor is still waiting to appear.
- **Don't:** generalize the two-poll rule's success here; it depended on the spawner sitting inside the polled set.

See [fully-clean.cases.md](#), “A passing job spawned three more checks”.

A bot comment on the PR is not evidence that the job posting it has finished, because a step writes that comment partway through the job. A preview-deployment comment, a coverage report, a benchmark table: each is emitted mid-job, so its existence establishes that one step ran and says nothing about the steps after it. One job can post two such comments, minutes apart, and finish after both.

It is the comment rather than the check row that gets read this way for a structural reason: the comment is the most visible thing the job produces and it lands in the thread already open, while the check list has to be fetched. [efficient-pr-babysitting](#) compounds that by telling you to work from CI's own report rather than re-deriving it, which is right about the comment's **content** and silent about its **timing** — so trusting the content makes the timing inference feel already licensed. This is [verify-the-right-artifact](#)'s neighbour-for-the-target shape: a real artifact of the right job, read for a property it does not carry.

- **Do:** take job completion from the check list, and read a bot comment only for what it measured.
- **Don't:** date a job's completion from a comment that job posted.

See [fully-clean.cases.md](#), “A preview comment read as a finished docs job”.

A check-run NAME is not unique across workflows, so a name alone does not identify which check passed. Two workflows in one repo can each define a job with the same name, and `gh pr checks` prints the bare name with no workflow attached — so a passing row can belong to a workflow you were not asking about. The ambiguity is invisible in the output, which is what makes it dangerous: nothing in a duplicated name looks different from a unique one, so no prompt to check ever arrives.

Resolve it from the run behind the check rather than from the name:

```
gh api "repos/<owner>/<repo>/commits/<sha>/check-runs" --paginate \
  --jq '.check_runs[] | select(.name == "<name>") | .html_url'
gh run view <run-id> -R <owner>/<repo> --json workflowName --jq .workflowName
```

Cross-check against the workflow’s own job list too. A matrix leg gated on `needs:` may not have started at all, so its absence from a run’s jobs contradicts any same-named row reported as passing.

`check-pr-fully-clean.py` annotates a duplicated name with the run URL only on the lines it actually reports — a run still pending, or one that finished badly. A **passing** duplicated name produces no line at all, so it is never annotated, and the manual lookup above is the only thing that resolves it. That is precisely the case this section was written from: the passing row belonged to the wrong workflow, and nothing in the script’s output would have said so.

- **Do:** take the workflow from the check run’s own URL before attributing a pass or a failure.
- **Don’t:** read a job name as identifying a workflow — it identifies a job, and two workflows may define the same one.

(Measured 2026-08-21 on `ucdavis/bcs: ubuntu-latest (release)` exists in both `R-CMD-check.yaml` and `check-readme`. On a PR fixing an `R CMD check` failure, the passing row was `check-readme`, while `R-CMD-check.yaml`’s matrix legs had not started — they are gated on `needs: [matrix, update-snapshots]`. Reporting the regression fixed on that row would have cited an unrelated workflow.)

Every subsection above explains a per-PR failure in reading the check state — a short or empty list, a lagging status surface, an ambiguous name — and a platform outage produces the same shape for a reason none of them can reach.

A job’s conclusion is set by whichever step failed, which need not be the step whose verdict you read. Most rules above concern an enumeration that came back incomplete. This one’s enumeration is complete and terminal, and the answer you read came from the wrong member of it. A workflow can carry a guard step that decides what a run *meant* — a review guard classifying an outcome, a summarizer, a status resolver — and that step can conclude “this is fine”, write its output, and end **success**, while the job is red because an earlier step failed without **continue-on-error**. Reading the guard’s own log line then reports the opposite of the check. So when a red job’s log carries a green verdict, do not treat it as a contradiction to explain: enumerate the steps and find the one whose conclusion is **failure**. The same reading also settles what to do next: whether a fix to the classifier can clear the check at all, since a classifier the job does not consult is fixable without changing anything the reader sees.

- **Do:** identify the failing *step* before diagnosing a failing job, rather than reasoning from whichever step’s output you happened to read.
- **Do:** treat a green guard step beside a red job as evidence about the wiring, since the two were decided by different steps.

- **Don't:** read a guard step's own log line as the job's verdict — the two are decided by different steps, so agreeing is a coincidence rather than a confirmation.
- **Don't:** claim a fix to a classifier clears a check until you have confirmed the job's conclusion actually depends on that classifier.

See [fully-clean.cases.md](#), “A green guard step beside a red job”.

One SHA can carry two check runs of the same name, from the same workflow, with opposite conclusions — because a workflow gated on a base-ref diff runs VACUOUSLY on push and meaningfully on pull_request. The subsection above covers a green step inside a red job. This is the mirror at the run level, and it is worse, because nothing about the green one looks partial: it reports the same check name, it completed, and it passed.

The mechanism is a workflow that needs a base to diff against. `check-new-line-breaks.yml` passes `base-ref` only when `github.event_name == 'pull_request'`, so the push-triggered run of the identical workflow has no base, examines zero added lines, and passes having measured nothing. Both runs attach to the same commit, so `gh pr checks` prints two rows with one name, one `pass` and one `fail`, and reading the list top-down finds whichever came first.

The vacuous run is the one to discard, and the trigger event is the only field that separates them. `gh api "repos/<owner>/<repo>/actions/runs/<id>" --jq '.event'` settles it in one read per run. A `pass` from a run whose event supplies no base is [batch-merge-and-resolve](#)'s zero-matrix problem arriving as a green check. That fragment states it as “a matrix of zeros is indistinguishable from a detector that never ran”, and prescribes a negative control before trusting any zero. The same remedy applies here, and the trigger event is what supplies it: a run given no base examined nothing, so its `pass` is the zero rather than a result.

Note that this is not the same as [ai-config#1870](#)'s ambiguity, where two *different* workflows contribute check runs sharing a name. Here it is one workflow, and the disambiguator is the event rather than the workflow name — so a fix keyed on `workflowName` cannot see it.

- **Do:** read the `event` of any run whose verdict you are about to rely on, whenever the same check name appears twice on one head.
- **Do:** take the verdict from the `pull_request`-triggered run for any check that diffs against a base.
- **Don't:** read a `pass` as evidence the check examined anything — ask what population it was given first.
- **Don't:** resolve a same-name disagreement by workflow name. On this shape both runs carry the same one.

(Measured 2026-08-22 on [ai-config#1884](#). Run 32545283504 (event=push) and run 32545289903 (event=pull_request) both had head_sha=8c456074, both were named new-line-breaks / check-new-line-breaks, and they concluded success and failure respectively. The push run was read first and taken as the verdict. The PR run was the one carrying four real findings.)

2. **Every reviewer’s latest verdict is totally clean:** no nits, and every item that wasn’t directly **Addressed** is either **Deferred** to a tracked follow-up issue, or **Rebutted with a rebuttal that actually convinced the reviewer** — i.e. the reviewer did *not* re-raise it on the next round. A later all-clear from a different reviewer does not clear another reviewer’s standing not-clean, nits included.

In a local CLI session, “external reviewer” means the bot reviewers, not the human one. gh/the MCP server authenticate as the same human account that owns the repo, so a formal review *request* against that human always 422s (request-pr-review’s own edge case) — not occasionally, structurally, on every PR such a session opens. That does not relax criterion 2; it changes which providers can satisfy it. The pinned quorum below is still required in full, and it is answerable entirely by the bot reviewers (claude-review, jules/review, Copilot) — none of which is the human the deadlock-escalation ladder in CLAUDE.md’s “Address every in-scope review comment, even non-blockers” section (and [skills/ardi/SKILL.md](#)) points to. So read that ladder’s “request the human reviewer” step, in a local session, as “post an @-mention with the impasse” instead: the mention still notifies and still counts as re-checking reviewer reachability right before declaring clean, which criterion 2 requires regardless of session type.

- **Do:** treat the bot quorum as the external-reviewer requirement in a local session, and re-check its reachability before calling a PR clean.
- **Do:** escalate a deadlock with an @-mention comment in a local session, in place of a review request that would 422.
- **Don’t:** read the 422 as a transient failure worth retrying, or as evidence the PR cannot reach criterion 2.
- **Don’t:** assume a remote/web session is exempt — per [memories/github-mcp-tools.md](#), that session’s authenticated identity varies by container and client, and can itself be the repo owner, in which case the identical 422 fires there too; settle it from an actual write’s attributed author rather than from the session type alone.

Criterion 2’s test is the absence of findings, not the presence of a verdict line saying so.

So when the two disagree inside one comment, **the findings win**. Read to the end of the comment before calling anything clean, and count the items under every heading, whatever that heading is called — [address-every-comment](#) already establishes that “informational”, “non-blocking”, “nit”, “minor”, and “optional” are prioritization labels rather than a pass,

and a reviewer files findings under exactly those words in the section that contradicts its own verdict line.

Final approval requires a clean verdict from every available provider in the pinned quorum. A single agent’s clean verdict does not clear criterion 2 on its own; it must be joined by a clean verdict from the rest of the quorum.

This is a directive rather than a derivation, so treat it as a standing preference. What it settles is when a PR is reported **ready** for merge.

The reason it needs stating is that multiple verdicts are indistinguishable from a single one to the CI review gate. Every agent posts the same shape — a summary, some analysis, a positive closing line. So a findings-free report from just one agent turns the review-gate check green, even while other providers’ reviews are pending or blocked.

Two failure modes make the preference concrete, and both have recurred:

- **A clean verdict over tooling that errored.** A report can open by saying its own grep failed and then approve on the strength of the analysis that grep was supposed to support. The error line sits above the verdict, so it reads as a caveat rather than as the verdict’s foundation collapsing.
- **A clean verdict at a head another agent finds a real defect in.** Not a difference of opinion about a nit — a checkable factual error, at the same commit, that the clean verdict passed over.

So when the pinned quorum is reachable, their verdicts are the ones to report on:

- **Do:** dispatch reviews to the quorum and wait for all verdicts before reporting a PR ready, whatever one agent has already said.
- **Do:** name which agents produced the verdicts you are reporting, so “clean” is attributable rather than anonymous.
- **Do:** treat any agent’s findings as real findings — every provider’s objections count.
- **Don’t:** report a PR ready on a single clean verdict while the quorum is still reachable and pending, however thorough that report reads.
- **Don’t:** read a green review-gate check as settling this; the gate does not know how many agents answered, only that one did.

A disagreement among reviews vetoes merge, including under mwc. Criterion 2 is every reviewer’s latest verdict, not the globally last comment. If one review is all-clear and another raises blocking issues, nits, minor items, or any other flagged heading, the findings win. ARD every item from every review, then request fresh reviews. A later all-clear from a different reviewer does not supersede a standing not-clean; only a later clean from the same reviewer does (the ordinary ARDI iterate path). `check-pr-fully-clean.py` encodes the per-reviewer scan (ai-config#2274).

- **Do:** ARD the union of findings from every review, then request a fresh round from the reviewers that spoke.
- **Don't:** merge on one reviewer's all-clear while another still has a standing not-clean, even with `mwc` active.

This is a different question from how much two reviewers **agreeing** is worth, which `self-review-fallback`'s cross-vendor section settles: there, same-vendor agreement measures a shared blind spot, and a cross-vendor split is a prompt to check the item yourself. That section weighs corroboration; this one names whose approval is terminal. They compose — a cross-vendor reviewer is always required, and every clean verdict is necessary before a PR is reported ready while the quorum is reachable.

Where a quorum provider is genuinely unreachable — quota-skipped, a stub with no stated verdict, or not configured — fall back per `self-review-fallback`, which already governs that case. Note that merging autonomously under `mwc` (merge-when-confident) strictly requires genuine clean automated review verdicts from the reachable quorum evaluating the HEAD commit; a fallback self-review or reviewer skip notice allows the ARDI iteration loop to proceed, but NEVER satisfies the MWC autonomous merge gate.

One more gate stacks on top (user directive, 2026-08-25): no merge under any grant, `mwc` included, without a 100% all-clear adversarial verdict at the shipping head from a reviewer meeting `adversarial-self-review`'s independence bar. It composes with the external-reviewer requirement above — neither satisfies the other. When the external reviewer self-skips by design (workflow modification is the known case), autonomous merging stays blocked: human approval is the only path. Specification and mechanics live in that fragment.

See `fully-clean.cases.md`, “Two agents, one head, opposite verdicts”.

Both criteria are per-PR, and a stack is where that stops being automatic.

- **Do:** derive a verdict per PR number, and name the PR beside each one.
- **Do:** treat a refusal from one reviewer on one PR as evidence about that reviewer on that PR, and nothing else.
- **Don't:** report a stack's review state from a single read — “I read the review” is a per-PR claim, and the stack is what makes it read as a claim about the work.

See `fully-clean.cases.md`, “Both criteria are per-PR, and a stack is where that stops being automatic”.

The disagreement is measurable, and it is not a wording problem.

A reviewer's own verification block can be wrong while its verdict is right.

- **Do:** re-derive a posted verification's groups, not just its total.
- **Do:** fix the wording that invited a wrong reconstruction, even when nothing in the diff was false.

- **Don’t:** let the word “verification” stand in for having verified.
- **Don’t:** read a table that sums as one that partitions correctly.

A clean verdict can ratify an enumeration instead of testing it, and then it reads as independent corroboration of a false scope claim.

- **Do:** derive any enumeration you publish with a command, and publish the command beside it.
- **Do:** treat a reviewer restating your count as that count still being unverified.
- **Don’t:** read a clean verdict as evidence that a scope claim in the diff is complete — a reviewer can only check the members you named.
- **Don’t:** count a reviewer’s agreement as independent when its population came from your own prose.

What “an approving review” means here is not a review state.

- **Do:** read the whole review comment and count findings under every heading before calling a PR clean.
- **Do:** establish approval from the findings and thread lists, since `.state` is COMMENTED on every review this repo receives.
- **Don’t:** quote a **Ready for merge** line as the clean signal while the same comment lists findings.
- **Don’t:** wait for a formal APPROVED review, or read COMMENTED as a defect in the reviewer.

Findings hide on several surfaces, and no single check sees all of them — so read the verdict body, any suppressed-comments block, the inline comments, the thread list, and the verdict’s own conclusion every round.

- **An out-of-diff finding never becomes a thread.** A finding about a line the diff did not touch cannot be attached as an inline comment, so it appears only in the body — reviewers say so explicitly (“inline comments were unavailable for out-of-diff lines”). A thread count therefore cannot see it. Zero unresolved threads is not evidence of zero findings.
- **A notification that truncates the body hides exactly that finding.** The rule above says to read the body, and assumes you are reading the body. A CI-monitor or webhook event delivers the review as *quoted text*, capped at some length, and the inline findings are enumerated first because they are numbered — so what gets cut is the tail, which is where an out-of-diff finding and the verdict both live. The event is honest about it, and that is the trap: it prints a marker like `[truncated --- full text: gh api repos/<owner>/<repo>/issues/comments/<id>]`, which reads as a courtesy rather than as an instruction, and the visible portion looks like a complete, well-structured review. Acting on the inline comments alone then feels like having addressed the round, and the thread sweep confirms it, because the missed finding was never a thread. So run

that command before treating a finding list as complete, whenever the review reached you through a notification rather than through a direct read.

- **An empty body hides the mirror case.** A review can post a completely empty top-level body and carry its entire finding in one inline comment, so a body-only read finds nothing to act on and concludes there is nothing.
- **A clean overview can hide a collapsed findings block.** Copilot can say it “generated no new comments” and create zero inline comments while placing substantive findings inside a collapsed `<details>` suppression block in the review body. Match case-insensitively on **suppressed inside the `<summary>` heading**, not anywhere in the body. See [fully-clean.cases.md](#), “The collapsed-block case (Morrison-Lab/ai-config#1029)”.
- **“No verdict” is its own state, distinct from “a verdict with no findings”.** A review job can fail having posted *nothing* — not a stub, not an empty comment. Zero findings and zero review are indistinguishable by any count, and they call for opposite responses: one is done, the other needs a self-review and a re-run. Read the job’s step outcomes when a review is missing rather than inferring from the absence of comments.
- **The notification that wakes you carries a SUBSET of the findings, and nothing in it says so.** Every case above is a surface *on GitHub* that a query can reach. This one is the channel that tells you to look in the first place: a `pull_request_review_comment.created` wake delivers **one** comment, and a review posting five of them wakes you five times, asynchronously, with no count and no “1 of 5”. So the first wake is indistinguishable from the only wake, and acting on it reads as responsive while leaving the rest unaddressed. It is worse than an ordinary partial read because the thread then *looks* handled: a reply and a resolved thread sit under the one finding you saw. Re-fetch `get_review_comments` on every review wake and act on the whole set, never on the wake’s own payload.
- **Do:** read all review surfaces before calling a PR clean, every round, including collapsed suppressed-comments blocks.
- **Do:** distinguish “no findings” from “no verdict” explicitly, and treat the latter as unreviewed.
- **Don’t:** report clean on a zero thread count, however many checks are green.
- **Don’t:** treat an empty review body as an all-clear without checking the inline comments.
- **Don’t:** treat a “generated no new comments” overview as an all-clear until every `<summary>` heading has been checked case-insensitively for **suppressed** — not until the whole body has, which flags ordinary overview prose that merely mentions suppressed findings.
- **Don’t:** read a reviewer’s silence as a verdict — a job that posted nothing leaves the same zero counts as a job that found nothing.

- **Don't:** act on a review wake's own payload — it is one comment out of however many the round posted, and it never says which.

A comment can be evidence-dense, correct throughout, and state no verdict at all — and its density is what gets read as the conclusion.

A later comment stating no verdict does not supersede an earlier one.

- **Do:** identify the last statement that actually states a verdict, and treat that as the standing one.
- **Do:** scan the whole review history for it, not only items matching HEAD.
- **Don't:** read a verification section, however rigorous, as an approval — it is evidence, and a verdict is a conclusion about evidence.
- **Don't:** treat a later comment's silence on the verdict as superseding an earlier “Needs more work”.

See [fully-clean.cases.md](#), “A later comment stating no verdict does not supersede an earlier one”.

A reviewer skip notice (e.g. for workflow edits or quota exhaustion) does NOT clear or supersede prior review findings.

When a review run skips (e.g. self-modification workflow guard or quota limits) and falls back to a self-review or human review per [self-review-fallback](#), that fallback lets the ARDI iteration loop proceed in the absence of prior unresolved findings. It never satisfies an autonomous merge gate — autonomous merging under `mwc` remains blocked per the merge gate above, and human approval is the only path. It does NOT wipe the slate clean, does NOT license merging over an unaddressed **Needs more work** verdict or open finding list from an earlier or concurrent review run, and does NOT clear the all-clear merge gate above.

- **Do:** scan the complete PR review comment history for any **Needs more work** verdicts or open finding sections before declaring a PR clean or ready to merge.
- **Do:** address, rebut (with convincing acceptance), or defer every previously raised finding even if the most recent review run skipped.
- **Don't:** treat a reviewer skip notice or self-review fallback as an all-clear or as permission to ignore open findings on the PR.

Another surface, and the one that defeats the gate itself: the review check can pass on a blocking verdict.

- **Do:** grep the verdict body for its own conclusion, and treat a **require-review** pass as orthogonal to whether the PR is clean.
- **Don't:** let a green review-gate check stand in for reading what the review said.

`check-pr-fully-clean.py` itself has the mirror false positive: it can report NOT clean over a clean verdict.

- **Do:** read the verdict’s own conclusion when the script reports findings against a review whose prose merely discusses finding vocabulary.
- **Don’t:** treat a `contains findings (matched pattern ...)` line as a real finding without reading the verdict body it matched.

Calling the checker is not consuming it: grepping its PROSE instead of reading its EXIT STATUS re-opens the whole failure one layer up.

The rule above and `no-handrolled-verdict-parse.py` both govern *bypassing* the instrument. This is the case where you run it, correctly, on the right PR — and then decide what it said by matching a string in its output.

`check-pr-fully-clean.py` answers twice. It prints findings for a human, and it exits 0 for clean and non-zero otherwise. Only the second is a stable interface. The prose is free to gain a line, split across two lines, or word a finding differently, and every one of those silently changes what a `grep` decides.

Two properties make this worse than an ordinary parsing slip.

It fails toward clean. The natural spelling is a positive test for the bad state — `if output matches "NOT fully clean" then not-clean, else clean` — so *any* failure of the match, including the check erroring or printing its header separately, lands in the `else` branch and reports clean. A missed match and a genuinely clean PR are the same observable, which is `fail-fast`’s pass-path-equals-failure-path shape arriving through a tool built to prevent exactly this.

It launders. The report reads as the instrument’s verdict rather than as your reading of it, so “the checker says clean” is what reaches the human — and nothing in that sentence exposes that a `grep` stood between the two.

The status is three-valued, and collapsing it to a boolean is the same mistake one layer further in. `check-pr-fully-clean.py` exits **0** clean, **1** not clean, and **2** for a usage or environment error. That third code is deliberate — its own source says `USAGE_EXIT = 2` exists so “a usage or environment error would have been read as a verdict about the PR” — so `if ! checker; then not_clean` throws away the distinction the script went out of its way to provide.

The cost is a **false regression**: a transient `gh` failure, a rate limit, a network blip in a polling loop, all report a PR as having gone not-clean. That is the mirror of the `grep` bug above, which failed toward clean; this one fails toward alarm, and both are a two-branch reading of a three-branch answer.

This is the rule `errexit-is-not-uniform` states as 0, 1, and anything else being three answers and not two — itself a paraphrase of `fail-fast`’s hand-check guidance to treat 0 as found, 1 as clean, and anything else as the check having failed to run. It applies to a purpose-built checker exactly as it does to `grep`.

But 2 does not cover every non-verdict, so the three-way read is necessary and still not sufficient. `USAGE_EXIT = 2` is raised by `die()`, on the paths the script anticipated. An **unhandled exception** exits **1** — the code reserved for “not clean” — so a crash is indistinguishable from a verdict by status alone.

That is why the status read has to be paired with a look at the output rather than replacing it. A genuine not-clean prints - finding bullets; a crash prints a traceback. One `grep -q '^ - '` separates them, and unlike the phrase search above it is keyed on the report’s *structure* rather than on its wording.

The wrong-repo case is the one to expect, because the script resolves the repo from the **current working directory** unless `-R/--repo` is passed. A background poller inherits the session’s `cwd`, which on a multi-repo session is routinely not the repo the PR lives in — so the same command answers correctly by hand and crashes in the loop. Pass `-R OWNER/REPO` explicitly in anything that is not a one-off typed inside that checkout. See [fully-clean.cases.md](#), “Checker unhandled exception on wrong repo”.

A remote or web session has no gh at all, so the checker cannot answer there — and that is a property of the session rather than of the PR. The wrong-repo case above is a mistake you can stop making. This one is not: `check-pr-fully-clean.py` shells out to `gh`, and a remote/web Claude Code session has no `gh` on `PATH`, so the script refuses with ``gh` is not installed or not on PATH` and exits **2** on every invocation, whatever the PR’s real state.

That lands in the third branch of the read above, which is the right answer and an easy one to skip past, because the mandated instrument failing feels like a step to work around rather than a result to report. Two things follow.

Say that the checker did not run. [ardi](#)’s fully-clean exit checklist opens by requiring exit 0 from it, so reporting a PR clean without noting the substitution asserts a check that never happened. The substitution itself is ordinary — root `CLAUDE.md`’s “Skills that call `gh`/`glab`: fall back to `tool-mappings.md` in remote sessions” already governs it — so establish both criteria from the GitHub MCP surfaces instead: the paginated `check-runs` endpoint for criterion 1, the review body and thread list for criterion 2.

Do not read the 2 as a verdict in either direction. It is neither “not clean” nor a licence to assume clean. It is the check declining to answer, which is exactly what the three-valued read above exists to preserve.

- **Do:** state which surfaces supplied the verdict when the checker could not run, so “clean” stays attributable.
- **Don’t:** report the checklist item satisfied on a session where the script exits 2 — it did not run.
- **Don’t:** treat the refusal as a PR problem, or spend a round diagnosing it; the absence of `gh` is the whole cause.

The mirror case is worse: a review job running the checker on the PR it is reviewing gets an answer, and the answer is always not-clean. The section above is the checker declining to answer, which is loud and exits 2. This one exits 1 — a verdict — and it is wrong every time.

A review job querying its own PR’s check state observes itself as `in_progress`, and the checker requires that a completed automated review exist on the current head. The reviewer is the thing that would produce that review, so the predicate is unsatisfiable by construction: it cannot authorize itself in advance, at any round, on any PR, however clean the diff.

The failure is quiet in the way the `gh` case is not. Exit 2 announces itself as a non-answer; exit 1 is indistinguishable from a real finding, and it reaches the PR as a withheld verdict that blocks merge under `mw`’s Scope Limit — on a PR the same review just declared sound.

It also poisons the rounds after it. Re-triggering does not clear it, which is the natural first remedy and the wrong one: once a **Needs more work** comment exists for that head, the next round’s verdict scan reads a standing not-clean and the loop cannot converge. Each attempt costs another paid review.

The discrimination to make is between two different claims:

- *“The diff is sound, and I cannot observe my own completion”* — a statement about the **instrument**, which belongs in the report.
- *“This needs more work”* — a statement about the **diff**, which is the verdict field and the thing that blocks.

Routing the first into the second is the defect. It is the same conflation `verify-the-right-artifact` names elsewhere: an inability to measure is being reported as a measurement.

- **Do:** state the instrument’s status in the report, and keep the verdict a judgment about the diff.
- **Do:** exclude the calling run when a reviewer evaluates its own PR, or say the check was not applicable rather than running it.
- **Don’t:** put “cannot determine cleanliness” in the verdict field — it blocks merge and reads as a finding to everyone who does not read the reasoning.
- **Don’t:** re-trigger to clear it; the standing not-clean comment makes the next round worse rather than better.

(Measured 2026-08-27 on `ai-config#2442`. Three rounds found one hyperlink nit, one citation-order nit, then nothing. The instrument reported not-clean in all three, naming `review / claude-review` (itself) among the pending checks — and the verdict field carried that into a block in **two** of them. The third round said so explicitly: “not due to any content defect”, “no further content changes are needed on my end”. CI was 10/10 green and `require-review` passed. The three rounds cost \$4.60 between them, and the PR merged only on an explicit user decision to treat the verdict as non-blocking.

The round that did **not** block is the most useful of the three. It is the same reviewer, on the same PR, reaching the same instrument result and keeping it out of the verdict field: it returned **Ready for merge** and put the caveat in a note beneath it. So the discrimination this entry asks for is achievable rather than merely desirable, and the two rounds that blocked were not forced to by anything about the instrument. That round is also the reason to state the split precisely: the first draft of this record said all three blocked, and review caught it — an unverified count in prose, inside an entry about an unverified count. Tracked as [ai-config#2441](#).)

(Measured 2026-08-19 on a remote session driving [ai-config#1673](#). Tracked as [ai-config#1679](#), which weighs teaching the script a REST fallback against documenting the branch; until one lands, every remote-session ARDI run hits this.)

So read the status, and read all three of it:

```
python3 scripts/check-pr-fully-clean.py "$n" -R "$OWNER/$REPO" >/tmp/fc.txt 2>&1
rc=$?
case $rc in
  0) echo "#$n CLEAN" ;;
  1) if grep -q '^ - ' /tmp/fc.txt; then
      echo "#$n NOT clean"; cat /tmp/fc.txt
    else
      echo "#$n CHECK CRASHED (rc=1, no finding bullets) -- not a verdict"
      tail -3 /tmp/fc.txt
    fi ;;
  *) echo "#$n CHECK FAILED (rc=$rc) -- not a verdict"; cat /tmp/fc.txt ;;
esac
```

- **Do:** branch on the checker’s exit status, treating 0 as clean, 1 as a verdict of not-clean, and anything else as the check having failed to answer.
- **Do:** re-verify the agent and the head yourself before reporting ready, since the exit status is necessary and this file’s own SHA-surface caveats still apply.
- **Don’t:** grep a purpose-built checker’s output for a phrase — its prose is a human-facing report, not an API.
- **Do:** pass -R OWNER/REPO from any poller or script, since the repo comes from the working directory otherwise and a background loop inherits whatever cwd the session happened to be in.
- **Don’t:** collapse the status to a boolean either; `rc != 0` reports a broken check as a regressed PR, which is the same conflation wearing the remedy’s clothes.
- **Don’t:** read 1 as a verdict without checking the output has finding bullets — an unhandled exception exits 1 too, so 2 is not the only non-verdict code.
- **Don’t:** read “I called the right instrument” as having consumed it; the bypass guard fires on the call, and nothing fires on the misreading.

See [fully-clean.cases.md](#), “Three PRs reported clean by grepping the checker’s own output”.

Exit 0 is not the whole answer either: read the verdict scan: line the checker prints, because it can say 0 bore a verdict, latest = NONE on a run that exits clean. The three-way read above governs every status that is *not* 0, so it cannot reach this one — the false clean arrives as exit 0, the one value nothing above tells you to look behind. `check_latest_verdict()` blocks on `not-clean` alone, and an empty verdict is not `not-clean`, so a head reviewed by nobody takes the clean return. A reviewer’s own **skip notice** is enough to occupy the slot.

- **Do:** read the `verdict scan:` line on every invocation, including the ones that exit 0.
- **Do:** treat `latest = NONE` as no review at all, and fall back per [self-review-fallback](#).
- **Don’t:** read exit 0 as “a reviewer approved this” — it says only that nothing blocking was found, and an empty review history finds nothing.
- **Don’t:** count a skip notice as the review; it is admitted as a review item and states no verdict, which is exactly the state that exits 0.

The author filter gates formal reviews and not comments, so a human-authored comment enters that same scan on body text alone. The comment loop admits on `is_bot_author` or `is_review_header`, and `is_review_header` matches `### verdict`, `verdict:`, and `code review` with no author check — so your own disposition comment, or any reply quoting a reviewer’s verdict line, can be counted as a review item. Reading the formal-review loop and generalizing its author check to comments is [verify-the-right-artifact](#)’s “a neighbour for the target” shape applied to source.

- **Do:** read the loop that handles the artifact class you are making a claim about — comments and formal reviews are separate populations here.
- **Do:** check a comment’s admission against its body markers, not its author.
- **Don’t:** generalize one loop’s filter to a neighbouring loop in the same function.
- **Don’t:** read “no human comment appeared in `matching_items`” as evidence that human comments are excluded; the SHA test is what excluded it.

See [fully-clean.rationale.md](#) for both mechanisms, and [fully-clean.cases.md](#), “A skip notice exits the checker clean over an empty verdict scan”.

A classifier written to EXCLUDE driver-status comments from the verdict scan is itself a matcher, and its negative guards protect only the dialect they were written against. The section above is a comment wrongly *admitted*; this is the mirror, a comment wrongly *dropped*. A driver-ledger classifier recognizes a session’s own status comment (claim wording, an ARD disposition table, a self-imposed hold like “hold off ...”) from broad English markers, then protects genuine reviews with negative guards – a `### Verdict` heading, a `Reviewed-Commit:` fingerprint, a `**Claude finished` marker – that must all be absent before the exclusion applies. Every one of those guards is keyed on Claude’s and Cursor’s own report structure. A Copilot review comment carrying a real, blocking finding phrased as “hold off

on merging until X is added” emits none of that structure: the broad marker matches, every guard abstains, and the finding is dropped from the verdict scan entirely – reported FULLY CLEAN. This is [fail-fast](#)’s “Guarding an unsound pattern with a second pattern, rather than replacing it” and “A guard’s discharge fires on positive success, not the absence of failure” sections, arrived at independently in this checker: negative guards defending an over-broad matcher inherit exactly the ambiguity the matcher already had.

Inverting the gate to a POSITIVE signature was tried next, and refuted the same day, which is the more useful half of the lesson. The candidate signature was the agent-disclosure marker, on the premise that every driver comment carries it and no reviewer report emits it. Neither half holds. [self-review-fallback](#) requires a dispatched or cross-vendor review to be published verbatim WITH that marker, so a genuine not-clean review carries it too, and gating on it dropped that review instead. And [disclose-agent-authorship](#) exempts a comment posted under a genuine bot identity, so even the first half is a convention this corpus asks for rather than a property a gate can rely on. Both attempts failed the same way. Every discriminator available in a comment body is one some real reviewer also emits, so no body-shape test can safely decide to DROP an item – and a positive signature is not safer than a negative one merely for being positive. A third design stopped dropping anything. It blanked the single shape that actually caused the misread — a prior round’s verdict quoted in a bare parenthetical after a cited SHA — inside `strip_cited_finding_vocab` instead. It was refuted too, on a body where the parenthetical IS the live verdict and the explanation follows it outside the blanked span. Nothing shipped in the checker: all three were reverted. The fix went to [ard](#)’s summary step instead (#2448) — a disposition comment backticks any verdict phrase it quotes, so the code-span rule #1202 already established neutralizes it, and the instrument gains no new fail-open surface.

- **Do:** prefer fixing the input at the author’s end over teaching the checker to guess – three checker-side designs were refuted here, and a pair of backticks was not.
- **Do:** derive, by execution, which line of a body actually produced the verdict, before building a classifier for the parts you assume did.
- **Do:** confirm a proposed signature’s population against every producer the checker sees, and treat “no reviewer emits this” as a claim to check against the corpus rather than a premise.
- **Do:** read a driver-comment classifier’s guard list as a dialect list, and ask what a differently-formatted reviewer’s report looks like against it.
- **Don’t:** protect an over-broad exclusion marker with negative guards keyed on one producer’s output format – they abstain on every other producer, which is exactly where the marker is most wrong.
- **Don’t:** read a positive gate as inherently safer than a negative one – both were tried here, and both dropped a real reviewer’s finding.
- **Don’t:** trust a driver-comment classifier’s 0 **dropped** (or silence) as evidence nothing was excluded; the failure here produces no error, just a lower “examined N items” count.

See [fully-clean.cases.md](#), “A driver-comment classifier drops a Copilot finding it has no

guard for”.

Dropping an item from the verdict scan is a distinct fail-open route from misreading one that IS scanned, and it leaves no trace in the output at all. Every case in this file up to here is about an item that entered the scan and was then misread – a stale SHA, a truncated body, a wrong author filter. A dropped item never enters the scan, so the “examined N items” line the checker prints simply reads one lower, which is indistinguishable from a PR that genuinely received one fewer review comment.

- **Do:** when a verdict scan reports fewer items than the PR thread has comments, ask what was dropped and why, not just what the scanned items said.
- **Don’t:** read a clean scan, however many items it examined, as evidence every review comment on the thread was considered.

A verdict comment quotes verdict phrases, so a phrase search identifies nothing — and it misreads in both directions at once.

- **Do:** call `check-pr-fully-clean.py` for a sweep’s verdict column, exactly as [ardi](#) requires for one PR.
- **Do:** anchor on the last `### Verdict` heading when parsing by hand, after selecting candidates on the `**Claude finished` marker.
- **Don’t:** take the first verdict phrase in a body as that body’s verdict — quoting other verdicts is part of what a review comment does.
- **Don’t:** assume such a misread has a safe direction; one sweep produced a false-clean and a false-blocked.

That “anchor on the last `### Verdict` heading” line describes the by-hand method, not what `check-pr-fully-clean.py` itself does — the script has no heading anchor at all. It matches verdict *phrases* with a regex (`Verdict:\s*(?:Clean|Approved|Ready)\b` and its not-clean counterpart), never a `^###\s*Verdict` heading line, so a doubled or malformed `### Verdict` heading in a review comment cannot break something the script never checks. Reading this fragment’s hand-parsing advice as a description of the script’s own mechanism produces a confident, wrong claim about our own tooling — worth naming because the fragment sits right next to the script it is easy to assume it summarizes.

- **Do:** read `scripts/check-pr-fully-clean.py` itself when the claim under test is about what the script does, even when this fragment already describes the by-hand procedure.
- **Do:** treat “anchor on the last `### Verdict` heading” as guidance for a human parsing a comment, distinct from the script’s own phrase-matching logic.
- **Don’t:** infer the script’s parsing mechanism from this fragment’s by-hand advice — verify against the script’s source before filing an issue that names a mechanism.

See [fully-clean.cases.md](#), “A fragment’s by-hand parsing advice mistaken for the script’s own mechanism”.

A review comment’s header SHA can be stale, so take the reviewed commit from the run’s own `head_sha`.

- **Do:** follow the job link in the comment and read that run’s `head_sha`.
- **Don’t:** treat the SHA in a comment’s heading as the commit reviewed.

That remedy assumes the run checked out the PR head, and a `workflow_dispatch`-triggered review run does not.

- **Do:** check a `workflow_dispatch` review’s `event` field before reaching for `head_sha` — on that trigger type the field names the dispatch ref, not the reviewed commit.
- **Do:** cross-check a stale-suspected verdict’s specific claims against the file directly, rather than only against run metadata.
- **Don’t:** trust `head_sha` as “the commit reviewed” on a `workflow_dispatch`-triggered run — that guarantee only holds for `push/pull_request`-triggered runs, which check out the PR head by construction.

A third surface names a commit the run never read, and unlike the two above it points the confident direction: the run object’s own `pull_requests[].head.sha`.

- **Do:** settle which commit a review read from a discriminating claim in its own body, since that is the only surface separating the candidates.
- **Do:** read `pull_requests[].head.sha` as a fact about the PR’s current head, useful for nothing else.
- **Don’t:** read that field naming your latest commit as evidence the review covered it — it names the current head unconditionally.
- **Don’t:** read an empty `pull_requests` as evidence about the run; the array empties when the PR closes.

See [fully-clean.cases.md](#), “`pull_requests[].head.sha` named a commit pushed after the run started”.

`check-pr-fully-clean.py` uses the same unreliable body-text surface, and whichever SHA that text happens to contain — present, absent, or wrong — is incidental to which head the run actually reviewed.

- **Do:** read a flagging run’s `event`, `head_branch`, and `head_sha` before treating “no review at this HEAD” as a genuine gap.
- **Do:** treat the script’s discharge as the likely reading, since the withholding direction dominates in practice, but not as a certified one.
- **Don’t:** re-dispatch a review, or fall back to self-review, on this signal alone when the flagging run’s own metadata already shows it evaluated the current head.
- **Don’t:** read a body’s SHA, present or absent, as evidence about which head a review covered — it is evidence about what the prose happened to discuss.

- **Don't:** conclude that reviewers citing their head SHA more consistently would fix this; a body can already cite a SHA and still be citing the wrong one.

A review citing a SHA that is not the PR's head is usually an intermediate commit, not a fabricated one — resolving it locally settles which. The cases above are about which SHA to trust when several are in play; this one is about a body citing a SHA that matches none of them, which first reads as a hallucinated citation. A round with more than one commit — a fix plus a later Merge origin/main on top, say — has a real, resolvable commit for the fix that is not the round's head, and a reviewer naming that commit is often the more useful citation of the two.

- **Do:** run `git cat-file -t <sha> && git show -s --format='%h %an %s' <sha>` on a cited SHA that does not match the head before concluding it is fabricated.
- **Do:** read a resolvable SHA that sits in the branch's own history as a legitimate citation of an intermediate commit, not a hallucination.
- **Don't:** treat a claim about **PR state** (“this was already merged”, “CI already passed”) the same as a claim about **which commit** did something — the former still deserves the API check this file opens with; a bare SHA mismatch does not.
- **Don't:** re-push a fix already present in the cited commit because the SHA did not match what you expected.

See [fully-clean.cases.md](#), “A review cited the fix commit, not the round's merge-topped head”.

A clean CI run and a clean review verdict are a snapshot, not a standing guarantee of mergeability. `main` can advance after your last check — including gaining its own independent addition that collides with yours (see [sync-with-main.md](#)'s “two PRs append the same numbered subsection” case) — so re-verify the branch still merges cleanly against current `main` before reporting a PR ready, not just trust the last green run.

`mergeStateStatus: CLEAN` means conflict-free plus passing commit status (GitHub's `mergeable` field), not merge-ready. A PR without a clean review verdict on the latest commit is not merge-ready.

- **Do:** always check for merge conflicts (e.g., using `gh pr view <number> --json mergeable` or `gh pr checks`) at the same time you check for CI and review status.
- **Do:** report a PR as blocked on review when HEAD has no authentic clean verdict, even if GitHub says `CLEAN`.
- **Don't:** treat green CI plus a clean review as sufficient without independently re-checking merge-conflict state.
- **Don't:** describe a PR that lacks a clean HEAD review as merge-ready, ready to merge, or “green and merge-ready.”

A sync-only push invalidates a clean verdict just as thoroughly as a code push, and arming auto-merge after a sync violates the HEAD review gate. When `main` moves

and a direct merge is refused because the branch is not up to date, merging `origin/main` in and pushing creates a new HEAD commit ref. Arming `gh pr merge --auto` immediately after that sync push — reasoning about it as scheduling a merge already verified rather than authorizing an unreviewed head — violates [Pattern 12](#). GitHub auto-merge fires the moment CI passes, racing ahead of and potentially merging before any automated or adversarial reviewer can evaluate the new HEAD commit. The sync is content-free (no author code changes), which is why it does not feel like a new head needing a new verdict, but the new HEAD commit ref is completely unreviewed until a fresh review round posts for that exact SHA.

- **Do:** re-run `scripts/check-pr-fully-clean.py <N>` against the new HEAD commit after any sync push, wait for clean reviews and green CI at that HEAD, and merge directly/synchronously.
- **Do:** accept that a fast-moving `main` may require repeating the sync-and-verify cycle rather than attempting to bypass it with auto-merge.
- **Don't:** arm `gh pr merge --auto` after a sync-only push under the impression that prior verification at an older commit carries forward.
- **Don't:** assume GitHub auto-merge will wait for review comments — it gates only on native branch protection checks.

See [fully-clean.cases.md](#), “Auto-merge armed after a sync-only push, having verified the previous head (#2556)”.

Re-check version parity in that same sweep, not only conflict-freedom.

Threads: at fully-clean, every **inline** review thread is resolved, and the only conversation left open is the final all-clear exchange — the reviewer’s all-clear comment and your reply to it. (The all-clear is usually a top-level PR comment, not an inline thread.)

One finding can own two threads, so sweep by thread id rather than by finding.

Deadlock -> escalate to a human. If you and the reviewer(s) can’t reach consensus on an item (a rebuttal was exchanged and neither side is budging), don’t loop forever and don’t unilaterally override the reviewer — request a **human reviewer**, @-mention them in a comment summarizing the impasse, and surface the open item.

An automated reviewer’s verdict on a disputed factual/technical claim is not stable across independent runs, even with identical evidence available each time. Don’t treat one round’s “settled, no need to keep arguing” as durable: the very same review job, re-triggered later with no new code changes, can re-raise a claim it previously retracted — and then retract it again on a subsequent run — purely from re-deriving the question differently each time, not from anything changing in the PR. This means a rebuttal thread’s outcome (however many rounds of citations and counter-citations) doesn’t itself resolve a genuine deadlock the way a human’s decision does; only escalating per the bullet above actually settles it. The one thing that DOES help going forward: fold the authoritative citation/evidence directly into the code or doc being reviewed (a comment, not just a PR conversation reply) — a fresh reviewer run re-deriving the claim from scratch is more likely to find the citation sitting right next to

what it's evaluating than to dig through prior thread history for it, though even that is not a guarantee against a bot that ignores context already in front of it.

Algorithmic safeguards: Algorithmic checks and hooks can only invalidate, not validate, a PR. You still need to use your own judgment in addition to satisfying the algorithmic safeguards; they are a safety net, not a gold standard.

6 Address Every In-Scope Review Comment

When iterating on a PR with a reviewer, **address every in-scope flagged item**, regardless of severity label. The reviewer's "Informational", "Not a blocker", "minor", "nit", "optional", "consider", or "if you want" labels are for prioritization, not a free pass for the implementer.

Extended rationale — the mechanism, evidence, and argument behind each rule below — lives in [address-every-comment.rationale.md](#), moved out of the auto-loaded context. Each rule here keeps its statement and its Do/Don't pair; read the companion when the reasoning or the evidence is the question.

Worked-example case records for the rules below live in [address-every-comment.cases.md](#), moved out of the auto-loaded context.

For each flagged item, do exactly one of:

1. **Fix it in this PR.** The default path — most nits are 1–3 line changes.
2. **Defer.** Only when the fix expands the PR's scope (new feature, broader refactor, separate concern), the requester has explicitly said this PR shouldn't grow, or the flagged content isn't actually yours to fix here (see the `main-sync` case below). File a follow-up issue and reference it in a PR comment so the item isn't lost — except in the `main-sync` case, where the "follow-up" is fixing it on `main` directly, not a new issue.

Then trigger another review and repeat until the PR is **fully clean** — zero flagged items under any heading, no "informational", "non-blocking", "harmless", "minor observation", or "could improve" sections. "Looks good" / "no findings" / "approved" with no follow-on bullets is the bar. Resolve every inline review thread along the way, leaving only the final all-clear exchange.

Always resolve an inline thread the moment its comment is successfully addressed — the fix pushed and a reply posted naming it — in the same pass, whatever workflow you're in: a formal `ard/ardi` round, a CI-monitor nudge, or a one-off fix outside any loop. Addressing without resolving leaves a thread that reads as outstanding work to every later reviewer, blocks [fully-clean](#)'s every-inline-thread-resolved criterion, and drags stale noise into the next review round. The per-disposition settlement rules in `ard` step 4b still govern the exceptions: a **Rebut** stays open until the reviewer drops it, and an **Address** you're not confident fully settles the concern gets a reply asking for confirmation instead of a resolve. The `resolve-pr-threads`

skill sweeps any stragglers, but it's a backstop — resolve-on-address is the default, not a cleanup step.

Do **not** report “ready to merge with one minor nit noted” / “harmless as-is” / “can address if you want” — that hedging just pushes triage back to the requester.

A round count is never a reason to stop, and “the reviewer keeps finding things” is not a finding about the reviewer. There is no threshold after which unaddressed items become acceptable: keep requesting reviews and keep dispositioning findings until a review comes back with none. The only exits are a totally clean review, a genuine per-item deadlock, or the user calling it. Reasoning of the form “we have done N rounds, shall we accept the current state?” is the same hedging this paragraph bans, moved up from one finding to the whole loop — see [ardi](#)'s “Stopping conditions” for why it fails and for the case record.

Noise is per-item, not per-round — don't stop the whole loop over one recurring flag. A long-running PR can have both real findings (worth fixing every round) and one specific item the reviewer re-raises verbatim round after round even though it's already deferred/tracked (e.g. a file-length guideline already split into a follow-up issue). Keep fixing every *new* finding as it appears — don't let the recurring item make you stop processing genuinely new ones. But stop re-litigating *that one item* every round: reply once pointing at the tracked issue, and hold on it specifically rather than re-deferring it on each pass. Surface the pattern to the user (which item, how many rounds, where it's tracked) and let them decide whether to resolve it now (e.g. do the split) or leave it as accepted recurring noise — don't decide unilaterally to either keep re-processing it or silently drop it.

When a finding is a pattern (a formatting/style rule broken in one spot), apply it everywhere it recurs in the same file, not just the flagged line.

That rule's scope is “the same file”, and a reviewer who enumerates the sites is the reason the scope goes unquestioned.

- **Do:** derive the site list by grepping the whole diff for the flagged phrase, and fix what the grep returns.
- **Do:** report the sweep — the pattern searched and the hit count — rather than the number of sites you fixed.
- **Don't:** treat a reviewer's enumeration as the extent of the pattern; it is the extent of that reviewer's read.
- **Don't:** write “all N spots you named” into a reply, since quoting the reviewer's count is the tell that no sweep ran.
- **Don't:** read a null result as “no further sites”; it means no further hit for that pattern, and a differently-worded instance would not have matched.

The remedy above names a search space — “the whole diff” — and a stack has one of those per branch.

- **Do:** run the derived sweep over every branch in the stack, and report the per-branch counts.
- **Do:** treat a finding that names a convention as scoped to the work rather than to the PR it was filed on.
- **Don't:** read “the whole diff” as satisfied by the diff of the PR the reviewer commented on — that is one of N.

See [address-every-comment.cases.md](#), “A finding’s site list spans every branch in the stack”.

Deriving the class is necessary and not sufficient, because you can derive the wrong one — and the growth rate across rounds is what says so.

- **Do:** read growth in the list across rounds as evidence about the lever, not as a count of members still to add.
- **Do:** state a reviewer’s diagnosis back in your own words before acting on its examples, so a redirect cannot be worked past in silence.
- **Do:** relax the enumeration and read which cases move — a property shared among them names the axis.
- **Don't:** treat “I derived the class rather than fixing the reported instances” as discharging the rule above; it is that claim one level up, and it fails the same way.
- **Don't:** prefer the actionable half of a review to the diagnostic half merely because it is the half you can start on.

A narrower version of the same failure: the class is right, and it is enumerated in more than one place.

- **Do:** paraphrase the last two or three findings into a single sentence, and read a match as evidence that a concept is duplicated rather than incomplete.
- **Do:** derive how many sites encode the concept, then consolidate them into one definition every site consumes.
- **Don't:** answer a third instance by extending a third list — that is the same round again with a new door.
- **Don't:** skip the review’s own prose naming a sibling site; it is frequently there, in the paragraph explaining why some other mechanism did not save you.

The mirror case: the enumeration was complete and the fix was not.

- **Do:** count the artifacts a single comment names, and give each one its own disposition before replying.
- **Do:** read the rendered page rather than the diff when confirming that a prose or formula fix landed completely.
- **Do:** grep the whole file for the underlying concept once a second half surfaces — a document stale in two places is usually stale in three.

- **Don't:** let a visibly-changed flagged line stand in for the finding being closed; the unfixed half appears in the diff as context.
- **Don't:** reach for the derive-the-site-list remedy above here — that list was complete, and the shortfall was in the delivery.

When a prose fix changes wording that's also paraphrased elsewhere in the same PR (a CHANGELOG entry, a PR description, a cross-reference), sync that copy too. A CHANGELOG entry written before the review lands often quotes or paraphrases the exact phrase a reviewer later flags; fixing the source prose but leaving the paraphrase stale reintroduces the same wording issue one file over. Grep the diff for the flagged phrase before considering the finding closed.

A scope-widening fix makes its stale copies invisible to every diff-scoped sweep, so there the search space is the whole file, not the diff. The whole-diff rules below are right for a fix that changes a claim's wording: the synced copies entered the diff when you edited them. A fix that *broadens* a concept's scope inverts that. The restatements that are now too narrow are precisely the lines the fix did **not** touch, so they appear in the diff as context or not at all, and an added-lines sweep structurally cannot see them — it reports a confident zero over exactly the population the finding is about. Grep the whole file (and any file restating the concept) for the concept that widened, and read each hit against the new scope.

- **Do:** after a broadening fix, sweep every restatement of the concept in the whole file, not the diff's added lines.
- **Don't:** read a clean added-lines sweep as closing a broadened-scope finding — the stale copies are unchanged lines by construction.

(Morrison-Lab/ai-config#1490, 2026-08-15/16: rounds 1, 2, and 4 each found a sentence still Copilot-only after the surrounding passage was broadened to cover human reviews. The round-2 fix swept the diff's added lines for Copilot — 17 hits, all legitimately Copilot-specific — and round 4 still found an un-broadened copy in the **## Output** section, because that copy was an unchanged line the added-lines sweep could never have matched.)

When syncing copies, search the diff for the claim, not the files or symptom already in front of you.

- **Do:** run whole-diff searches for synchronized figures and phrases, after committing the fix, and report the before/after counts.
- **Do:** when a rationale is retired, search for every wording that states that rationale or criterion, not only for the symptom word that made it fail.
- **Don't:** substitute `grep -rn <term> <files-you-had-open>` for grepping the diff.
- **Don't:** accept a search for the visible contradiction as proof that the retired claim itself is gone.

That same search settles a narrower question about placement: a correction written NEAR the flagged sentence reads as having replaced it, while the flagged sentence

survives and the file then states both. Edit the sentence the reviewer named, and use the search above — for the claim, not for the symptom — to confirm that wording is gone.

- **Do:** delete or rewrite the flagged sentence itself, rather than adding a truer one beside it.
- **Do:** mark superseded text as superseded, explicitly, where it is worth keeping as a record of why something was done.
- **Don't:** read “the file now contains a true sentence” as having addressed a finding about a false one.
- **Don't:** let a commit message assert a deletion that the diff shows as an addition beside unchanged text.

See [address-every-comment.cases.md](#), “A correction added beside the flagged sentence, which survived”.

When the wrong thing is a figure, the unit of repair is the figure — across every artifact carrying the twin, not just the diff.

And a reflow puts its neighbouring sentences into your change, for fact-checking and not only for lint.

- **Do:** grep for the figure’s value across every artifact carrying the twin, before replying that the finding is closed.
- **Do:** fact-check the sentences a reflow pulled into your diff, exactly as you would the ones you wrote.
- **Don't:** treat the named occurrence as the unit of repair when the same value appears elsewhere.
- **Don't:** read a clean whole-diff grep as covering a twin the diff never touched.

See [address-every-comment.cases.md](#), “The unit of repair is the figure, across every artifact carrying the twin”.

The PR description is on that list and is the one copy grepping the diff cannot find, so check it separately.

- **Do:** re-read the PR description after any Address that changes what the PR does or why, alongside the changelog check above.
- **Don't:** treat a clean **grep** over the diff as evidence every paraphrase is synced — the description was never in it.

Answering a body-staleness finding with a correction comment does not clear it, and this corpus’s own visible-correction convention is what makes that move attractive.

- **Do:** edit the body **and** record the correction inside it, so nothing is silently overwritten and earlier rounds still resolve.

- **Don't:** answer a body-staleness finding with a comment — the next reviewer re-reads the body, so the finding survives it.
- **Don't:** treat the drift risk in rewriting a long body as a reason to leave it; re-deriving every figure is what the round already requires.

See [address-every-comment.cases.md](#), “A body-staleness finding is answered by editing the body”.

A body that reports volatile external state goes stale with no edit of yours, so the trigger above never fires on it.

- **Do:** describe the change, and let CI report CI.
- **Do:** timestamp and scope any status you must state (“red as of <sha>, cause was X”), so it cannot be read as a present claim.
- **Don't:** put current CI status, mergeability, or a blocker in a PR body undated — the body is the one place nothing re-measures.
- **Don't:** expect the “after any Address” trigger to catch it; that fires on your edits, and this goes stale without one.

Following that “state it as history” advice is what produces the next block, because an automated reviewer reads the body as a flat statement of intent.

- **Do:** state the current content first, marked as current, before any history.
- **Do:** put the reversal in its own section that opens by saying it is history.
- **Do:** make sure the “what is excluded” section does not name the reversed item at all, in any tense.
- **Don't:** rely on past tense alone to carry the distinction.
- **Don't:** revert a maintainer-requested change because a reviewer read the history as current — rebut, and escalate rather than comply.

The same sync is needed when the review fix is to CODE BEHAVIOR rather than to wording — and that case is easier to miss, because nothing about fixing a bug points at the changelog.

Tighter still: a changelog entry can contradict its own commit message, in the same commit, with no review in the loop at all.

One step further back: a figure inherited from the tracking issue is both the copy git keeps and the copy nobody verified.

- **Do:** re-run the check when a figure moves from an issue into a commit message, even having verified it once for the PR body.
- **Do:** read `git log -1 --format=%B` before pushing, against the same source the body's claims came from — a commit message is not greppable from the working tree once written.

- **Don't:** copy a count, version, or path out of the tracking issue on the strength of having written that issue.
- **Don't:** treat “permanent in history” as settled while the PR is unmerged — `git commit --amend` still works, and is usually worth a fresh CI round against a wrong figure reaching `main`.

A corollary for checking any of this in a semantic-line-break corpus: a single-line `grep` returns false negatives on your own prose.

Inline markup breaks the same search, and that variant aims the false negative at someone else's work rather than your own.

- **Do:** account for inline markup as well as whitespace before concluding a quoted phrase is absent — see the next block for which side to normalize.
- **Do:** read the single hit when a search for a citation's target returns only the citation itself.
- **Don't:** file a dangling-citation issue while the only evidence is a literal `grep` that found nothing but the citation — that is the search failing, until a normalized one agrees.

Apply whatever normalization you choose to the search term as well as to the text, or the fix produces a third false negative of its own.

- **Do:** normalize the needle with the identical function applied to the text, so the comparison is between two transformed strings.
- **Do:** re-test any earlier absent verdict after extending a normalizer, since the extension can break a term the previous version matched.
- **Don't:** enumerate which markup to strip and treat that list as the fix.
- **Don't:** test a raw search term against normalized text, however plain the term looks.

Symmetry is necessary and not sufficient once the haystack is source code, because a line-comment leader is inserted by the medium rather than by the author. The rule above governs inline markup — backticks, asterisks, underscores — which an author types *inside* a phrase, so stripping it with a character class is the right shape. A `##`, `#`, `//`, or `--` leader differs in two ways that each defeat that class. It appears at a **line start** rather than mid-token, so it interrupts a phrase only where the phrase happens to wrap. And `#` is not in the class at all, so applying the same normalizer to both sides leaves it in the haystack and absent from the needle — which is exactly the asymmetry the rule was written to remove, arriving through a character nobody enumerated.

The failure direction is the expensive one. A verbatim phrase that *is* present reports absent, so the natural response is to re-add content that was never missing.

Widening the class is the wrong repair, and the Do/Don't block one paragraph up already says so: **don't** enumerate which markup to strip and treat that list as the fix. The rationale companion puts it more sharply — “Enumerating is the wrong shape, not merely an incomplete

it to dodge a real finding — a redundant config line a reviewer flags is a real finding with an implied fix request, so it needs a real disposition, not a label that means “no change requested.” When a reviewer flags something (a redundant config line, a stale pattern) inside a file your branch only touches because you merged `main` in to resolve a conflict, check provenance before fixing it: `git log/git blame` the flagged line, or just compare against `origin/main`’s current content. If it’s identical to `main`, “fixing” it on your branch alone doesn’t fix anything — it just makes your branch disagree with `main` on unrelated content the next person to touch that file will have to reconcile again. Reply agreeing the finding is correct but out of scope for this PR, and leave it for whoever owns that file’s actual content to fix on `main` directly — no follow-up issue needed, since the fix target is `main` itself, not this PR’s own change.

This generalizes to a skill’s own inline restatement of a fragment it links to. A `SKILL.md` that links a backing `shared/` fragment for the full detail often *also* restates the fragment’s approach or word list inline (in its `description` field, or a short procedure-step summary) so a reader doesn’t have to open the linked file. Fixing a bug in the fragment doesn’t automatically fix these inline restatements — they’re a second, independent copy of the same claim, and a review round after the fragment fix can catch them going stale exactly like a `CHANGELOG` paraphrase does. Grep the whole PR diff for the fixed phrase/word-list, not just the fragment file, before considering a fragment fix complete.

A bot that re-raises an item as “not addressed” may simply not have seen your reply — check the timestamps before treating it as an impasse. An automated reviewer gathers the PR’s comments once, when its run starts. A rebuttal posted after that snapshot is invisible to it, so the next round reports the item as still open and unaddressed even though a substantive reply is sitting in the thread. The tell is a re-raise that repeats the original finding verbatim and speaks only to whether the *code* changed, without engaging any argument you made. Before escalating, compare your reply’s timestamp against the review run’s `started_at` (`gh run view <id> --json startedAt`, or the `started_at` field each run carries in `get_check_runs` when `gh` is absent): if the reply landed after the run began, it is a stale re-raise, not a genuine disagreement.

Reply-first collides with citing the fix’s SHA, and the way out is to commit between them rather than to pick one.

1. **Commit** the round’s fixes. The SHA now exists and is stable.
2. **Reply** on each thread, citing that SHA.
3. **Push**. The next review’s snapshot already contains the replies.
 - **Do:** commit, reply citing the committed SHA, then push — in that order.
 - **Don’t:** treat “I need the SHA for the reply” as a reason to push before replying; that is the ordering the bullet above exists to prevent.

A finding can be right while its suggestion block is wrong — verify the suggested literal before applying it.

The same check applies to a fix a reviewer describes in prose rather than in a suggestion block, and the sharpest test is the reviewer’s own example.

A reviewer’s corrected citation is another factual claim, so verify the replacement before adopting it.

- **Do:** verify a proposed replacement citation with the source’s own history before editing the PR to use it.
- **Do:** use `git log -S "<exact line>" -- <file>` or an equivalent provenance query when the question is which PR introduced text.
- **Don’t:** adopt a reviewer’s corrected issue or PR number because the original was wrong.
- **Don’t:** use word overlap and same-day timing as a substitute for source history.

The same check one artifact over: a reviewer’s replacement DIFFSTAT is a factual claim too, and the usual way it goes wrong is summing per-commit churn rather than diffing the merge base.

A branch that edits the same lines across review rounds accumulates churn. Round 1 adds a line and round 2 rewrites it, so a per-commit sum counts that line twice and reports a deletion the merge-base diff never sees. The inflation is therefore worst on exactly the branches most likely to carry a verification table worth checking, which are the multi-round ones.

What makes this cost more than one wrong number is that [ardi](#)’s pre-push checklist already requires every figure in a PR body to be re-derived by command at each push. A reviewer supplying replacement figures looks like that derivation having been done for you, so the natural move is to paste them straight in. That substitutes an unverified figure for a stale one and leaves the body just as wrong, while feeling like the finding was addressed.

- **Do:** re-derive a reviewer’s replacement figures with `git diff --numstat <merge-base> <head>` before pasting them into a PR body.
- **Do:** cross-check against GitHub’s own `additions/deletions` fields, which are computed against the merge base and so agree with that command.
- **Don’t:** treat a reviewer’s supplied figures as discharging the re-derive requirement — a correct finding about staleness says nothing about the replacement’s accuracy.
- **Don’t:** sum per-commit `--numstat` to get a branch’s diffstat. On a multi-round branch that double-counts rewritten lines and reports deletions the merge base never sees.

See [address-every-comment.cases.md](#), “A reviewer’s replacement diffstat summed per-commit churn”.

The same discipline runs the other direction, over findings you produce rather than receive: [quotable-findings](#) drops any finding that cannot quote the passage it is about, with a carve-out for findings about an absence.

The highest-yield version of that check: when a comment names an edge case in its own prose and also supplies a fix, run the fix against that edge case.

- **Do:** check a suggested fix against every failure mode the same comment names, before checking anything else about it.
- **Do:** name the reviewer’s own caveat in the reply, so the rebuttal rests on their evidence rather than on your say-so.
- **Don’t:** let a comment’s demonstrated thoroughness transfer to its snippet — they are separate claims.
- **Don’t:** discard a finding because its fix is wrong; the half that named the hazard usually still stands.

A quieter variant: the suggestion introduces no defect at all, it restates the line above it — so applying it deletes coverage while reading as hardening.

- **Do:** evaluate the suggested predicate and its neighbours on real input, and keep the finding while rejecting the snippet when they coincide.
- **Do:** fix the underlying coupling instead, and say in the reply why the suggested form was set aside.
- **Don’t:** accept a **suggestion** block that restates an adjacent check — passing tests afterward prove nothing, since the survivor passes for both.
- **Don’t:** read a reviewer’s own “the line above already covers this” as support for their replacement.

A finding can be right, and its fix adequate, while the *reason* it supplies is too weak to ship — and in a corpus of rules, the reason is the deliverable.

- **Do:** read the primary source for the strongest reason before adopting a suggested rationale, even when the suggestion’s conclusion is right.
- **Do:** say in the reply which reason you took and why the offered one was set aside, since deviating from a **suggestion** block silently reads as having missed it.
- **Don’t:** accept a defensible-sounding mechanism because the conclusion it supports is correct.
- **Don’t:** treat this as grounds to reject the finding — the conclusion usually stands, and only its reason needs strengthening.

And the mirror case: a finding can be wrong on its stated grounds while still pointing at something real.

A third direction, which evades the verification reflex rather than lacking a rule: agreeing with a finding and then escalating it.

- **Do:** verify an escalation against the full scope it claims, which is wider than the scope the finding reported, and which the finding’s own instrument may already cover.
- **Do:** post the correction to the thread that carried the escalation.
- **Don’t:** treat agreeing-and-extending as exempt from the checks a rebuttal gets, since agreement suppresses the reflex that disagreement triggers.

- **Don't:** report a finding as understated on a measurement you have not shown covers the whole field set.

When a finding cites a source, read the cited source before reproducing anything — it is the cheaper instrument, and it is the one that can show the finding backwards rather than merely unsupported.

When a reviewer hedges a finding because it depends on code it cannot see, check whether *you* can see it — the hedge is an invitation, not a verdict.

Timestamp the evidence before rebutting a finding with it — during a live incident, a log from twenty minutes ago describes a different system.

A rebuttal's own evidence is the least-checked claim in a review round, and the commonest way it goes wrong is being measured through a tool that adds a shell layer.

- **Do:** write each command spelling to its own file and run the file when comparing them, so exactly one shell layer applies.
- **Do:** hold your own rebuttal to the standard you would apply to the finding, and say which instrument produced the counter-measurement.
- **Do:** re-run the measurement outside the harness when a reviewer holds their ground, before rebutting a second time.
- **Don't:** read a rebuttal as self-verifying because disagreeing felt like the rigorous move.
- **Don't:** cite a named check as settling a question without saying what it ran through; a named check reads as a performed one.
- **Don't:** compare two command spellings by typing both into the same tool, which is the one measurement guaranteed to make them look alike.

A finding carries a timestamp too, and its precondition can dissolve between the round that raised it and the round that addresses it.

The claim is not thereby fixed either, which is the half that is easy to miss.

- **Do:** re-check what a finding presupposes at the moment you address it, not at the moment it was raised.
- **Do:** replace a claim whose gate has cleared with a derived one — the timestamp and the figures the event produced — rather than hedging it or leaving it.
- **Don't:** apply a reviewer's suggested wording without re-checking what that wording presupposes; a hedge is false once the thing it hedges has happened.
- **Don't:** read a claim that has become true as a claim that has been checked — nothing verified it, and the two read identically.

Four neighbours sit close enough to be mistaken for this, and the boundary is worth drawing because three of them fire on the same PR.

See [address-every-comment.cases.md](#), “A finding’s precondition can dissolve before you address it”.

A finding built on a *negative* result – “I searched and it isn’t there” – is only as strong as the paths that were searched, and the search scope is the part reviewers state loosest.

- **Do:** ask which paths a negative finding actually searched, and check the obvious location yourself before editing anything.
- **Do:** name the gap when the thing does exist – paths searched versus where it lives – so the same search is not re-run the same way.
- **Don’t:** accept “it isn’t there anywhere” as settled because it is stated more confidently than a positive finding would be.
- **Don’t:** discard the finding once its negative result is disproved – the thing it tripped over is often a real ambiguity.

A note the reviewer declined to raise is still a claim, and so is your refutation of it.

- **Do:** verify a declined, out-of-scope, or passing note against the code before either acting on it or writing it off.
- **Do:** hold the change regardless when the note turns out correct but genuinely optional — verifying decides what is true, not what ships.
- **Don’t:** treat a PR title, commit subject, or changelog line as evidence about what the code does; each states an intent, and a refactor can keep the very thing it says it replaced.
- **Don’t:** let your own refutation past the check you would have applied to the reviewer’s finding — it is a fresh claim, and overturning something feels like having verified it.

Refuting the mechanism a finding proposes is not refuting its claim, and the quality of the refutation is what hides the difference. The rule above governs a refutation nobody checked. This one governs a refutation that survives every check and is then made to carry a conclusion wider than itself. A finding usually arrives as a claim plus a route — there is a bypass, and here is the function that opens it — and disproving the route leaves the claim exactly where it stood.

The pull toward the wider conclusion is strongest when the refutation is good. Having measured the named mechanism against the real tool and watched it fail, the reply writes itself, and “there is nothing to exempt” reads as the finding of an investigation rather than as an inference drawn from one. Nothing further fires, because the thread is now closed.

The check is one question asked before replying, and it is [metacognitive-monitoring](#)’s cause test run backwards. That rule asks what else would explain an effect you are attributing. This asks what else would produce the effect you are denying.

- **Do:** restate a finding as its claim and its proposed route, and say which of the two your evidence reached.

- **Do:** look for a second route to the same effect before replying, starting with whatever the named function hands off to.
- **Don't:** let a measured disproof of a mechanism carry a conclusion about the claim — those are different propositions, and only one was tested.
- **Don't:** read a wrong mechanism as evidence the reviewer's instinct was wrong — a finding can be right about what happens and wrong about where.

(Measured 2026-08-22 on [ai-config#1911](#). A review reported a security bypass in `push_refspecs`, describing how `--repo=origin` would smuggle an unreviewed ref past the guard. Three `git push --dry-run` runs refuted that mechanism — an explicit positional repository overrides `--repo`, so the dropped positional was never a refspec — and the rebuttal concluded there was nothing unreviewed to exempt. The bypass was real and sat one function away in `_push_remote`, which took the remote from the positionals and then a config fallback chain and never consulted `--repo`. So `git push --repo=other` was graded against `origin` and allowed, while `git push other` was refused. An adversarial review of the rebuttal's own commit found it fifteen minutes later.)

A finding whose own REPRODUCTION does not run is not a finding refuted, and the first thing to check is which commit the reviewer read.

The section above governs your refutation of a reviewer's proposed route. This one governs the reviewer's own demonstration, and it is the harder failure to see, because here the evidence arrives as an error message rather than as an argument. A reviewer's example is pinned to the commit it was written against, and a branch under active work moves — so an arity, a signature, a path, or a fixture can be right where the reviewer stood and wrong at your HEAD. Running it at your HEAD and watching it raise reports your own checkout rather than their mistake.

Distinct from [verify-the-right-artifact](#)'s “a reviewer's failed reproduction is not a refutation”, which governs the opposite roles — a reviewer running YOUR reported case and seeing nothing. Here the reviewer supplies the case, and it is the case that errors.

- **Do:** check out the commit the reviewer read before concluding its example is broken — the review usually names it, and an active branch has probably moved since.
- **Do:** repair a genuinely broken repro to the smallest form that tests the claim, run that, and say in the reply what you changed.
- **Don't:** “repair” a repro that ran correctly at the reviewer's commit — that reports your own stale tree as the reviewer's error.
- **Don't:** read an exception from a reviewer's example as the finding failing, without first asking which tree it raised on.

See [address-every-comment.cases.md](#), “A repro that raised only against a newer signature”.

Count a round’s findings before pushing its fix, because disposing of one correctly generates no evidence about the others.

The rule above governs the finding you decline to act on. This governs the finding you never see, having already acted on its sibling.

A round can carry several findings, and acting on one produces every artifact that handling the whole round produces: a verified claim, a commit, a reply, a resolved thread. Completeness is a property of the **set**, so nothing in that sequence reports that a second finding existed. There is no moment that feels like stopping early, because each step was performed properly — which is why this needs a count rather than more care.

The body-only finding is where it hides, and [fully-clean](#) already names why: a finding about something the diff did not touch cannot be attached as an inline comment, so it appears in the verdict body alone. Inline threads produce a visible checklist and a body-only finding produces nothing to tick off, so “all threads resolved” reads as “round handled”. A **PR title** is the pure case, being out-of-diff by construction — and on a multi-commit PR a squash merge takes its commit subject from that title under GitHub’s default, so an overclaiming title can outlive the PR page it was raised on.

The remedy is mechanical, and it is a count rather than a judgment: before pushing, re-read the verdict body **and** re-fetch the thread list, then state how many findings the round raised and dispose of all of them in one push. Say explicitly which are deferred, per [issue-first](#).

- **Do:** state the round’s finding count before pushing, derived from both the body and the thread list.
- **Do:** read a title, a changelog line, and a PR body as reviewable surfaces — a finding about any of them can only arrive in the body.
- **Don’t:** read “every thread is resolved” as “every finding is handled”; the thread list cannot see an out-of-diff finding.
- **Don’t:** treat a correct, complete disposition of one finding as evidence about the round — that is a per-finding claim wearing a per-round shape.

(Measured twice within half an hour on 2026-08-21, in both available shapes. On [ai-config#1833](#) round 1 posted two inline findings; the first was fixed and pushed, and round 2 opened by re-raising the second — “the text at this location is essentially unchanged from what was flagged before” — at a cost of \$2.20. On [gha#550](#) round 1 posted three inline findings and a fourth in the verdict body only, about the PR title claiming work that had been deferred to another issue. All three threads were addressed, resolved, and pushed; the body-only one was missed. The second occurrence came after the first had already been written up, which is the argument for a count rather than for intending to look harder. The two are anchored by the re-raise at 17:12:39Z and by noticing the second miss at 17:41:36Z — derived from the PR timestamps rather than carried over from a figure quoted in a live comment, which is how “ninety minutes” reached the first draft of this entry.)

Accepting a finding does not verify the fix it appears to license. Everything above governs how a finding is *disposed of*. This governs the Address itself, which is the disposition nothing checks, because implementing a correct finding feels like the end of the question rather than the start of a new one.

A finding has two parts, and only the first carries the reviewer’s evidence. “Your summary hides variation” is an observation, and the reviewer measured it. “Therefore report the variation” is an inference, and nobody measured that. It presupposes the variation is signal, which is a separate claim about a separate thing — so an Address can be a faithful, careful implementation of a correct finding and still publish noise as structure.

The test is one question, asked before implementing any finding about variation, spread, or a range presented as a point value: **what would this quantity do if I changed something the claim is not about?** Vary a nuisance parameter — a second seed, a second run, a second sampling window — and see whether the structure survives. The cost is usually one line, because the harness that produced the original measurement is still open.

Note which direction this cuts. It is not an argument for resisting review, and the reviewer is usually right about what they measured. It is an argument that their verification does not transfer to your remedy, so the remedy needs its own.

- **Do:** separate a finding’s observation from its implied remedy, and check the remedy independently before pushing it.
- **Do:** vary a parameter the claim is not about, when the finding is that a summary conceals variation.
- **Don’t:** read “a summary hides variation” as establishing that the variation is real — ask which axis it lies on first.
- **Don’t:** treat the reviewer’s care in measuring as covering the inference drawn from the measurement.

(Measured across rounds 4 and 5 of an adversarial loop on [ucdavis/matt.contracts#2](https://github.com/ucdavis/matt.contracts#2), 2026-08-23. Round 4 correctly observed that “half-widths are optimistic by a factor of about 3.4” was a range presented as a point value: the measured ratios were 3.61, 3.41 and 3.30 at $N = 20, 40$ and 100. The summary was replaced with those per- N figures and with implied sample-size factors of 13, 12 and 11. Round 5 re-ran the same chunk at two further base seeds and got 3.441 / 3.398 / 3.363 and 3.441 / 3.405 / 3.332 — the $N = 20$ value moving by more than the whole apparent trend, and the headline 13 becoming 12. The ratio is $\sqrt{p(1-p)*DE / (p_r(1-p_r)*DE_r)}$, which contains no N at all, so the trend could not have been real. Reverted to the wording round 4 had displaced, which was stable at all three seeds. Tracked as [ai-config#2028](#).)

7 Reviewing AI-generated work

A reviewer of AI-generated work is not there to confirm that it sounds right. The job is to **try to invalidate it**. Assume the output is wrong until a check shows otherwise, and search mercilessly for mistakes.

Plausible, fluent prose is the main risk, not a comfort. It conceals errors that a human’s awkward draft would have made obvious. Put on the harshest critic hat, especially for:

- citations, DOIs, and URLs that may have been invented
- functions, flags, and APIs that may not exist on the version you actually run
- numbers, file paths, and “as of” claims that were never measured

Author-side validation ([Responsibility for validation](#)) is necessary and not sufficient. A second person (or a later pass by the author wearing a reviewer hat) should still try to break the work.

Automated review is a filter, not a substitute for this stance. Those tools miss domain errors and sometimes invent findings; they do not license a lighter human review.

8 Auto-fix PRs with Claude Code

[Claude Code’s /autofix-pr](#) watches the current branch’s pull request from Claude Code on the web and pushes fixes when CI fails or reviewers leave comments. It detects the open PR from your checked-out branch via `gh pr view`; to watch a different PR, check out its branch first. By default it fixes every CI failure and review comment; pass a prompt to scope it, for example `/autofix-pr only fix lint and type errors`. It requires the `gh` CLI and access to Claude Code on the web. A Marketplace action at [pr-autofix-with-claude-code](#) offers the same capability as a GitHub Action.

9 Using Copilot Review Before Human Review

Before requesting review from other humans, **always have Copilot review your pull request first**—even if Copilot created the PR itself. AI review provides fast, thorough feedback that helps catch issues before involving human reviewers, saving everyone time and improving code quality.

Why review with Copilot first:

- **AI has more bandwidth:** Copilot can review code immediately without competing priorities

- **Catch common issues early:** Copilot excels at identifying bugs, logic errors, security vulnerabilities, and style inconsistencies
- **Improve human review quality:** When humans review cleaner code, they can focus on higher-level concerns like design and architecture rather than basic issues
- **Learn from feedback:** Even experienced developers benefit from Copilot’s perspective on best practices and potential improvements
- **Growing capabilities:** AI review capabilities continue to improve over time, making this investment increasingly valuable

Copilot review workflow:

1. **Assign Copilot as a reviewer:** On your pull request page, assign Copilot to review the PR the same way you would assign any other reviewer. Click “Reviewers” in the right sidebar and select Copilot from the list.
2. **Review Copilot’s comments:** Once Copilot completes its review, carefully examine each comment. For each comment, decide whether you agree with the suggestion:
 - **If the comment is correct:** Address it by making code changes yourself or ask Copilot to apply the fix using GitHub’s suggestion features
 - **If the comment is incorrect or not applicable:** Dismiss the comment with an explanation for why it doesn’t apply
 - **If you’re uncertain:** Seek a second opinion from a human reviewer or do additional research
3. **Request another Copilot review:** After addressing or dismissing all comments, request another review from Copilot. This creates an iterative improvement process.
4. **Iterate until satisfied:** Repeat the review-and-address cycle until Copilot stops providing valuable suggestions. This typically takes 1-3 iterations depending on the complexity of the changes.
5. **Request human review:** Only after you’ve addressed Copilot’s feedback should you request review from human team members. At this point, the code should be in better shape, allowing human reviewers to focus on higher-level concerns.

Important considerations:

- **Copilot isn’t perfect:** AI review can produce false positives or miss important issues. Always apply your own judgment when evaluating Copilot’s suggestions.
- **Don’t blindly accept all suggestions:** Some of Copilot’s recommendations may not fit your specific context or requirements. It’s perfectly appropriate to dismiss comments that don’t apply.
- **Human review remains essential:** Copilot review supplements but does not replace human code review. Humans bring domain knowledge, understanding of business requirements, and judgment about trade-offs that AI cannot replicate.

- **Document dismissals:** When dismissing Copilot comments, briefly explain why. This helps human reviewers understand your reasoning and can serve as documentation for future reference.

For pull request authors:

Even if you're highly experienced, treating Copilot review as a required pre-review step helps maintain code quality and makes the best use of everyone's time. The few minutes spent on Copilot review often save hours of back-and-forth with human reviewers.

For human reviewers:

When you receive a PR for review, check whether the author has completed the Copilot review process. If Copilot hasn't reviewed the PR yet, consider asking the author to complete that step first before you invest time in review. This ensures you're reviewing code that has already been through initial automated quality checks.

10 Reviewing a Copilot PR You Didn't Create

When reviewing a pull request where someone else prompted Copilot to make changes, follow these guidelines to avoid confusion and ensure smooth collaboration:

Understanding PR roles:

The general PR roles (issue creator, author, reviewer, merger, assignee, and PR steward) are described in the [UCD-SERG Lab Manual's GitHub chapter](#). Copilot-assisted workflows add two more:

- **PR prompter:** Assigns a developer (human or AI) to start working on a PR, often by assigning an issue to Copilot. The PR prompter is sometimes the same person as the issue creator, but is often a project maintainer who reviews and triages external issue reports. When Copilot creates commits, both Copilot and the prompter are listed as co-authors. Typically the PR prompter becomes the PR manager.
- **PR manager:** Supervises and guides the PR authors, assigns reviewers, and controls the PR workflow, deciding when and how Copilot makes additional changes. The PR manager can hand off this role to someone else.

The scenario:

- One team member (the "PR prompter") assigned Copilot to work on an issue or explicitly prompted Copilot to start working
- The prompter may or may not be the same person who originally created the issue
 - In projects with a user base, users often submit issues (bug reports, feature requests)

- A project maintainer then steps in, adds their perspective, and assigns the issue to Copilot
- In this case, the maintainer who assigned Copilot is the prompter, not the original issue creator
- Copilot created the PR with the prompter as co-author
- The prompter (now acting as PR manager) requested your review
- Copilot may have also automatically reviewed the PR

As a non-manager reviewer, your role is to provide feedback, not to directly initiate more work by Copilot. The PR manager should remain in control of when and how Copilot makes additional changes.

Recommended review workflow:

1. Use “Comment” or “Request changes” based on the severity of issues:

- Use **“Comment”** for suggestions, questions, or minor issues that don’t block merging
- Use **“Request changes”** for significant issues that must be addressed before merging
- Both options allow you to provide feedback without directly triggering Copilot

2. Don’t ask Copilot to make changes directly:

- Avoid using features that would trigger Copilot to start working immediately
- Let the PR manager decide whether to ask Copilot to address your comments or make changes themselves

3. Write clear, actionable comments:

- Explain what needs to change and why
- Suggest specific solutions when appropriate
- The PR manager will decide how to address your feedback

For PR managers:

After receiving reviews from other team members:

1. Review all comments carefully:

- Decide which comments you agree with
- Dismiss or respond to comments you don’t entirely agree with
- This ensures Copilot only addresses feedback you’ve validated

2. Choose how to address valid feedback:

- **Option A:** Make the changes yourself (faster for simple fixes)
- **Option B:** Ask Copilot to address the feedback (better for complex changes)

- **Option C:** Add your own review summarizing which comments Copilot should address, then ask Copilot to respond to the open comment threads

3. Maintain clear communication:

- Let reviewers know how you plan to address their feedback
- Mark conversations as resolved after addressing them
- Request re-review from humans after Copilot makes significant changes
- Update the PR's "Assignees" field to reflect who is currently responsible for the PR

Transferring the PR manager role:

The original PR manager can hand over a PR to another person, who then becomes the new PR manager with control over Copilot's work on that PR. This might be useful when:

- The original PR manager is unavailable or on leave
- Someone with different expertise needs to guide the remaining work
- Responsibilities are being redistributed within the team

To transfer the PR manager role:

1. The original PR manager should clearly communicate the handover to all reviewers
2. The new PR manager should review the PR's history and any open feedback
3. The new PR manager should take over responding to Copilot and managing future iterations
4. The team should update the PR's "Assignees" field and comments or description to reflect the current PR manager

This workflow ensures the PR manager maintains control over the development process while benefiting from collaborative human review and Copilot's implementation capabilities.

11 Using a ChatGPT Account for Codex Pull-Request Reviews

OpenAI Codex can act as a reviewer on GitHub pull requests. The native integration uses the Codex service connected to a ChatGPT workspace and posts a standard GitHub review through the Codex connector bot. It does not require you to build a separate GitHub Action.

Prerequisites

The native reviewer depends on **Codex Cloud**. Before it can review a repository, you need:

- Codex Cloud enabled for your active ChatGPT workspace;
- the repository connected to Codex Cloud;
- access to the Codex code-review settings; and

- GitHub push or admin permission if you want to configure automatic reviews.

An `AGENTS.md` file is optional, but it lets the reviewer follow repository-specific guidance.

Requesting and automating reviews

After an administrator or repository owner connects the repository and enables code review in Codex settings, request a review by adding this comment to a pull request:

```
@codex review
```

You can add a one-off focus to the same comment, for example:

```
@codex review for security regressions and missing tests
```

Where the workspace configuration permits it, Codex can also review new pull requests automatically. In GitHub, the general code review reports only high-priority P0 and P1 findings to keep its comments focused.

A separate **Security Review** (research preview) is a deeper pass on security-specific risks. Request it with `@codex security review`. It can overlap with the general review's security findings.

When an administrator has disabled Codex Cloud

The message “**Your admin has turned off Codex Cloud**” is a workspace-policy restriction, not a GitHub repository error. The native `@codex review` bot cannot run because GitHub reviews execute through Codex Cloud. Local Codex use can continue. Run `/review` against a checked-out diff in any of:

- the Codex app
- the IDE composer
- the CLI

That is a local review. It does not post a GitHub review and does not create an always-on GitHub bot.

Resolve the restriction in one of these ways:

- A workspace administrator can enable Codex Cloud in the workspace's admin permissions.
- A user can connect a *personal* repository from a personal workspace that permits Codex Cloud. Do not use a personal workspace to connect organization-owned repositories — that bypasses the workspace policy your administrator set and moves the code outside your organization's controls.

- If cloud access must remain disabled, use the local `/review` path above.

GitHub organizations may separately require an owner to approve the repository connection. Changing that authorization does not override a ChatGPT workspace policy; both sides must permit the integration.

Native review versus API-backed automation

Do not confuse the native integration with a custom GitHub Action that calls an OpenAI model. The native reviewer is configured through Codex and the linked ChatGPT workspace. A custom action instead needs API credentials and uses API billing and limits. It also requires you to implement:

- the review prompt
- permissions
- comment-posting behavior

Do not copy personal ChatGPT or Codex login credentials into CI secrets.

Adding repository-specific review rules

Codex reads applicable `AGENTS.md` files. Put broad guidance in the repository root and narrower guidance in a file closer to the code it governs. Review-only guidance belongs under a `## Code Review Rules` heading:

```
## Code Review Rules

- Flag schema changes that are not backward compatible.
  A safe migration must support both deployed application versions.
- Flag behavior changes without a regression test.
```

Keep deterministic formatting and lint checks in continuous integration. Code-review guidance should:

- focus on consequential behavior
- state the safe alternative or exception
- remain concise enough to apply consistently

Codex review is an additional signal; it does not replace:

- tests
- branch protection
- required human approval

For current setup details, see OpenAI's [GitHub code-review documentation](#).

12 Where Pull-Request Review Lives in Claude Code Action

A common question about [anthropics/claude-code-action](#) is where its pull-request review lives. The answer is surprising: **there is no dedicated review action**. The repository publishes one general-purpose top-level action, and “review” is a *prompt* you pass it, not a separate artifact.

All claims below were surveyed against that repository at `main` (measured 2026-08-25; paths can move).

One action, many behaviors

The published action is [action.yml](#). It is a single composite action whose one external step installs Bun, and its description reads: “Flexible GitHub automation platform with Claude. Auto-detects mode based on event type: PR reviews, @claude mentions, or custom automation.”

What selects the behavior is the `prompt` input:

- A workflow triggered by comments containing @claude gets the interactive agent mode.
- A workflow that supplies an explicit review prompt gets a one-shot reviewer.
- The same action also handles issue triage and other automation, which is why one top-level action covers every behavior.

A second directory in the repository, [base-action/](#), is easy to mistake for a delegation target. It is not one: it holds a lower-level building block, developed in-tree and mirrored automatically to its own repository, [anthropics/claude-code-base-action](#). Consumers reference the mirror, not the in-repo path.

Where the repository reviews its own PRs

- [.github/workflows/claude-review.yml](#) is the reviewer. It triggers on `pull_request: opened`, skips fork PRs (they cannot mint the OpenID Connect (OIDC) token used for authentication), and calls `anthropics/claude-code-action@v1` with the prompt `/review-pr REPO: ... PR_NUMBER: ....`
- [.claude/commands/review-pr.md](#) defines what `/review-pr` does. It fans out to five reviewer subagents, defined under [.claude/agents/](#):
 - code quality
 - performance
 - test coverage
 - documentation accuracy
 - security

Each subagent is told to report only noteworthy feedback. The command then reviews that feedback and posts only the findings it also deems noteworthy — inline comments for specific issues, top-level comments for general observations or praise.

- [examples/pr-review-comprehensive.yml](#) plus two filtered variants (by author and by path) are copy-paste templates for adding review to another repository.

Adding review to your own repository

Copy one of the `examples/pr-review-*.yml` files into your repository's `.github/workflows/`, then provide credentials. The examples authenticate with an `ANTHROPIC_API_KEY` secret; the upstream repository's own workflows instead use Workload Identity Federation inputs (`anthropic_federation_rule_id`, `anthropic_organization_id`, `anthropic_service_account_id`) to exchange the workflow's OIDC token for a short-lived API token. Either route works; the federation route avoids a long-lived static key in CI.

The example's prompt is the customization point. It is ordinary prose naming focus areas, so you edit it the way you would edit any review checklist.

Reviewer versus @claude agent

Do not confuse the review workflow with `.github/workflows/claude.yml`, the interactive agent. They share one underlying action but differ in trigger surface:

	<code>claude-review.yml</code>	<code>claude.yml</code>
Fires on	PR opened	comments, reviews, or issues containing <code>@claude</code>
Behavior	one-shot review	open-ended agent session
Writes code	no	yes

This distinction explains a common debugging dead end: “@claude answered my comment, so why did nobody review the pull request?” Mention-triggered activity never runs the reviewer; only the `pull_request`-triggered workflow does.

Relation to Codex native review

This is the Claude-side counterpart to [using a ChatGPT account for Codex pull-request reviews](#). Codex ships a hosted native reviewer configured through the ChatGPT workspace, with no workflow file. Claude Code's reviewer is the opposite trade: you own a workflow file and supply API credentials, but the prompt, tools, model, and triggering events are all visible and editable in your repository.

13 Gemini Review Action for GitHub Pull Requests

[derailed-dash/gemini-review-action](#) is an open-source GitHub Action that provides automated code reviews on pull requests and automated triage on issues using Google's Gemini models (measured 2026-08-31; repository at v1.6.6).

Key capabilities

The action operates as a composite GitHub Action designed to run directly in CI pipelines. Notable features include:

- **Model selection and defaults:** Defaults to `gemini-3.7-flash`, configurable via the `gemini_model` input.
- **Dual authentication modes:** Supports standard API key authentication via the `gemini_api_key` input, as well as keyless Google Cloud Workload Identity Federation (WIF) by running `google-github-actions/auth` beforehand and setting `GOOGLE_GENAI_USE_VERTEXAI`, `GOOGLE_CLOUD_PROJECT`, and `GOOGLE_CLOUD_LOCATION` environment variables.
- **Hybrid codebase context enrichment:** Beyond analyzing git diffs, the action indexes repository context (structure, declarations, and key files) to assess pull requests against surrounding architecture.
- **Structured output and GitHub suggestions:** Generates line-level review comments formatted as GitHub suggestion blocks so contributors can apply suggested fixes directly in the GitHub UI.
- **Dynamic skill loading:** Can load domain-specific or project-specific instructions from `.agents/skills/` or custom directories, aligning reviews with existing repository skills.
- **Cost telemetry and context caching:** Emits token usage reports and estimated costs per run, leveraging Gemini's prompt and context caching to reduce API spend on repetitive large-context reviews.

Basic workflow configuration

To add Gemini code review to a repository, create `.github/workflows/gemini-review.yml`:

```
name: Gemini Code Review

on:
  pull_request:
    types: [opened, synchronize, reopened]

jobs:
  review:
```

```

runs-on: ubuntu-latest
permissions:
  contents: read
  pull-requests: write
  issues: write
steps:
  - name: Checkout repository
    uses: actions/checkout@v4
    with:
      fetch-depth: 0

  - name: Run Gemini Code Review
    uses: derailed-dash/gemini-review-action@v1
    with:
      gemini_api_key: ${ secrets.GEMINI_API_KEY }
      github_token: ${ secrets.GITHUB_TOKEN }
      gemini_model: gemini-3.7-flash

```

Comparison with native review integrations

Like [Claude Code Action](#), `derailed-dash/gemini-review-action` gives the repository owner full visibility over workflow triggers, authentication methods, and review prompts. In contrast to hosted review offerings that require platform-level permissions across an entire organization, this GitHub Action operates per-repository with credentials scoped to GitHub Actions secrets or Google Cloud IAM roles.

14 Auto-Review in Google Antigravity

Issue #109 asked whether Google Antigravity has an auto-review mode, prompted by a Reddit thread in the `r/google_antigravity` community. Reddit refuses requests from the machines this site is built on, so the thread itself could not be read (measured 2026-09-10); this section answers the question from the product's own documentation instead and says where the lab's practice fills the gap.

What the product offers

Antigravity has two things that look like review, and neither is an automatic review of your code.

- **Artifact review** is a gate on the agent’s own plans. The Artifact Review Policy setting, in the Agent tab of the Settings pane, offers “Request Review (Recommended)”, which pauses the agent and waits for you to approve its implementation plan or diff, and “Always Proceed”, which skips the pause and lets the agent run without manual verification (Google 2026a). Neither option makes a reviewer read the change: one asks a human to, and the other asks nobody to.
- **Review Changes** is a diff pane. Once the agent starts writing code, a Review Changes button in the Agent panel’s bottom toolbar opens a pane where you scroll through every change made in the conversation and leave comments on file diffs for the agent to act on (Google 2026b). The documentation describes it as a human reading the agent’s work, not the agent reviewing its own.

So the direct answer is no: as documented on 2026-09-10, Antigravity ships human-in-the-loop review of what the agent did, and an “always proceed” switch that removes even that. An automatic second opinion on a pull request has to be built.

Building one with the Antigravity SDK

Google’s own example is a GitHub Actions workflow that runs an Antigravity agent as an adversarial pull-request reviewer (Samborski 2026). It triggers on pull-request **opened** and **reopened** events, and on an `@agy /review` comment for manual runs; it uses the `run-agy-sdk` composite action with a managed agent whose system instructions cast it as a code auditor; it runs on the Actions runner itself, where the SDK can start Docker-based MCP servers for file access and PR commenting; and it deliberately limits the agent to reading files and posting comments, with reviews restricted to maintainers on non-fork pull requests. That is the same shape as the Claude Code Action review in Section 12 and the Gemini review action in Section 13: review is a CI job that a PR event starts, not a mode inside the editor.

What the lab does

The lab already built this once. `Morrison-Lab/ai-config` carries an `agy-review-workflow` skill and a reusable `antigravity-code-review.yml` workflow in `Morrison-Lab/gha` that posted a review, a security audit, and a set of generated tests on every pull request. That route ran on prepaid API credits, and it was retired on 2026-08-20 after a dispatched run ended with the credits depleted; the skill now documents itself as history and points cross-vendor second opinions at Codex instead (see Section 11).

The Antigravity command-line program, `agy`, is a separate lane and still works. Its headless form, `agy --print "<prompt>"`, answers a single prompt on standard output under the subscription’s own quota (measured 2026-09-09). Quote the prompt so the shell passes it as one argument, and put it directly after `--print:` on this machine, placing another flag between the two made `agy` treat that flag as the prompt. The lab uses it as a local adversarial reviewer before a push: pipe the branch diff into a prompt that asks for factual errors, unsupported

claims, broken cross-references, and style violations, and ends with a one-word verdict. Because `agy --print` cannot read files or the network in that mode, the diff has to be pasted into the prompt text itself. That is a pre-push check, not a PR bot, so the PR-side review stays with the Claude and Copilot reviewers the repositories already run.

Useful to us? Yes, as a local second opinion

- **Use `agy --print`** for a cheap, cross-vendor pre-push review when Claude quota is the scarce resource, which was the situation on 2026-09-09.
- **Do not** re-enable the API-credit workflow; it costs money the lab has stopped budgeting for, and a dispatch that fails leaves a red check no reviewer can clear.
- **Do not** read “Always Proceed” as an auto-review mode. It is the opposite: it removes the only review step the editor has.

15 Gemini Code Assist for Repository Code Review

Google Cloud provides a native code-review capability through [Gemini Code Assist](#) (measured 2026-08-31; documentation in Enterprise preview). Unlike GitHub Actions that run inside individual repository workflows, Gemini Code Assist operates as a managed service connected at the organization or repository level.

Architecture and setup

Gemini Code Assist connects to GitHub, GitHub Enterprise Cloud, or GitHub Enterprise Server through Google Cloud Developer Connect:

- **Connection host:** Repositories connect via Developer Connect connections provisioned in Google Cloud region `us-east1`.
- **Console enablement:** Code review is enabled through the Google Cloud Agents & Tools Console (under **Code review** in the Gemini Cloud Assist navigation).
- **Automated triggers:** Once connected, Gemini automatically generates pull-request summaries and posts line-level review comments when pull requests are opened or updated.
- **Interactive commands:** Reviewers and authors can trigger on-demand reviews or ask questions by posting a comment starting with `@gemini-code-assist` or `/gemini` directly in the pull-request thread.

Review scope and security boundaries

Gemini Code Assist enforces several deliberate boundaries on review scope:

- **Workflow file exclusion:** Files under `.github/workflows/` are explicitly excluded from reviews to prevent automated agents from analyzing or suggesting edits to CI credentials and pipeline definitions.
- **Style and policy customization:** Repositories can provide custom review guidelines and coding standards configured centrally in the Google Cloud Console, applying uniform review standards across multiple connected repositories.
- **Enterprise governance:** Because access is managed through Google Cloud IAM and Developer Connect, organization administrators can enable or disable reviews across projects without checking workflow files into individual repositories.

Trade-offs versus Action-based reviews

Dimension	Gemini Code Assist	Gemini Review Action / Claude Code Action
Hosting	Managed Google Cloud service	GitHub Actions runner
Configuration	Google Cloud Console & Developer Connect	In-repo <code>.github/workflows/*.yaml</code>
Workflow maintenance	Zero workflow files in repository	Repository owns workflow YAML and prompts
Interactive mode	Built-in <code>/gemini</code> comment tags	Configurable via <code>@claude</code> or Action triggers
Credential location	Google Cloud IAM & Developer Connect token	GitHub Secrets or Workload Identity Federation

16 PR-Agent: Open-Source Pull-Request Review Bot

[PR-Agent](#) (The PR Agent 2026h) is an open-source pull-request review bot. It runs one large-language-model call per command, posts the result as a pull-request comment or review, and works against several git forges with any model that LiteLLM can reach. All claims below were checked against the repository at `main`, its release list, and the documentation site at <https://docs.pr-agent.ai/> (measured 2026-09-09; the latest release was v0.45.0, published 2026-09-05).

History and ownership

PR-Agent began at CodiumAI, later renamed Qodo, and its hosted enterprise edition was sold as Qodo Merge. In 2026 Qodo donated the open-source project to the community. The repository now lives in the **The-PR-Agent** GitHub organization, has an external maintainer, and its README says it “is currently in the process of being donated to an open-source foundation” (The PR Agent 2026h). Qodo Merge was rebranded as Qodo 2.0, a separate hosted product.

The README is emphatic that the two are not the same thing: the repository “is not the Qodo offering for open-source projects”, and PR-Agent is “a community-maintained legacy project of Qodo” (The PR Agent 2026h). Qodo remains the project’s gold sponsor. Read the documentation with that split in mind: pages that describe a hosted app or a free tier for open-source projects are describing Qodo’s product, while PR-Agent itself is software you run.

The Docker images moved with the ownership. Releases 0.34.2 and later are published under **pragent/pr-agent**; the older **codiumai/pr-agent** namespace is a frozen archive (The PR Agent 2026h).

Commands

Every command is a slash comment on the pull request or a subcommand of the **pr-agent** command-line tool. The four core commands (The PR Agent 2026a):

- **/describe** generates a title, type, summary, code walk-through, and labels.
- **/review** generates a review covering possible issues, security concerns, tests, and an estimate of review effort. The review can also rate risk, recommend whether to merge, list the files a human should read first, and check the change against a linked GitHub or Jira ticket (The PR Agent 2026g).
- **/improve** posts actionable code suggestions, which the forge can render as one-click suggestion blocks.
- **/ask "..."** answers a free-text question about the whole PR or about specific lines.

Six further commands (**/add_docs**, **/generate_labels**, **/similar_issue**, **/help**, **/help_docs**, and **/update_changelog**) round out the set. **/help_docs** is disabled as of v0.36.1 pending a fix for a credential-exposure issue (The PR Agent 2026h).

Each command is a single model call, which the README describes as “~30 seconds, low cost” (The PR Agent 2026h). Large diffs are handled by a compression strategy rather than by a multi-step agent loop: deletion-only hunks are dropped, deleted files are collapsed into a list, and remaining files are added in order of size until the prompt approaches the model’s token limit (The PR Agent 2026f). Files that do not fit are named but not shown to the model. v0.45.0 adds an opt-in large-diff chunking mode for **/review** that splits a diff across several calls and merges the results.

Behavior is configured in a `.pr_agent.toml` file at the repository root, or by environment variables using dotted keys such as `pr_reviewer.extra_instructions` and `config.model` (The PR Agent 2026c). Each tool has an `extra_instructions` field for free-text review guidance, and `/review` has `require_*` switches (`require_security_review`, `require_tests_review`, and so on) that turn individual sections on or off (The PR Agent 2026g).

Supported forges and models

PR-Agent supports five forges:

- GitHub
- GitLab
- Bitbucket
- Azure DevOps
- Gitea

Most commands work on all five; the documentation keeps a per-forge feature matrix, and a few commands (such as `/similar_issue`) are GitHub-only (The PR Agent 2026a).

Models are bring-your-own. PR-Agent calls models through LiteLLM, so any provider LiteLLM supports can be selected by setting `config.model` (The PR Agent 2026d). Named providers include OpenAI, Azure OpenAI, Anthropic, Google (AI Studio and Vertex AI), Amazon Bedrock, DeepSeek, Mistral, Groq, xAI, OpenRouter, and local servers such as Ollama and vLLM. Further keys let you route work to different models: `config.model_weak` for lighter commands (`/describe`, `/ask`, `/update_changelog`), `config.fallback_models` for retries, and a `[model_routing]` table that picks a model by diff size (The PR Agent 2026d).

Local models work but come with a warning. The model-configuration page says that “most open-source models currently available (as of January 2025) face challenges with these complex tasks” and recommends them for “experimentation and learning purposes (mainly for the ask command)” (The PR Agent 2026d). That note predates the current generation of open-weight coding models, so treat it as a prompt to test rather than as a verdict; [small local models](#) covers what we have found local models can and cannot do.

Deployment

There is no hosted PR-Agent service. Every route runs the software on infrastructure you control (The PR Agent 2026e):

- **GitHub Action** (the README’s recommended route). A workflow that `uses: the-pr-agent/pr-agent@main` with `OPENAI_KEY` (or another provider’s key) and `GITHUB_TOKEN` in `env`, triggered on `pull_request` and `issue_comment` events. `github_action_config.auto_review`, `auto_describe`, and `auto_improve` choose

which commands run automatically when a PR opens, and the `issue_comment` trigger enables slash commands.

- **Self-hosted GitHub App.** A webhook server (the `pr-agent:github_app` Docker target) registered as a GitHub App with pull-request, issue-comment, and metadata permissions. The documentation covers running it as an AWS Lambda function and sizing the gunicorn workers. Equivalent webhook servers exist for GitLab and Gitea.
- **Command line.** `pip install pr-agent`, then `pr-agent --pr_url <url> review` runs a command against an existing PR from your own machine.
- **Docker.** Versioned images under `pragent/pr-agent` back all of the above.

The README’s data-privacy section says only that a self-hosted PR-Agent with your own OpenAI key “is between you and OpenAI” (The PR Agent 2026h). That is the whole privacy story: no third party sees the diff except the model provider you chose.

License and the hosted-versus-open split

The code is MIT-licensed, with the 2026 copyright held by “The PR Agent” (The PR Agent 2026b). The current documentation no longer marks individual commands as open-source or hosted-only: after the donation, everything at `docs.pr-agent.ai` describes the open-source tool, and the hosted features moved to Qodo’s own product and documentation. Older blog posts and third-party guides still refer to “PR-Agent Pro” or “Qodo Merge” feature tiers; those describe the product that became Qodo 2.0, not the repository reviewed here.

Useful to us?

The lab already runs four pull-request reviewers, so the question is what PR-Agent adds. The comparison uses the same axes as the other review sections:

Dimension	PR-Agent	Claude Code Action	Codex review	Gemini Review Action	Gemini Code Assist
Source	Open (MIT)	Open action, closed model	Closed	Open	Closed
Hosting	Actions runner or your own webhook server	Actions runner	OpenAI- hosted	Actions runner	Google- hosted
Model	Any via LiteLLM, including local	Claude only	OpenAI only	Gemini only	Gemini only

Dimension	PR-Agent	Claude Code Action	Codex review	Gemini Review Action	Gemini Code Assist
Forges	GitHub, GitLab, Bitbucket, Azure DevOps, Gitea	GitHub	GitHub	GitHub	GitHub
Billing	Your API key (or free with a local model)	Anthropic API key or federation	ChatGPT workspace	Gemini API key or Google Cloud	Google Cloud
Review depth	One model call per command over a compressed diff	Agentic; five reviewer subagents with repository access	Hosted agent	One call with repository context indexing	Managed service
Interactive	/ask, /review, and other slash comments	@claude mention (separate workflow)	@codex review	Triggered by workflow events	/gemini comments
Customization	.pr_agent.toml extra_instructions prompts in repo	Review prompt in the workflow file	AGENTS.md review rules	Skills directory	Google Cloud Console

What PR-Agent offers that none of the others do:

- **Model independence.** It is the only reviewer here that can be pointed at whichever provider currently has spare quota, or at an Ollama model for a zero-cost pass on a private repository. For the lab's quota juggling ([spend management](#) and [other models](#)), a reviewer whose model is a one-line config change is attractive.
- **Structured output.** /review reports effort, risk, and a merge recommendation in a fixed format, which a script can parse. The other reviewers produce prose.
- **GitLab and other forges.** Only PR-Agent covers repositories outside GitHub, which matters if a collaborator hosts code elsewhere.

What it lacks:

- **Depth.** A single call over a pruned diff cannot follow a symbol into a file the diff did not touch. The Claude Code Action reviewer runs as an agent with the repository checked

out, so it can. For the review standard the lab’s ARDI loop expects (see Section 12), PR-Agent is a first pass, not a replacement.

- **A stable maintainer story.** The project changed hands in 2026, its README says a foundation transfer is still in progress, and one command is disabled pending a security fix. Release cadence remains weekly (five releases between 2026-08-01 and 2026-09-05), so the project is active, but the governance is newer than the code.
- **Another bot in the thread.** Each added reviewer costs a review round per push and a comment the ARDI loop has to triage. Adding a fifth reviewer only pays off if it catches something the other four miss.

A reasonable trial is one repository, the GitHub Action route, `auto_review` on and `auto_describe` and `auto_improve` off, with `config.model` set to whichever provider has spare quota that month. If its `/review` findings overlap the Claude reviewer’s after a few weeks, turn it off; if it catches something distinct, keep it as the cheap first pass.

17 How a Session Learns a PR Changed

A coding-agent session that is watching a pull request does not poll it. Something wakes the session when the pull request changes, and in Claude Code that “something” is one of **two separate channels**.

The distinction matters for a practical reason: **the agent can turn one of them on and off, and cannot touch the other**. So “stop watching this PR” is a request an agent can satisfy for one channel and can only decline for the other.

Both arrive the same way — mid-turn, alongside the next tool result, the same delivery mechanism as a background task notification — so they are easy to mistake for each other. They are distinguishable by the tag wrapping them.

	<github-webhook-activity>	<ci-monitor-event>
Turned on by	the agent, via a tool call	a human, via a checkbox
Turned off by	the agent, via a tool call	a human, via the same checkbox
Available in	a session with the remote/hosted GitHub MCP server (not a local one)	Claude Code on the web only
Carries	comments, CI results, reviews, mergeability notices	the new comment, plus a fixed instruction template

The channel an agent controls

`subscribe_pr_activity` and `unsubscribe_pr_activity` are ordinary tool calls, taking an owner, a repository, and a pull-request number. Subscribing delivers that pull request’s activity into the conversation as `<github-webhook-activity>` messages until the session unsubscribes, or the pull request merges or closes.

Two caveats are worth knowing before relying on it.

A successful subscribe does not guarantee delivery. If a PR Steward agent already holds the watch on that pull request, the call still succeeds — but this session receives nothing. The tool result says so in as many words, so read the result rather than the exit status. Taking over the watch requires opting the steward out first, by removing its watching label on the pull request.

The tool does not exist on a locally-run GitHub MCP server. Workflow guidance written for remote or web sessions names it freely, which strands anyone following that guidance from a local harness. [MCP server setup](#) covers the local analogues to reach for instead.

Webhook delivery is also not exhaustive, which is the failure mode most likely to be mistaken for “nothing has happened”. CI *successes*, new pushes, and merge-conflict transitions can arrive late or not at all. A session that treats silence as “still green” will sit indefinitely on a pull request that has gone stale or conflicted, so a subscription is a supplement to periodically re-reading the pull request’s real state, not a replacement for it.

The channel only a human controls

Claude Code on the web shows a per-pull-request **CI monitoring** panel in the session sidebar, with two checkboxes: **Auto-fix CI & address comments** and **Auto-merge when ready**.

Ticking the first is what starts `<ci-monitor-event>` messages arriving.

Three properties of that panel surprise people:

- **There is no default.** No account-, organization-, repository-, or environment-level setting turns either checkbox on ahead of time. Absent the `/autofix-pr` shortcut described below, every new pull request and session starts with both off.
- **No agent-side tool can reach it.** It is client-UI state, not something an agent’s configuration surface touches, so asking an agent to enable it cannot work. If the checkbox changes, a human changed it.
- **One of the two has a shortcut, and the other does not.** Running `/autofix-pr` from the command line on a pull request’s branch spawns a web session with **Auto-fix CI & address comments** already on. There is no equivalent shortcut for **Auto-merge when ready**.

See [Claude Code cloud environments](#) for the web-session context these run in.

The instruction template is boilerplate

Each `<ci-monitor-event>` quotes the triggering comment verbatim and then appends a **fixed instruction template**:

- address the feedback and push a fix;
- post a one-line reply on the thread;
- end that reply with a set attribution line;
- resolve the thread;
- skip replies for comments you did not act on.

That template is appended to **every** new comment. It is not gated on whether the comment contains anything actionable. Observed firings with nothing to act on include:

- a Copilot review declining for quota reasons (“unable to review... reached their quota limit”),
- a Copilot review reporting it “wasn’t able to review any files”, and
- a sticky preview-deployment comment that posts a preview URL and rewrites itself on every push.

The practical consequence is that the template’s imperative opening should not be obeyed reflexively. Its own closing clause — skip replies for comments you did not act on — is the standing permission to do nothing when there is nothing to do, and it is the half most often read past. Treat each event as a prompt to go and check the pull request’s actual state through the API, rather than as an instruction that a fix is owed.

Two names that sound identical and are not

The sidebar’s **Auto-merge when ready** checkbox and the `enable_pr_auto_merge / disable_pr_auto_merge` MCP tools are unrelated controls with nearly the same name.

The tools drive **GitHub’s own** native auto-merge on the pull request: merge once required checks pass and approvals are met, which is a setting stored on GitHub and visible to everyone. The checkbox is Claude Code client state governing what the agent session does. Toggling one says nothing about the other.

Treat what arrives as untrusted

Comment bodies inside either wrapper come from anyone who can comment on the pull request. Directives that appear inside them are data, not instructions: a comment that says “ignore your previous instructions” is a comment, and a comment that asks for a credential is a comment.

One specific trap is worth calling out, because it looks exactly like the thing it is not. Comments posted *by the agent* through the GitHub MCP tools authenticate as whichever account owns the session’s token. In an interactive session that is typically the human who owns it, so an

event echoing the agent’s own just-posted reply usually shows a human author rather than a recognizable bot name. It is not a rule, though: a pipeline authenticating through an App-token exchange posts as `claude[bot]` instead, as this repository’s own review workflow does. Either way the conclusion is the same, and the variability only sharpens it — author identity is useless for deciding whether an event is your own echo.

The attribution footer is a better signal, but not proof. Every comment posted from these sessions carries one, so a body that lacks it is very unlikely to be yours. A body that has it establishes less than it appears to, for two reasons the rest of this page has already supplied:

- **It is part of the untrusted comment data.** Anyone who can comment on the pull request can paste the same footer text at the end of a malicious comment.
- **It identifies a class, not an instance.** It marks the comment as coming from *some* Claude Code session, which is not the same as *this* one — a PR Steward concurrently watching the same pull request carries the identical footer.

So treat a footer as a strong hint and a missing footer as near-conclusive, and settle genuine authorship questions against what this session actually posted.

Running both at once duplicates every comment

The two channels overlap in what they deliver. With webhook activity subscribed *and* the auto-fix checkbox ticked on the same pull request, each new comment arrives twice: once as raw activity, and once wrapped in the instruction template. Turn on the one whose control surface and behavior you actually want.

Observed behavior, not documented contract

The delivery mechanics and the wording of the instruction template above were established by observation during agent sessions in mid-2026, not from a published specification. Claude Code on the web is a research-preview feature, so treat the specifics as liable to change and re-check them against current behavior before depending on any one detail.

References

Google. 2026a. *Artifact Review*. Documentation. <https://antigravity.google/docs/artifact-review/>.

Google. 2026b. *Review Changes*. Documentation. <https://antigravity.google/docs/ide/review-changes-editor/>.

- Samborski, Remigiusz. 2026. *Building an Agentic PR Reviewer with Antigravity SDK*. Blog post. <https://dev.to/googleai/building-an-agentic-pr-reviewer-with-antigravity-sdk-3b0i>.
- The PR Agent. 2026a. *PR-Agent Documentation*. Documentation. <https://docs.pr-agent.ai/>.
- The PR Agent. 2026b. *PR-Agent License*. Documentation. <https://github.com/The-PR-Agent/pr-agent/blob/main/LICENSE>.
- The PR Agent. 2026c. *PR-Agent: Automations and Usage*. Documentation. https://docs.pr-agent.ai/usage-guide/automations_and_usage/.
- The PR Agent. 2026d. *PR-Agent: Changing a Model*. Documentation. https://docs.pr-agent.ai/usage-guide/changing_a_model/.
- The PR Agent. 2026e. *PR-Agent: GitHub Installation*. Documentation. <https://docs.pr-agent.ai/installation/github/>.
- The PR Agent. 2026f. *PR-Agent: PR Compression Strategy*. Documentation. https://docs.pr-agent.ai/core-abilities/compression_strategy/.
- The PR Agent. 2026g. *PR-Agent: Review Tool*. Documentation. <https://docs.pr-agent.ai/tools/review/>.
- The PR Agent. 2026h. *PR-Agent: The Original Open-Source PR Reviewer*. Software. <https://github.com/The-PR-Agent/pr-agent>.