

Pull-Request Workflow with Agents

Last modified: 2026-08-04 17:01:52 (PDT)

The practices below apply whether you are driving an agent’s work or doing the work yourself. They keep parallel sessions from colliding and keep a pull request moving toward a clean, mergeable state.

File an Issue Before Starting

When starting a **new** piece of work, go **issue-first**: before branching, editing, or opening a PR, make sure a tracking issue exists. Search the tracker first; if no open issue covers the task, **file one** (`gh issue create` / `glab issue create`), then proceed. Never jump straight into a PR without a tracking issue behind it.

The issue is the durable record of intent, scope, and “done” criteria — it gives reviewers context, lets the PR auto-close it via `Closes #N`, and keeps the work discoverable even if the PR stalls. Skip only when the task is already tracked by an open issue.

This rule settles *whether* something is tracked, not *where* it goes. An item whose deliverable is a decision rather than a diff belongs on the discussion board instead, per [choose-issue-or-discussion](#) — so read “file one” here as “file one in the right venue”, which for actionable work is the tracker.

When the issue is a **bug report**, include a minimal reproducible example (a `reprex` — <https://reprex.tidyverse.org/>) whenever you can. A `reprex` is what a maintainer needs to confirm and fix the bug, and it’s what they’ll ask for anyway, so providing it up front saves a round trip. The `reprexes` skill helps reduce the problem to a minimal, self-contained example.

When filing an issue that contains a list of independent subissues, file each subissue as a child issue linked under the parent (GitHub sub-issues feature: `mcp__github__sub_issue_write` in remote sessions, or `gh api` with the sub-issues endpoint in local sessions).

That splitting rule has teeth, and they are worth stating: a PR’s `Closes #N` closes the whole issue, including every item in it the PR never addressed. Read as tidiness,

the rule is easy to skip when the second item feels like a footnote. The actual consequence is that GitHub cannot partially close an issue, so the residual items are not deferred and not reopened — they are silently gone, and nothing in the merge, the PR, or the closed issue reports that anything was dropped.

It is worse than an ordinary lost to-do, because a closed issue is *evidence that the work was handled*. A later reader searching the tracker finds it closed and reasonably concludes every item in it was dealt with, so the loss is not merely silent but actively misleading.

So before writing **Closes #N**, re-read #N and confirm the diff covers all of it. When it doesn't, either split the remainder into its own issue first, or reference the parent with **Refs #N**, which links without closing.

- **Do:** split at filing time, or at the latest before the closing PR merges.
- **Do:** use **Refs #N** when a PR advances an issue without completing it.
- **Don't:** let **Closes #N** ride on an issue whose scope is wider than the diff.

(Morrison-Lab/ai-config#847, 2026-07-29: an issue was filed carrying a primary bug and a secondary note, and the PR fixing the first said **Closes #847**. The second item survived only because the maintainer asked about it before the merge, which is not a mechanism; it was split into #852 and shipped as #853, and both PRs merged within the following half hour. The splitting rule directly above already existed and was simply not applied when #847 was filed, which is the argument for stating its consequence rather than only its instruction.)

Claim a PR or Issue Before Working on It

Before starting a work session on a GitHub PR or issue — i.e. before fetching the branch, making edits, or invoking an automated review cycle — post a brief comment on the PR/issue so other people and any automated review bots know not to start a conflicting parallel session.

Use:

```
gh pr comment <N> --body "Working on this --- paws off until I'm done."
gh issue comment <N> --body "Working on this --- paws off until I'm done."
```

Then proceed with the work. After the session ends (PR merged, issue closed, or work otherwise paused), follow up with a closing comment so the PR/issue is unclaimed for the next person.

Skip the claim step if the most recent comment already says you are working on it. This applies to any task that will push commits to a PR branch or run iterative review loops. It does **not** apply to read-only inspection (showing a PR, checking status, explaining a diff) — those don't risk a parallel session.

This includes a PR **you opened yourself**: in repos with an active @claude agent (`claude.yml`), the agent can push commits to your branch on PR activity — e.g. merging `main` in — and collide with your in-flight push, so claim early to flag the branch as actively worked. (See `memories/claude-bot-workflows.md`, “(claude?) CI action”, for the collision-recovery steps.)

When starting work from an issue, follow the claim comment with an immediate draft PR — see [pr-on-claim](#) for the mechanics. An open PR is a stronger “in-flight” signal than a comment alone.

Verify a mid-task “already done” claim against real PR state before trusting or redoing it. A PR you claimed and are actively driving can still gain commits from a **second, independently-running session** under the same account — a `<github-webhook-activity>` review-comment-reply event can describe work (“Addressed... Pushed in `<sha>`”) that this session never did. Don’t assume it’s fabricated or injected, and don’t reflexively redo the same fix: cross-check the PR’s actual commit list (`gh pr view --json commits / pull_request_read get_commits`) and review threads before either (a) trusting the claim, or (b) starting the same fix yourself. If a commit with that SHA genuinely exists, authored close to when the event arrived, treat it as confirmation a live parallel session owns this PR right now — stop pushing further speculative fixes yourself, and, if genuinely in doubt, ask whether to keep driving or step back, rather than racing the other session’s pushes. This gap is distinct from the initial claim check above: it’s not about claiming a PR before starting, but about **re-verifying you’re still the sole active driver** once work has been under way for a while — especially when you picked up the PR mid-session (e.g. by answering a diagnostic question about it) rather than through the normal claim-then-branch flow, so no fresh “paws off” check ever ran right before you started pushing. (d-morrison/gha#286, 2026-07-24: a webhook event delivered a review-comment reply attributed to d-morrison reading exactly like a Claude-authored reply, claiming a fix “Addressed... Pushed in 3fb8c5b” that this session hadn’t made; verified real via `get_commits` before proceeding — a second live session, not injection.)

The git-level variant of that check: a rejected push whose remote commit is byte-for-byte what you were about to push. The section above covers a *comment* claiming work was done. Here the parallel session makes no claim at all. Your `git push` is simply rejected because it pushed first, and what it pushed is the same merge you just made. The reflex on a rejected push is to merge again, which would stack a redundant merge commit on top of an identical one.

Four reads settle it before you touch anything:

```
git rev-parse HEAD^{tree}          # your merge's tree
git rev-parse origin/<branch>^{tree} # theirs
git show -s --format=%P HEAD        # your merge's parents
git show -s --format=%P origin/<branch> # its parents
```

An identical tree plus identical parents means the two merges are the same merge, so the right action is `git reset --hard origin/<branch>`.

- **Do:** compare trees and parents before deciding what a rejected push means.
- **Do:** discard your local merge with `git reset --hard origin/<branch>` once both match.
- **Don't:** re-merge reflexively on a rejected push — that is what produces the redundant merge commit.
- **Don't:** force-push over the other session's commit.

(Morrison-Lab/ai-config#965, 2026-07-31: main moved one commit, a local `git merge origin/main` was made, and the push was rejected. The remote carried `b8d2273`, a merge of the same two parents, with tree `1bda1bc`, identical to the local merge's.)

Handing off mid-task to another agent, on user request (“finish what you’re doing, then relinquish holds; I’ll put another agent on them”): don’t just stop — leave the next agent a clean starting point. On each claimed PR/issue: (1) post a status comment on the PR itself distinguishing what’s **done** from what’s genuinely **not done** (the actual point of the issue, not just the side-fixes found along the way) and any blocker still open, so the next agent doesn’t have to re-derive it from the diff; (2) post the closing/unclaim comment on the issue per the pattern above; (3) `unsubscribe_pr_activity` (or stop babysitting locally) so you don’t keep auto-fixing a PR you no longer own; (4) stop any background watch/poll task tied to that work (e.g. a `ScheduleWakeup` or a `Monitor/background-Bash wait`) so it doesn’t fire into a session that’s moved on. A merge-conflict-free `git status` and a pushed branch are not enough on their own — the status comment is what makes the handoff legible. (ucdavis/bcs gia session, 2026-07-06: handed off PRs #310 and #311 mid-implementation this way, each blocked on the same slow `renv::restore()`.)

Keep Your Branch Synced with Main

Whenever `main` has moved ahead of a PR branch you’re working on, **merge main into the PR branch** before the next push or review trigger. Don’t wait for a conflict to surface or for someone to ask.

This fragment covers the single-branch-vs-main case. When orchestrating a multi-agent `ultracode` session, merges can happen at more points than that — see [ultracode-merge-conflicts](#) for the broader check (worktree-isolated agent branches, concurrent `parallel()` results) and the note on GitHub’s mergeable indicator not evaluating custom `.gitattributes` merge drivers.

Always check for merge conflicts with main before pushing results to remote. Run this before every push, not just before triggering a review:

```
git fetch origin main
git log --oneline ..origin/main | head    # any commits? main is ahead --- merge it in
git merge origin/main
```

If the push is rejected because `main` has moved (! [rejected] with (fetch first) or (non-fast-forward)), fetch and merge before retrying — don't force-push.

Always do this before triggering a fresh review too, so the reviewer evaluates the PR against current `main` rather than a stale snapshot.

Don't rebase or squash-rewrite a published PR branch unless explicitly asked — a merge commit is the right move because it matches GitHub's "Update branch" button and preserves the PR history.

If the merge has conflicts, resolve them, run the project's standard pre-commit checks (render / lint / spell / tests), commit, then push. Don't push a half-resolved merge.

After merging main, re-check version parity. In R packages with a `version-check` CI job, the branch's `DESCRIPTION Version:` must *exceed* main's. A conflict-free merge can silently put them at parity — main advanced (e.g. another PR merged between when you last bumped and now). After every merge of main, compare versions:

```
git fetch origin main
git show origin/main:DESCRIPTION | grep ^Version
grep ^Version DESCRIPTION
```

If they match, bump the branch's `Version:` by one patch level before pushing.

Re-check main again right before the final push, not just at the start of a merge. Resolving a conflict (rerunning generators, fixing prose, updating a `CHANGELOG` entry) can take long enough for `main` to advance a second time. A `git fetch origin main` immediately before `git push` — after conflict resolution is done, not only before it started — catches that case; an earlier CI failure on a commit you thought was current is a symptom of skipping this second check.

A conflict-free merge does not mean derived artifacts are in sync. If your branch regenerates a generated tree (e.g. `codex-skills/`, a lockfile, rendered docs) and `main` added a new *source* input the generator consumes (a new skill, a new dependency), git merges both cleanly — but the generator never ran against the new input on your branch, so its output is missing or stale and the sync check fails on `main` after both land. After merging `main`, re-run the generator and commit the result whenever main touched the generator's inputs — don't trust the absence of conflicts. (Concretely: merge the PR that adds the new skill *first*, then sync the wrapper-regenerating branch and rerun `scripts/sync-codex-skill-wrappers.py` before merging it.)

A CI failure on a brand-new PR's very first commit (e.g. the empty claim-commit from pr-on-claim) is a signal to check main's position before debugging the failure itself. A local checkout that sat around since before the session started can already be many commits behind `main` — the failure (a stale generated-tree check, a check `main` has since added or dropped) often isn't a real problem with your change at all, just `main` having moved. `git fetch origin main && git log --oneline ..origin/main` first; if `main` is ahead, merge it in and re-run the checks before treating the failure as something to fix in the diff. (`stack-prs` #359: an empty-commit draft PR failed `validate` on a stale `codex-skills/` generated tree, and the `require-changelog` job on a newly-added `CHANGELOG.md` requirement from PR #354 — both were `main` having advanced past a checkout that predated the session, not a defect in the new skill.)

The same staleness trap has a silent variant with no CI failure to flag it: a worktree/branch named after a PR's followup can still be based on a main from before that PR actually merged. A worktree directory or branch name suggesting “after PR #N” (e.g. `pr-N-followup-...`) is not proof the branch's actual base commit postdates #N's merge — it can have been created earlier and simply named for its intended purpose. Trusting that naming, then reasoning from `git show <hash>` for a commit found via `git log --all` (which lists every reachable commit across all refs, not just your branch's ancestry) can make content look present when it isn't actually in your branch yet. Verify with `git log --oneline HEAD..origin/main` or by reading the actual blob your branch would produce (`git show HEAD:<path>`, or the working tree itself before assuming what it contains), not a commit hash pulled from `--all`. If `main` has moved, merge it in before building further edits on the assumption the missing content exists. (`ai-config`#637: a worktree named `pr-636-followup-...` was cut from a `main` snapshot that predated #636's own merge; an edit referencing “the bullet above” — added by #636 — was written and committed before the bullet actually existed on the branch, caught only when `git push` reported `main` has moved and the subsequent merge produced a real conflict.)

A real conflict inside a file whose logic is also copied elsewhere (an extracted script, a doc example) needs the copy re-synced too, not just the conflicted file resolved. When a PR extracts inline logic (e.g. a workflow step's shell block) into a standalone script for testability, and `main` independently changes that same inline logic while the PR is open, resolving the merge conflict in the workflow file is not enough — the extracted script must be updated to match `main`'s new logic exactly, or the PR silently reverts `main`'s fix the moment it merges. Diff the extracted copy against `main`'s current inline version line-for-line (strip indentation, `diff`) to confirm an exact match, not just “looks about right.” If the PR carries tests against the extracted copy (fixtures, unit tests), add regression coverage for whatever `main`'s change fixed — the merge is the natural moment to catch a gap the original PR's tests didn't anticipate, and to prove the new fixtures actually catch the regression (temporarily revert the fix, confirm the test fails, then restore). (`gha`#176: `main` landed #173's lenient verdict-matching fix to `claude-code-review.yml`'s inline fail-check logic while a PR extracting that same logic to `scripts/check-review-execution.sh` was still open; the conflict resolution

updated the script to match verbatim and added two new fixtures for #173's specific fix, verified to fail against the pre-fix logic.)

When `main` DELETES a file your branch references, resolving the marked conflict is not enough — grep the whole tree for the deleted path. Git only conflicts where both sides edited the same lines, so a merge that brings in a deletion flags the file that *used* the thing, and nothing else. Any other reference to the deleted path — a docstring citing it as precedent, a comment, a doc cross-reference — merges cleanly and silently becomes a dangling reference, because those files were never in the conflict's scope. After resolving any merge that removed a file, run `grep -rn "<deleted-path>"` across the repo and re-point or reword each hit; then distinguish live references (must be fixed) from historical citations of the removal itself (correct as-is, leave them). This is the deletion counterpart to the extracted-copy case above: there the logic moved and a copy went stale, here it vanished and the pointers went dead. (ai-config#696: `main` retired `scripts/check-new-line-breaks.py` via #703 while the PR was open. The `validate.yml` conflict was visible and resolved, but `scripts/check-memory-file-size.py`'s docstring cited the deleted script as its advisory-exit-code precedent — a file the conflict never touched, caught only by grepping for the path afterward.)

A textual conflict in a skill file can be the symptom of a conceptual duplicate, not just competing edits to the same line. When merging `main` into a branch that's authoring a new skill, if the conflict lands in a `## Relationship to other skills` section (or `main` added an entirely new skill in the same territory), that's a signal to re-run `skill-builder`'s Step 0 judgment — not just resolve the diff mechanically. Compare the new skill against whatever landed on `main`: are they the same concern (fold into one, redirect), or genuinely distinct (cross-link both directions so neither reads as an unexplained near-duplicate)? `skill-builder`'s in-flight-work scan only runs once, at the start; `main` can grow a colliding skill in the time a PR is open, so the check has to be repeated at merge time too. (PR #352's `check-info-quality` landed alongside #344's independently-authored `fact-check-prose` this way — distinct enough to keep both, resolved by adding an explicit boundary in each skill's Relationship section rather than consolidating.)

The same collision can land before you write a line, and then it produces no conflict at all — just duplicated work nobody flags. The bullet above catches a duplicate at merge time, via a conflict. When `main` gains the colliding content while your change is still *planned* rather than written, there is nothing to conflict with: you write the duplicate, push it, and the review has to argue you out of content that was already redundant on arrival. So re-run the dupe check after any fetch that brings in new commits, not only at merge — a plan researched an hour ago was researched against a different `main`.

The cheap version is to read what actually arrived rather than only the count: `git log --oneline <old>..origin/main` plus `git diff --stat` over the same range, then ask whether any of it covers something still on your list. In a session that loads skills or plugins from the repo, a new one appearing in the session's own skill listing is the same signal arriving for free.

This is a *timing* gap, and it composes with the *scope* gap rather than replacing it. [check-open-prs-before-duplicating](#) covers work that is still in flight, unmerged, and therefore invisible to any check against `main`; run that one too, since a duplicate is just as wasted whether the collision has landed yet or not. Both checks share the same weakness — each runs once, at the start, and answers for the moment it ran.

Dropping the planned work is the cheap outcome, so record why in the issue and the PR body rather than deleting it silently — otherwise the next person re-proposes it. (ai-config#774, 2026-07-28: a planned `profile-before-optimising` fragment was mooted by `skills/measure-performance`, which merged via #762 during the session and covered the same two chapters — including the specific gap the fragment was meant to fill. It surfaced only because the new skill appeared in the session’s skill list after a routine fast-forward; the plan had been written before it existed. Dropped before implementation, with the reasoning recorded in both the issue and the PR body, and the neighbouring fragments cross-linked to the skill instead.)

A routine merge from `main` can create the duplicate inside your own diff. The collision above lands before you write, so the duplicate is redundant on arrival. A later `main` merge is quieter: both branches were non-duplicative when they were written, and the duplicate appears only when you bring the other branch’s text into yours. Git reports a clean merge because the two copies sit in different files. Diff-scoped added-line checks do not help either, because the duplicated lines already existed on one side or the other. So after merging `main` into a prose branch, run the duplicate check against the branch’s full current diff and the neighbouring corpus, not only against lines added by the merge commit.

- **Do:** after a `main` merge, re-run a cross-file duplication check over the merged branch’s whole prose diff.
- **Do:** treat a reviewer finding on such duplication as correct even when each copy was independently right before the merge.
- **Don’t:** assume a conflict-free `main` merge preserved DRY, or that the duplicate would have appeared in an added-lines-only scan.
- **Don’t:** answer by asking which branch “introduced” the duplication; the merge introduced the state that made both copies coexist.

(Morrison-Lab/ai-config#969, 2026-08-01: #969 added `shared/workflow/batch-merge-and-resolve.md` with a blockquote generalizing that an added-lines-only instrument is unsound when a defect can be introduced by deleting a line. PR #966 independently added the same generalization to `shared/workflow/sync-with-main.md` and merged after #969’s branch was written. `git show 50afe818:shared/workflow/sync-with-main.md`, normalized for whitespace and markup, did not contain the phrase, so the duplication did not exist at #969’s pre-merge head. The round-2 merge from `main` brought #966’s copy in, and the round-3 review correctly flagged the two uncited copies.)

Two PRs that each append a new terminal numbered subsection to the same file (e.g. `### 5. ...` in a `CLAUDE.md` review-guidelines list) will conflict on merge even when neither side’s content actually disagrees. This isn’t an editorial clash — it’s two

authors both writing to “the next number” at the same insertion point. Resolve by keeping **both** additions and renumbering sequentially from the collision point on, not by dropping either side; then grep the file for any other place that names the old numbering (a cross-reference, an index). This is also a reason `fully-clean`’s CI-green-and-review-clean verdict is a snapshot, not a mergeability guarantee — `main` can pick up its own append in the same spot after your last review round, so a PR can go from “reviewed clean” to “needs a merge conflict resolved” with no defect in its own diff. Before reporting a PR ready to merge, re-check with `git fetch origin main` plus the `git merge-tree` command from `resolve-conflicts`, not just a cached mergeable flag or an earlier green CI run. (gha#211: `main` merged #209’s own new ### 5. Check for AI-generated prose tells subsection between this PR’s clean review and its actual merge — `git merge-tree` surfaced a real conflict that neither PR’s own CI nor review status had flagged, since neither had rerun since `main` advanced.)

After merging a PR that extracts an inline block into a reusable unit (a composite action, a shared script/function), check other open PRs that still edit that same inline block — your merge just broke their textual diff, even though their intended change is usually trivial to re-apply to the new location. This is the mirror image of the case above: there, you’re the one resyncing after `main` moved a copy of your logic; here, *you* are the one who moved the logic, so the burden of noticing and fixing the resulting conflict falls on you, not on the sibling PR’s author waiting to hit it. Don’t wait for that PR’s own merge/CI to surface the conflict — check every open PR touching the same file right after your extraction merges: `git merge-tree "$(git merge-base origin/main origin/<sibling-branch>)" origin/main origin/<sibling-branch>` (or `gh pr diff <N>` against the new `main`) shows whether it still applies cleanly. Re-apply the sibling PR’s actual semantic change (not a mechanical `--theirs`) to the new location, verify with a direct diff that the extracted unit now differs from `main` by exactly that PR’s intended change and nothing else, then push to their branch and flag what you did in a PR comment. (gha#201 extracted `claude-code-review.yml`’s `claude_args` block into a new `run-claude-review-attempt` composite action to support a retry; gha#202, open in parallel, edited that same inline block to allowlist `WebFetch/ Bash(curl:*)`. Proactively rebasing #202 and re-applying its allowlist change to the new composite action — rather than leaving its author to discover a conflict — let it merge within the hour instead of stalling.)

An add/add conflict on a *shared config file* usually means two PRs independently fixed the same root cause — reconcile the reasoning, don’t just pick a side. This generalizes the skill-file case above beyond skills: a repo-wide CI/lint/build config fix (a new tool config file, a workflow tweak) is exactly the kind of change multiple sessions or bots are likely to attempt in parallel once a check starts failing on `main` for everyone. When the conflict is a whole-file add/add (not just competing edits to an existing file), read both sides’ reasoning — code comments, commit messages, the PR discussion — before resolving; usually one side’s explanation is more complete (covers a case the other missed, cites the tool’s actual constraint) and should win outright rather than mechanically merging fragments of both. Re-diff the PR against `origin/main` after resolving to confirm the PR’s remaining changes are its own original scope, not a reintroduction of what the other, now-merged PR already

added. (d-morrison/altdoc#7 vs #18: both independently added a `jarl.toml` excluding the same fixture directory for the same `jarl-check` failure; #18 merged first, #7's merge conflicted on the new file, resolved by keeping #18's more detailed comment and re-confirming #7's diff against `main` was back down to just its own four files. This same "append-collision" pattern struck a third time one insertion point over: this bullet and the two above it were each added by independent PRs landing in quick succession, all appending after the same "PR #352's `check-info-quality...`" paragraph — resolved, per the guidance above, by keeping all three rather than picking one.)

A dirty mergeable_state on a bot-opened PR can mean a sibling PR already closed the same issue, not just that main drifted. An issue-triggered @claude workflow can fire twice on the same issue in quick succession (a duplicate dispatch, or two people independently routing the same request), producing two independent PRs that both fully resolve it — including adding the identical new file. The second PR's merge conflict is an add/add on that new file, and it looks like ordinary main-drift, but treating it that way and mechanically resolving in favor of "ours" silently reintroduces a duplicate the other PR's merge already published. Before resolving, check the PR's linked issue for **other** cross-referenced PRs/closing events — if one already merged and closed it, diff the conflicting file against `main`: if it's the sibling PR's already-published version, keep `main`'s content and keep only this PR's genuinely distinct remainder (a piece the sibling PR never did), rather than re-adding a second copy. (ai-config#501: issue #500 was independently resolved twice — #502 merged first, adding `shared/writing/math-derivation-steps.md` and closing #500, but never wiring it into `CLAUDE.md`; #501 added a second copy of the same fragment plus the missing `CLAUDE.md` wiring. Resolved by keeping `main`'s published fragment and #501's wiring, turning a dirty merge into a clean +8/-0 diff.)

The same parallel resolution can be a whole-file split, and then files can vanish from your diff with no deletion hunk to read. The add/add and duplicate-issue cases above both say to keep `main` when a sibling PR already published the same new file. The split case adds a second check, because resolving the one conflict can also make other files disappear from your PR's diff entirely. Those files look harmlessly gone, and there is no deleted line for `ardi`'s pre-push deletion sweep to inspect. Two causes are indistinguishable from the final diff alone: `main` absorbed your cross-reference edit, or the merge dropped your work. So verify each vanished file against the pre-merge head before calling the collapse correct. For each file that left the diff, compare the original head against the merge-base to recover what your branch intended, then confirm current `main` now carries that same change. Only after that per-file check is it safe to treat the smaller diff as a successful conflict resolution rather than as lost work.

- **Do:** save or read the original pre-merge head, list the files that left the PR diff after the merge, and verify each one's intended change is already on `main`.
- **Do:** keep `main`'s version for the overlapping split file when the sibling PR has already published the same refactor, then carry forward only this PR's distinct remainder.

- **Don't:** infer that a vanished file was safely absorbed merely because the final diff got smaller.
- **Don't:** rely on the deleted-lines sweep for this case; content that left the diff has no deletion hunk for that sweep to show.

(Morrison-Lab/ai-config#966, merging `origin/main` after #973 landed as `ea11bc9a`, hit `CONFLICT (add/add)` on `memories/github-mcp-tools.md`. Both branches had split that file out of `memories/github.md`; the two split versions were byte-identical apart from #966's own 49-line addition, so keeping `origin/main:memories/github-mcp-tools.md` was correct. The final PR diff collapsed from 13 files, 1093 insertions, and 661 deletions to 8 files, 429 insertions, and 6 deletions. Five files disappeared entirely: `memories/claude-bot-workflows.md`, `memories/claude-code.md`, `shared/workflow/efficient-pr-babysitting.md`, `shared/workflow/fully-clean.md`, and `skills/purge-hallucinations/SKILL.md`. Each had contained only a cross-reference repointing from `memories/github.md` to `memories/github-mcp-tools.md`, and `git show origin/main:<file> | grep -c github-mcp-tools` returned 1 for each.)

A merge into a growing numbered list (e.g. gha's `CLAUDE.md` “Code review guidelines” section) can produce zero blank lines between two adjacent headings even with no textual conflict — `lint` catches it, `git` doesn't. When a section is a hotspot several PRs independently append items to (each PR adding its own `### N.` block at the end), a clean three-way merge can still splice one PR's closing line directly against the next PR's heading with no blank line between them — this doesn't produce a `<<<<<<<` conflict marker (`git` resolves it as a straightforward insertion), so it's easy to push without noticing. `markdownlint`'s MD022 (blanks-around-headings) is what actually catches it, as a CI failure with no proximate code change to explain it. Re-run the repo's markdown lint (or at minimum re-read the diff around every `### N.` boundary you didn't personally write) after any merge that touches a shared growing list, not just after a merge with conflicts. (gha#208: an out-of-band merge from `main` — done by a different session, not the one that opened the PR — landed a new item 7 directly against the PR's own item 6 with no blank line; `lint-markdown`'s MD022 failed with no conflict marker anywhere in the diff to point at.)

The same splice happens to `LIST ITEMS`, and there `markdownlint` most likely does **NOT** catch it — so nothing turns red at all. The case above is a heading spliced against preceding text, which MD022 decides. The changelog case is a *bullet* spliced onto the previous item's continuation line:

```
`data-raw/precompute-true-effects-chunk.R` (#429).
* The `docs` workflow's "Build site" step no longer times out intermittently.
```

That is a valid **tight** list item, so a list in which every other entry is blank-line separated silently starts mixing tight and loose items and renders inconsistently. `markdownlint`'s blanks-around-lists rule governs the boundaries of a list, not the gaps *between* its items, and no default

rule enforces consistent looseness within one list — so unlike the heading case, CI stays green and only a human reading the merged section notices.

Check it mechanically instead of by eye; one line decides it:

```
awk 'prev !~ /^[[:space:]]*$/ && /^[*+\\-] / {print FILENAME":"NR": "$0" {prev=$0}}' NEWS.md
```

Use `[[[:space:]]*` rather than a bare `/^$/` — a whitespace-only preceding line is not a violation and produces false positives. The pattern `[*+\\-]` covers all three common Markdown unordered-list markers; `^\\*` alone would miss `-` and `+` bullets.

Two consequences. Run this after any merge into a changelog or other growing bulleted list, alongside the heading check above. And note that a `merge=union` driver on such a file (see [configure-gitattributes](#)) *increases* the rate of this defect, since union resolves an append collision by keeping both sides with no conflict to review — so confirm a detector is wired into CI before enabling one, not after. (ucdavis/bcs#422/#430, 2026-07-26: a clean three-way merge spliced one PR's `## Bug fixes` bullet against another's; the check above then found **four** pre-existing instances in the same `NEWS.md`, in a repo that had no Markdown linting at all.)

Run that check as a whole-file count, and compare it before and after — scoping it to the lines you added cannot see this defect at all. The check above is the right instrument; the natural way to apply it is the wrong one. Having found the file's spliced bullets, the obvious next question is which of them are yours, and the obvious way to answer is to intersect them with the lines the branch added. That question is unanswerable, because the defect is a **deleted blank line** before a bullet that was already there. The bullet is *context* in the diff, never an addition, so the intersection is empty by construction and the check reports a confident zero.

Note how this differs from the scope failures elsewhere in this corpus, where a check's **inputs** were too narrow — a glob, a missing flag, a two-dot range. Here the inputs were right and the **question** was wrong, which no widening fixes. And it fails in the direction that reads as an all-clear, on the one file a reviewer will not re-derive.

The sound form is a count delta over the whole file, which needs no judgment about ownership and no diff at all:

```
git show origin/main:NEWS.md | awk '...' | wc -l # before
awk '...' NEWS.md           | wc -l             # after
```

A merge must not increase the count. That is an [algorithmatize-checks](#) instrument in the strict sense — two integers decide it — and it holds whoever authored the surrounding lines.

Generalize past changelogs, because the property is about the defect rather than the file: **when a defect can be introduced by deleting a line, any instrument keyed on added lines**

is unsound. Ask instead whether a whole-file measurement got worse. The version-parity rule above is the same shape — a conflict-free merge leaves the branch at parity with `main`, `version-check` goes red, and there is nothing in the diff to point at — which is why that rule is a direct comparison of two `DESCRIPTION` versions rather than a diff inspection.

The shared trigger is the practical part: **a conflict-free merge is exactly when nothing prompts anyone to look.** Both defects arrive through one, both are invisible to diff-scoped checking, and the merge reports success.

- **Do:** measure the whole file before and after a merge, and treat any increase as the merge’s fault regardless of who wrote the lines.
- **Do:** ask, of every diff-scoped check, whether the defect it targets could be caused by a deletion — and replace it with a count if so.
- **Don’t:** intersect a whole-file finding with the branch’s added lines to decide ownership; for a deletion-caused defect that always returns zero.
- **Don’t:** read a conflict-free merge as a merge that changed nothing beyond what the diff shows.

(ucdavis/bcs#534, 2026-07-31: merging `main` spliced its `NEWS.md` bullet directly beneath the branch’s own with no blank line between. markdownlint stayed green, since `blanks-around-lists` governs a list’s boundaries rather than the gaps between its items. A pre-push check asking which flagged bullets were among the branch’s added lines returned 0 — `git diff -U0 origin/main...1b74899d -- NEWS.md | grep -c '^+\\* Say that the'` — and that zero was reported to the user as “none of these are mine”, which was false. The count delta settles it: 14 spliced bullets on `origin/main`, 15 at `1b74899d` after the merge, 14 again at `9f5dab34` after the fix. Evidence filed on ucdavis/bcs#437; the `merge=union` driver proposed in ucdavis/bcs#438 would raise this defect’s rate, so it is explicitly blocked on the detector landing first.)

A commit claiming “I’ve pulled main and resolved the merge conflicts” can be lying — verify it actually merged before trusting the claim. A genuine conflict-resolution commit is a merge commit (two parents); a commit that just hand-edits files to *look* resolved, without running a real `git merge`, is an ordinary single-parent commit — and it never actually incorporates whatever new state of `main` prompted the “resolve conflicts” request in the first place. This is easy to miss because GitHub’s own `mergeable/mergeStateStatus` fields don’t distinguish the two: both look identical from the PR page until you check the commit graph. Verify with `git show -s --format="%P" <commit>` — one hash means no real merge happened, regardless of what the commit message says. If a branch needed conflict resolution more than once and each attempt claimed success but the PR still shows `CONFLICTING`, check every “resolved conflicts” commit in its history this way before trying yet another resolution attempt on top of a foundation that was never actually re-merged. (Lacaedemon/sparta#1070, 2026-07-27: an automated PR-authoring agent pushed two consecutive commits both titled “Resolve merge conflicts”, each claiming to have pulled `main`; both were single-parent commits that never touched `main`’s actual current state, so the PR kept showing `CONFLICTING` no matter

how many times the agent “fixed” it. A real `git merge origin/main` — the first one actually run against this branch in three attempts — surfaced the genuine conflicts and resolved them for good.)

Driving a Pull Request to Clean

Whenever you are working a PR/MR, run the full **ARDI** loop by default, without being asked: **A**ddress every flagged item, **R**ebut findings that are wrong, **D**efer out-of-scope items to tracked issues, then **I**terate with a fresh review — repeating until the latest review is **fully clean**. Don’t stop at “review-clean, just needs approval” and hand triage back; keep the cycle going until it’s genuinely clean.

Continuously monitor every PR/MR you are actively working until it reaches that terminal state. At every periodic check-in, and again after any push or base-branch advance, query the current head for all three surfaces: mergeability (including conflicts), every CI workflow/check run, and both formal reviews and top-level/inline review comments. A conflict, CI failure, or newly posted finding is ARDI work immediately — sync and resolve the conflict, investigate and fix or track the CI failure, or disposition the finding — not merely a status item to hand back to the user. Keep polling while a review or check is in progress; do not call the PR clean from an earlier head or from green CI without a current-head review verdict. This applies transitively to PR-driving workflows such as `gi`, `gii`, and `ardia`; only monitor PRs the session owns or has explicitly claimed, so the rule does not authorize changing someone else’s work.

The loop’s terminal action is to **report the PR ready, not to merge it**. Merging is human-gated — it happens only on an explicit human “merge it” (the `merge-it` skill), never as a step ARDI takes on its own. So when you carry a PR across a `ScheduleWakeup` or `/loop` wait, **never** bake a self-merge directive like “if clean and CI green, merge it” into the wakeup/loop prompt: a scheduled prompt fires back as a user-role turn, so a self-authored “merge it” only *looks* like human approval (and Claude Code’s auto-mode classifier will rightly deny it as a self-authored merge). Drive to fully clean, report ready, and leave the merge — and any other destructive one-off, e.g. a `gh workflow run` that force-pushes — for explicit human authorization.

Because the loop ends there, **the clean verdict is also where ums runs** — don’t hold the pass for the merge, which is on the human’s clock rather than this session’s and may land after a `/clear` or not at all. See `CLAUDE.md`’s “Run UMS proactively, as learnings accumulate”; the merge-time pass in `post-merge` then only has to cover what the merge itself taught.

The one exception: if the human has explicitly granted the `mwc` (merge-when-confident) session permission, that grant is a live human instruction, not a self-authored one, so baking a self-merge step into a wakeup/loop prompt is fine for the rest of that session. See `mwc` for the grant’s scope and limits.

In the **clear-all family** (*ardia*, *gia*, *gii*, *gip*), “report ready, don’t merge” gates only the merge — it does **not** pause the sweep. A clean-but-unmerged PR is not a stop; move to the next item, and stack it when it isn’t naturally independent of that PR. See [stack-dont-pause](#).

Self-review against the project’s own stated conventions before every push, not just the first — and don’t just re-read the criteria, actually run the applicable review skills against your own diff and iterate on what they find, the same ARD cycle you’d run against an external reviewer’s findings. Don’t treat the review bot as the mechanism that discovers a project’s documented conventions — self-apply them first. When a project’s own `CLAUDE.md` (or equivalent agent doc) already states specific criteria — a DRY/no-duplication rule, a doc-sync checklist for a new input, a changelog-category rule, a citation requirement, a “new logic needs test coverage” norm, a prose-quality check like **fact-check-prose**, **fix-forward-references**, or **detect-informal-definitions** — a first-pass implementation checked only against feature correctness forces the review loop to spend a round re-deriving what the project’s own docs already said. Before every push, re-read the project’s own stated review criteria and actually invoke the review skills/checks it names against the diff (not just recall them from memory), the same way an external reviewer would apply them. Address every finding your own self-review surfaces — fix, rebut, or defer, exactly like the ARD step above — before the push goes out; a self-review that finds issues and pushes anyway has only moved the round to the external reviewer instead of skipping it. Repeat until your own self-review pass is clean, then push. ([gha#219/#220](#): one review round surfaced five findings — a DRY duplication, an incomplete-coverage doc overclaim, a wrong changelog category, an uncited claim, and missing test coverage for new logic — all catchable this way, since each was a direct match against gha’s own `CLAUDE.md` conventions, not new information the review surfaced.)

Pre-push checklist

Pause point: after committing, before git push. Do-Confirm — the items are independent, so work in whatever order suits the round and confirm all seven here. Per [skill-checklists](#); every item below exists because the bullets in this fragment record it failing at this exact boundary.

- ☐ **The whole test suite ran**, not the files you predicted the change touches, and the tests/failed/**skipped** triple was read — a non-trivial skip count means re-running with the gating flags set (`NOT_CRAN=true`, and whatever else un-gates a conditional skip).
- ☐ **Generated trees were regenerated** if the diff (or a `main` merge) touched a generator’s inputs, and the PR body states how many changed files are generated.
- ☐ **Added lines were scanned** for banned punctuation and multi-sentence lines, run *after* committing and with the three-dot range (`origin/main...HEAD`) — a pre-commit run reports on the wrong tree, and a two-dot range re-attributes whatever `main` deleted to you.

- **The changelog entry and the PR description were re-read** against the new behavior, not just the code — neither is in the diff, so no reviewer and no grep will catch a stale one.
- **The diff's deleted lines were read** (`git diff origin/main...HEAD | grep '^-'`), and each one was a decision rather than collateral from an edit's blast radius — a reviewer reads every deletion as deliberate and will rationalize an accidental one.
- **main was merged in** if it moved, with version parity re-checked afterward, so the round costs one review run rather than two — and any whole-file count a merge can worsen (spliced changelog bullets) compared before against after, since a defect caused by a *deleted* line is invisible to every added-lines check ([sync-with-main](#)).
- **Killer item: the push landed.** `git rev-parse HEAD origin/<branch>` agree before any reply asserting a fix. This one is marked because its failure is not an omission but a **false claim about state**, which a reviewer has no reason to doubt: CI reports green because it correctly validated the older head, and the session's own recollection agrees with the reply.

A clean verdict does not certify that your diff contains only what you meant, because a reviewer cannot tell an accident from a decision. Every check above tests whether the diff is *correct*. None of them tests whether it is what you *intended*, and those come apart whenever an edit does something extra — a replacement string that drops neighbouring lines, a global substitution that rewrites more than the flagged occurrence, a stray hunk carried in from another branch.

A reviewer reads the diff as a set of deliberate choices, since that is the only thing a diff can present. So it does not report the extra change; it *explains* it, and often well — constructing a plausible rationale, grading the result appropriate, and moving on. That is worse than silence. Silence leaves the change unexamined, while a reasoned endorsement converts it into a decision the thread now records as settled, and any later reader finds an accident with an argument attached.

The tell is reading a review that justifies something you have no memory of choosing. Treat that as a prompt to check the diff rather than as confirmation, and note that the review's argument may be perfectly sound — the question is not whether the change is defensible but whether anyone decided it.

Deletions are where this concentrates, because an addition is something you wrote and a deletion is usually something that got displaced. `git diff origin/main...HEAD | grep '^-'` lists them in one command, and on a prose diff the list is normally short enough to read in full.

- **Do:** read your diff's deleted lines before pushing, and confirm each one was a decision rather than a casualty of an edit's blast radius.
- **Do:** say plainly, in the thread, when a review has blessed something unintended — the reviewer cannot know, and its verdict will otherwise stand as the record.
- **Don't:** treat a clean verdict as evidence about intent; it is evidence about correctness only.

- **Don't:** keep an unintended change because the reasoning offered for it turned out to be good.

(Morrison-Lab/ai-config#922, 2026-07-30: a replacement string written to add one block silently dropped the three `Do/Don't` bullets belonging to the entry above it, leaving that entry with prose and a case record but no labelled pair — which `CLAUDE.md`'s “Record both the pattern and the anti-pattern” specifically asks for. `claude-review` returned Ready for merge, analysed all three deletions, and concluded they were “appropriate”, reasoning that two restated surviving prose and the third was superseded. The third point was right and the other two were not; the bullets were restored, one reworded, and the deletion count fell from seven lines to two.)

A clean verdict does not discharge the self-review against project conventions either, and the reviewer's own “not a finding” is where that shows up. The section above covers an **accident** a reviewer explains rather than reports, and its check is reading the diff's deleted lines. Here that check finds nothing. The choice is deliberate, nothing was displaced, and the diff contains exactly what you meant — it simply violates a rule stated verbatim in the repo's own `CLAUDE.md`.

A reviewer can notice such a choice, analyse it correctly, and still grade it acceptable, because it is judging whether the code is defensible rather than checking it against the project's written rules. That verdict arrives under a heading like “Observations (non-blocking)” and closes “Not a finding”, which is stronger than the severity labels `address-every-comment` already warns about: “nit” downgrades an item, while “not a finding” retires it. So the part of a review most likely to be skimmed is the part a genuine convention violation is most likely to sit in.

The pre-push self-review this fragment already requires is the only thing that catches it, and a clean external verdict is exactly what makes that step feel finished.

- **Do:** re-run the project-conventions check against your own diff after a clean verdict, not only before the push.
- **Do:** read a reviewer's “observations” and “not a finding” items as candidate violations, and grep `CLAUDE.md` for whatever they discuss.
- **Don't:** let a reasoned “belt-and-suspenders is fine” settle a question the repo already answered in writing.
- **Don't:** treat a non-blocking label as deciding whether an item gets checked at all.

(Morrison-Lab/ai-config#965 at b85941c, 2026-07-31: a diagnostic block ran both `git merge-base --is-ancestor HEAD origin/main` and `git rev-list --count origin/main..HEAD`, where `CLAUDE.md`'s `wrap-up-sweep` section says verbatim to “resist adding an ancestry check beside the first of those”, since the two confirm one thing twice rather than two things once. `claude-review` returned Ready for merge, called the pair “logically equivalent (both express `HEAD <= origin/main` in ancestry)”, judged that “presenting both is reasonable as belt-and-suspenders for a diagnostic block”, and closed the item “Not a finding”; a second reviewer comment returned Ready for merge at the same head. The same block's `--is-ancestor ...`

&& echo "pure upstream history" was graded the same way, although [address-every-comment's own ai-config#868](#) case record already establishes that `--is-ancestor` exits 2 or higher on a pruned ref and && fails on any non-zero status — measured here, a bogus ref gives `rc=128` and the two-arm form still reports “not ancestor”. Both verdicts were wrong; fixed in `0c19d3c`.)

Proactively self-correct a technical claim you already told a reviewer, the moment further testing shows it was wrong — don't wait for the reviewer to catch it. If you stated a rationale (an approach is safe, a risk doesn't apply, a backstop exists) and then discover through your own follow-up verification that it's false, post the correction with the actual evidence immediately, rather than leaving the stale claim standing until a review round re-raises it. This keeps the review loop converging instead of churning on a claim you already know is wrong. ([d-morrison/rme#989](#) / [ucdavis/epi204#363](#): after telling both reviewers `references.bib` didn't share `CLAUDE.md`'s union-merge corruption risk, a follow-up merge simulation showed it does — posted the correction with repro steps on both PRs before either reviewer re-raised it.)

A fix is not “pushed” until it is on the PR's head commit — verify with a SHA comparison before telling a reviewer you pushed it. From inside a session, an edited working tree and a pushed commit feel identical, so a round that edits the files, writes the reply, and never runs `git push` produces a reply asserting a fix that does not exist on the branch. Nothing contradicts it: CI reports green, because it correctly validated the older head; the next review round reviews code without the fix; and the session's own recollection of having made the change agrees with the reply. That makes it worse than an ordinary wrong claim — it is a false statement about *state*, which a reviewer has no reason to doubt and no cheap way to check. Before posting any reply that asserts a push, compare `git rev-parse HEAD` against the PR's own `head.sha` (`pull_request_read get`); if they differ, push first, then reply naming the real SHA. Run the same comparison in every periodic check-in on a PR you are babysitting, since the failure is silent and survives each round until something explicitly looks for it. This is the [algorithmatize-checks](#) rule applied to your own claims: two SHAs decide it exactly, so never substitute recollection. ([d-morrison/altdoc#54](#), 2026-07-25: two review fixes were edited locally and a PR comment said they were “addressed in the latest push”; the head sat at the pre-fix commit for over an hour, with 14 green checks validating a branch carrying neither fix, until a scheduled check-in compared the SHAs.)

A SHA you put in a PR body or a reply must be read, never recalled — and the PR body is where an invented one survives longest. The bullet above asks whether the right commit reached the branch. This asks the prior question: whether the commit you named exists at all. A short SHA is seven plausible hex characters, so writing one from memory feels like recalling a fact rather than asserting one, and the result is indistinguishable from a correct citation — nothing renders differently, and GitHub neither linkifies nor validates a SHA with no commit behind it.

The PR body is the worst host for it, for the reason [address-every-comment](#) gives about stale paraphrases there: the body is in no diff, so no reviewer reads it as part of the change and no

`grep` over the diff finds it. It is also what a maintainer reads while deciding whether to merge, so an invented SHA misdirects the one reader most likely to act on it.

One command against the value you are about to paste settles it:

```
git rev-parse --verify <sha>^{commit}    # or: git cat-file -e <sha>
```

When a wrong SHA has already been published, correct it **visibly** rather than overwriting it silently — a reader who saw the original cannot otherwise tell a revised body from one that always said this, which is the same reasoning the withdraw-a-stale-blocker bullet below applies to a retracted caveat.

This is the commit-SHA case of the rule `report-mistakes-proactively` states for issue numbers (“never name an issue number before the issue exists”). Same defect, different artifact: an identifier guessable enough to assert casually, with nothing in the repository to contradict it.

- **Do:** read every SHA you cite out of `git rev-parse` or `git log`, and confirm it resolves before pasting it.
- **Do:** correct a published wrong SHA with a visible note naming the real one.
- **Don’t:** write a short SHA from recollection because it looks like the commit you just made.
- **Don’t:** expect review to catch it — a reviewer has no reason to suspect a citation, and the body is not in the diff they are reading.

(ai-config#871, 2026-07-30: the PR body credited a sentence-boundary fix to `1f79a4a`, which existed nowhere in the branch or the repository — `git cat-file -e` returned `Not a valid object name`. The real commit was `fc605f`. Two review rounds read that body without flagging it; it surfaced only when the body was re-read against the diff before declaring the PR ready, which is the `address-every-comment` check above doing work its own rule did not anticipate.)

The read side of that comparison can lag a push by a few seconds, so test the two *local* refs against each other before concluding anything failed. The rule above is an `algorithmatize-checks` case because two SHAs decide it exactly. That holds only as far as both numbers are current, and one of them is fetched over the network: `gh pr view <N> --json headRefOid` can still report the **previous** commit immediately after a successful push.

The failure direction is a false alarm rather than a false all-clear, so it is the safe one — but the reflexive response to it is wrong twice over. The rule’s own remedy is “if they differ, push first”, which is a no-op here, and treating a healthy branch as broken invites an amend or a force-push that manufactures the problem the check was watching for. A check that cries wolf on a clean branch also stops being run, which is the objection `algorithmatize-checks` raises against any instrument whose threshold cannot be trusted.

Local refs cannot lag this way, so they settle it:

```
git rev-parse HEAD origin/<branch> # both local reads
```

- Equal — the push landed, and any disagreement with the PR API is a read-side artifact. Re-read it, preferably on a different surface, rather than re-pushing.
- Different — a genuinely unpushed commit, which is the case the rule exists for.

Note also that `git push` answering **Everything up-to-date** is itself evidence the remote branch already carries the commit.

- **Do:** compare `HEAD` against `origin/<branch>` first when the PR API disagrees, and re-read rather than re-push when those two agree.
- **Don't:** amend, force-push, or re-commit on the strength of an API SHA alone.

(Morrison-Lab/ai-config#845, 2026-07-29: `git rev-parse HEAD` and `git rev-parse origin/<branch>` both read 9a3e722 and `git push` said **Everything up-to-date**, while `gh pr view --json headRefOid` still returned the prior commit 4bf5063. `pull_request_read` got returned 9a3e722 moments later, so the two surfaces disagreed and the git-native one was right.)

A brand-new branch can read back at the wrong commit, so the local two-ref comparison above is not sufficient there. That bullet offers `git rev-parse HEAD origin/<branch>` as the pair that settles the question, on the grounds that local refs cannot lag. They cannot. What they can do is agree with a remote ref that reads back at the wrong commit, and then the pair reports the push as landed.

The failure direction inverts, which is what makes this worth separating. A read-side lag is a false alarm, correctly called the safe one there. This is a false all-clear, arriving on the one instrument that rule offers for deciding the question.

The gap is in the trigger rather than in the remedy. The original SHA comparison does catch this whenever it is run. Nobody runs it after a `git push -u` that printed `* [new branch]`, set the upstream, and exited 0, because that is the least suspicious moment in the round.

`git ls-remote origin <branch>` reads the remote directly rather than through a tracking ref, so it is the instrument that catches it. The corrective push is an explicit refspec: `git push origin HEAD:refs/heads/<branch>`.

The likeliest explanation is a local one, and it reproduces offline. `git push -u origin <branch>` pushes the **branch ref**, not `HEAD`. So a local branch left behind at `main`'s tip, while `HEAD` carries the new commit, produces this whole signature with nothing on the server going wrong. Reproduced in a local bare repo, where no replica and no race exists: the push printed `* [new branch]` and exited 0, `git ls-remote` and the tracking ref both read `main`'s tip, `main..<branch>` held zero commits, and `git push origin HEAD:refs/heads/<branch>`

then reported a real range. That last command is a test as much as a fix, since it answers **Everything up-to-date** when the branch ref and **HEAD** already agree.

Two things weigh against the read-side story, which an earlier draft of this entry weighted equally against the write-side one. A lagging replica cannot invent a value for a ref that never existed before, so its failure mode is the ref reading **absent** rather than reading one specific wrong commit. And a tracking ref is set from what the push sent, which makes its value a client-side fact rather than a later network read.

What stays genuinely unsettled is narrower than either reading claimed: the branch ref's own value at push time was never recorded, so the local explanation is the best supported one rather than a proven one. Note the shape of that, since it is the failure this entry is about. The entry has now over-claimed twice, first asserting a write-side fault, then asserting a parity between two hypotheses that the record does not support either. The practical advice survives all three readings, because the checks below are cheap whichever is right.

Two things about diagnosing one of these. The downstream error misdirects, because opening the PR fails with **No commits between main and <branch>**, which names a base-versus-head relationship and sends you to check the wrong argument. And *** [new branch]** here is the ordinary output for a branch that did not exist before, **not** the deleted-underneath-you signal CLAUDE.md's "Use the existing PR branch" section describes. There the line is diagnostic precisely because the branch had already been pushed to; here it is expected, so the two cases must not be conflated.

The wrong value is the informative part, and it reads at first like noise. A race or a server fault has no reason to land on **main**'s tip in particular, whereas a branch ref cut from **main** and never advanced sits there by construction. So read a wrong value that happens to equal **main**'s tip as pointing at the local ref rather than at the network.

- **Do:** run `git ls-remote origin <branch>` after the first push to a new branch, and compare its SHA against `git rev-parse HEAD`.
- **Do:** run `git rev-parse HEAD <branch>` first when those two disagree, since a branch ref left behind accounts for the whole signature (and note that `--short` rejects a second revision, so pass neither).
- **Do:** re-run plain `git ls-remote` as well, so a ref that self-corrects stays distinguishable from one a re-push repaired.
- **Do:** re-push with `git push origin HEAD:refs/heads/<branch>` when the mismatch persists, and read the SHA range it prints as the confirmation.
- **Don't:** treat a `git push` that exited 0 and printed *** [new branch]** as evidence the commit reached the remote.
- **Don't:** assume `git push -u origin <branch>` sent the commit you just made – it sends the branch ref, which **HEAD** may have moved past.
- **Don't:** credit a corrective re-push with having repaired a remote-side fault when neither of those two controls was run.

- **Don't:** answer a No commits between main and <branch> error by re-checking the base branch argument before checking where the head ref actually points.

(Morrison-Lab/ai-config#985, 2026-07-31: `git push -u origin ums/prose-count-adjacent-to-block`, carrying commit 1611ccc, printed * [new branch], set the upstream, and exited 0. `git ls-remote` showed that ref at 98102a2, which was main's tip. The local `origin/ums/prose-count-adjacent-to-block` agreed with the wrong value, so the two-ref comparison reported the push as landed. `create_pull_request` then returned a 422 reading No commits between main and ums/prose-count-adjacent-to-block. `git push origin HEAD:refs/heads/ums/prose-count-adjacent-to-block` reported 98102a2..1611ccc. Neither `git rev-parse ums/prose-count-adjacent-to-block` nor a second plain `git ls-remote` was run, so the branch ref's own value at push time is the fact the record is missing. Describing that push as "carrying commit 1611ccc" was an inference from the commit just made, not a reading of the ref that was pushed.)

The same false claim arrives as *incoming* state when you pick a PR up mid-flight, and there the SHA comparison usually has nothing to compare. The bullet above governs a claim you are about to make. Its mirror is the claim already sitting on the PR when you arrive: the latest comment says which findings are fixed, so starting from it means starting from a summary rather than from the branch. The rule is the same either way, and only the *instrument* differs, because a claim written for a human rarely carries a SHA to check. "Three fixes landed in two commits" names files, not commits.

So compare the claim against the file list, which decides it in one call:

```
gh pr diff <N> --name-only # DIFF_PR
```

A named file absent from that list was not touched, whatever the comment says. Nothing else in the PR contradicts a claim like this, and the trap is that green CI reads as corroboration when it is nothing of the sort. The checks that would have exercised the claimed fixes are frequently the very checks those fixes were supposed to add, so their absence is the reason CI is green.

The right response is to do the work rather than to dwell on the discrepancy. Say once, in the round's summary, that the earlier claim was not true of the branch, since a reader who saw it will otherwise assume the round was redundant.

- **Do:** run `gh pr diff <N> --name-only` against any inherited "already fixed" claim before deciding a finding is closed.
- **Do:** state plainly in your own summary that the prior claim did not hold, and name the head it was false at.
- **Don't:** treat green CI as evidence that a claimed fix landed.
- **Don't:** infer that a finding is stale because a comment says it was addressed.

(Morrison-Lab/ai-config#804, 2026-07-29: the PR’s own review workflow posted “Three fixes landed in two commits”, naming `bootstrap.sh`, `validate.yml`, and `validate-skills.py`. The head was still the original commit plus a `main` merge, and `gh pr diff --name-only` returned four paths, none of them those three. All ten checks were green, because `validate` did not yet check out the submodule the fixes were about. This is the ownerless cousin of the parallel-session case in `claim-pr`, which assumes a real commit exists to cross-check; here there was none.)

Run that same command before *any* readiness claim, not only against an inherited one — a PR whose branch carries no implementation is green on every check. The bullet above uses `gh pr diff <N> --name-only` **differentially**: it has a claim naming files, and it asks whether those files are in the list. That test needs a claim as input, so when nobody claimed anything it never runs, and the readiness path is exactly the case where nobody has. The **existential** question — does the list have anything in it at all — is the one no rule was asking.

`pr-on-claim` manufactures the hazard by design, and is right to: it opens the PR from `git commit --allow-empty` so the branch has a diff before any code exists. Merge `main` in later and the branch carries two commits, a real history, and no implementation. Every instrument then works perfectly and certifies nothing, because a check that finds no fault in an empty diff and a reviewer that raises no finding against one are both answering a narrower question than the one being asked.

Note which nearby checks *pass* on such a branch, since their passing is what makes the state feel verified. `git rev-parse HEAD origin/<branch>` agree, so the pre-push checklist’s killer item is satisfied. The branch is not behind `main`. `get_commits` returns two, so a sweep keyed on zero commits does not flag it. `fully-clean`’s two criteria are each satisfied **maximally** by an empty diff, since neither has a term about content.

- **Do:** run `gh pr diff <N> --name-only` before reporting a PR ready, and read the returned paths against what the PR says it does.
- **Do:** treat an empty return, or a return holding only a `main` merge’s incidental paths, as the PR carrying no implementation.
- **Don’t:** count the claim commit or a `main` merge as work — neither is implementation, and both give the branch a plausible history.
- **Don’t:** read all-green CI plus a finding-free review as evidence a PR contains anything; on an empty diff that is the expected result.

(2026-07-30/31, a `ucdavis/bcs` session: a PR was reported **CLEAN** with every check passing, on a branch holding the empty claim commit plus a `main` merge and nothing else. Nothing had gone wrong with any instrument. The implementation had never been pushed, and no check, no reviewer, and no rule in this file was asking whether there was one.)

When the change affects downstream consumers, validate it against a real consumer repo before reporting the PR ready — a package’s own test fixtures are built to

exercise its code, not to resemble the packages that will actually use it. Fixtures are minimal by construction and tend to share one shape, so whole branches of new code can be structurally unreachable from them. A real consumer brings the input variety fixtures lack, and it is usually one clone plus one command to check.

This bullet, and the two further down that also turn on fixtures, are all about **coverage** — a fixture too thin to reach the code. **fixtures-are-not-evidence** covers the opposite direction: a fixture that works perfectly, and an inference drawn from its behaviour back to the real system it stands in for.

Three classes of gap this catches, none of them findable in a fixture:

- **Input shapes no fixture happens to contain.** A real package carries metadata the fixtures never needed — an entry of a different kind, an extra tag, an unusual name — so a branch written for it has never actually run on real input.
- **Message formatting under real counts.** Fixtures usually trip the plural path; a real repo hitting the same code with exactly one item exercises the singular wording, which no test asserted.
- **The migration/upgrade path, as opposed to the fresh-install path.** This is the one fixtures can never reach: a fixture is created new by the test, so it always gets the current templates. An existing consumer has the *old* config, and whether the feature reaches it at all is a different question from whether it works. Verify the claim in the changelog by running the documented migration step, rather than describing it.

Do it against a throwaway copy and push nothing to the consumer; the deliverable is evidence in the PR, not a change there. Record what the run covered in a PR comment, so a reviewer can see which paths real input reached. (d-morrison/altdoc#34: running the new reference-index generator against d-morrison/rpt covered a `\docType{package}` topic, the singular form of a missing-topic warning, and the documented “existing settings files do not pick this up automatically” caveat — confirmed by the page generating while `grep -c reference.html docs/index.html` returned 0. None of the three were reachable from the repo’s own fixture packages.)

Verify a blocker you assert in a PR body or a reply, with the same rigor you apply to a reviewer’s claims — a stated blocker becomes a premise other people build on. The reviewer-facing checks above all point outward: verify the suggestion, verify the literal, verify the push landed. The inward case is easier to miss, because a limit you hit yourself feels like an observation rather than a claim. It is still a claim, and writing it into a PR body publishes it as settled fact: a reviewer reading “the pinned tool is unavailable in this sandbox” will reason from it, recommend a follow-up around it, and never re-test it, so one unverified sentence quietly redirects the review. Before asserting that something is unavailable, blocked, or impossible here, actually attempt it once — an install, a fetch, a single command — and say what you tried. A negative result from one incidental symptom (a failed version query, a single 403) is evidence the thing is not *already set up*, not evidence it cannot be. When a blocker you published turns out to be false, correct it where it was published, not only in the thread

that surfaced it. (d-morrison/altdoc#76, 2026-07-27: the PR body said roxygen2 8.0.0 — the version DESCRIPTION pins — was unavailable, inferred from one failed `packageVersion()` call with no install attempted. The review built a “this may need a follow-up” recommendation on top of it. A single `install.packages()` disproved it, and the regeneration landed in the same round the finding did.)

Name the specific gate when you report a blocker, not a category word that happens to be one of several. The rule above governs *whether* something is blocked, and its remedy is to attempt the thing once. This governs *why*, and it fires after that remedy has already succeeded: the call was attempted, it genuinely failed, and the blocker is real. Only the attribution is wrong, which is why nothing about it feels like an unverified claim – the part that usually goes unchecked has, this time, been checked.

The hazard is a platform with two gates whose refusals read alike. `resolve_review_thread` on a transferred repo fails under either spelling of the owner, saying `Access denied` both times, for unrelated reasons: the old owner trips a comparison between the thread’s node and the declared `owner/repo` string, and the new owner trips the session’s own repository allowlist. Only the second of those is scope. So “blocked for scope reasons” is not a loose summary of the first. It names a mechanism that was not involved, and it names one that genuinely exists on that platform, which is what lets it survive re-reading.

That last point is the whole cost. A category word that is also the proper name of one mechanism cannot double as the generic term for its family, because a reader cannot tell which you meant, and the wrong reading is actionable: someone told a call failed on scope will reach for the other owner, which fails too. Quote the error’s distinguishing clause instead of classifying it. The quote is usually shorter than the paraphrase, it is checkable, and it stays correct even when your model of the platform is not.

- **Do:** quote the clause that distinguishes the failure, and name the gate it belongs to.
- **Do:** re-read a blocker you have restated several times, since a paraphrase repeated across status reports hardens into the record.
- **Don’t:** use one mechanism’s own name as a generic word for its category.
- **Don’t:** treat having verified *that* something is blocked as having verified *why*.

(2026-08-01, Morrison-Lab/ai-config worked from a d-morrison-scoped session: an unresolvable review thread was reported as blocked “for scope reasons” across roughly six status updates, while the failure actually observed under that spelling was the node-versus-declared-string comparison. `memories/github-mcp-tools.md` records both gates and their verbatim errors.)

A blocker that was true when you published it can stop being true while the PR is open, and withdrawing it is your job, not the reviewer’s. The verify-a-blocker bullet above covers a blocker that was never true, and the gate-naming bullet between it and this one covers a real blocker whose mechanism was misnamed. This is the harder case, because the caveat was correct and diligent when written, so nothing about it reads as a defect later — and

a sentence saying “this could not be checked” is one nobody re-checks, least of all the reviewer, who has no way to know the environment moved. It keeps steering the review regardless: a verdict can repeat the caveat back as an accepted limitation, which makes the stale claim look corroborated. So when the cause of a blocker changes — a host unblocked, a tool installed, a quota reset, a dependency published — re-run the check and withdraw the caveat where it was published, saying explicitly that it is withdrawn rather than quietly deleting the sentence. A reader who saw the original needs to know it was retested, not be left wondering whether it was ever true. (ai-config#774, 2026-07-28: the PR body said four `adv-r.hadley.nz` anchors could not be verified because the host was egress-blocked, which was accurate when written. The host was unblocked mid-session, and all 16 URLs then verified 200 with every anchor resolving. The review had already absorbed the caveat — it listed those anchors as “unverified per the PR body’s own caveat ... not a new finding” — so leaving it would have shipped a limitation that no longer existed, blessed by a reviewer who could not have known.)

Landing a fix falsifies whatever prose documented the defect, and that prose is never in your diff — so grep for it rather than expecting to be reminded. The bullet above covers a caveat **you** published on **this** PR, which the environment then moved out from under. This is the case where you moved it yourself, and where the stale text lives in the standing corpus rather than on the PR: a memory bullet describing the hazard, a README warning about it, a docstring asserting the behaviour you just changed.

The sync rules in [address-every-comment](#) all end at the PR’s own artifacts — the changelog, the PR body, a skill’s inline restatement — and each prescribes grepping the diff. Documentation of a defect cannot be found that way, because not being in the diff is the entire property that makes it survive. So the trigger has to be the fix itself, and the search has to leave the files you edited.

Two shapes, and the second is worse because it was never true.

- **Prose staled by the fix.** It was accurate when written, so nothing about it reads as a defect, and a workaround it prescribes becomes active misdirection the moment the thing it worked around is gone. Keep the entry where the old behaviour explains something — most of a corpus is written against it — but mark plainly that it is history and name the change that ended it.
- **Prose asserting conformance to a reference.** A docstring saying the code “follows” some reference implementation is a claim about two artifacts, and your own divergence falsifies it. This one is not staleness at all: it was false before you arrived, and it is load-bearing, because a reader checking the code against the reference stops at the sentence saying someone already did.
- **Do:** grep the repository for the defect, the workaround, and the behaviour you changed, before calling a fix complete.
- **Do:** mark a superseded entry as history and name the change that ended it, rather than deleting it, when the old behaviour still explains other text.

- **Don’t:** treat a clean grep over the diff as coverage — the stale prose is outside it by construction.
- **Don’t:** leave a doc asserting conformance to a reference standing when the code diverges; correct the claim in the same change that establishes the divergence.

(ucdavis/bcs#534, 2026-07-30/31: standardizing a G-computation CIF over the observed age distribution falsified two documents at once. `compute_gcomp_cif_ab507bs()`’s roxygen had read “this function follows the SAS pipeline’s plug-in-at-the-mean approach”, which the SAS program does not do — false before the fix, and quoted by the fix’s own changelog entry as such. A row in `inst/docs/program_steps.qmd` described the retired behaviour and was refreshed in a separate commit. Concurrently, ai-config#951 diff-scoped `scripts/semantic-line-breaks.py`, which falsified `memories/tools.md`’s entry prescribing “format new prose by hand” as the workaround; that entry was kept and marked ****Fixed in ai-config#951.****, which is the first shape handled correctly.)

An instruction’s own suggested code is not exempt from the project-conventions self-review above. The self-review rule assumes you wrote the diff; a snippet handed to you in an issue, a task description, or a design doc slips past it, because adopting someone else’s suggestion does not feel like authoring. It is authoring — once pushed, it is your diff, and the project’s conventions bind it exactly as they bind anything you wrote yourself. Run the same convention check over borrowed code before pushing it, especially when the suggestion is a plausible-looking one-liner and the convention it breaks is documented rather than linted. (d-morrison/altdoc#73: the issue proposed ending a function with a bare trailing `hashes`, which reads as a fix for the fragility it names but is still an implicit return, so a statement added after it silently becomes the return value. The lab manual asks for an explicit `return()` regardless. Review caught it; the project’s own stated convention would have, one step earlier.)

When the code path under test has a staging or transform step between input and output, a passing unit suite is not evidence it works — exercise the real path once. Fixtures instantiate the shape the test author had in mind, so a wrong assumption about *where* the code runs is invisible to every one of them: the tests and the bug share the assumption. This is the same gap the downstream-consumer rule above covers, one level in — there the missing variety is the consumer’s input, here it is the pipeline’s own directory layout, timing, or intermediate representation. One real invocation is usually cheap, and it tests the assumption the fixtures encode rather than re-confirming it. (d-morrison/altdoc#76: a guard checked for the copied logo under `docs/`, but the `quarto_website` path stages into `_quarto/` first, so the logo line was dropped on every render of the one generator the feature wired up. Seventeen unit assertions passed throughout; one throwaway render found it immediately.)

When new code branches on a third-party tool’s behavior, read that tool’s own config or docs for the specific behavior — don’t infer it from what the tool broadly does. The bullet above covers your own pipeline’s layout; this one covers the tools that pipeline drives. An inference of the form “it builds HTML, so link to `.html`” is exactly the shape that feels too obvious to check, and a tool’s defaults routinely contradict it. Two properties make this

worse than an ordinary wrong guess. The inference usually lands in a branch your own fixtures cannot reach — you have no fixture for someone else’s renderer — so the test suite agrees with you. And it produces output that is well-formed and plausible (a link, a path, a flag), so a reviewer skimming the diff has nothing to catch, and the failure surfaces only in a consumer’s published site. Name the setting you are relying on, and check its actual default before writing the branch. (d-morrison/altdoc#78, 2026-07-27: a generator-to-extension map gave mkdocs `.html`, reasoning that mkdocs compiles Markdown to HTML. Its `use_directory_urls` default is `TRUE`, so it serves `/man/foo/` and never `/man/foo.html` — every reference link the feature emitted for that generator would have 404’d. Caught in review, not by the 39 tests.)

A regression test written alongside a fix can lock the bug in rather than catch it — assert the two paths that diverge, not the one you just touched. A test authored in the same pass as the code tends to record what the code *does*, because you run it, see it pass, and move on. That is usually harmless. It becomes a lock when the fixture is thin enough that the buggy and the correct path produce the *same* output: the assertion then encodes the degraded result as intent, and every later reviewer reads a green suite as evidence the behavior was chosen. The next round’s finding lands on your test, not just your code.

The tell is the same each time: **a fixture missing the input variety that makes the two paths differ.** So when a bug is an asymmetry — nested versus top-level, second render versus first, one generator versus another — build the fixture so both sides are present and assert them together. Either side alone is unfalsifiable, since the case that reveals the bug is the *comparison*. Then prove it: revert the fix and confirm the new test actually fails. A regression test never seen to fail is a guess about what it covers. (d-morrison/altdoc#78, 2026-07-27: twice. A `.pdf` vignette test asserted the entry’s extension but never its label, so an extension leaking into the label passed; and a nested-article test built no source tree, so top-level and nested resolved identically and a nested-only title bug was pinned as expected output. Both were found by review reading the test, not the code.)

A systematic audit done by skimming is worse than the one-at-a-time version it replaces. Batching a check — “rather than wait for the next round to find divergence number four, compare all four at once” — is the right instinct, and it inverts if each lookup gets less care than it would have alone. Two things make the batched form more dangerous, not less. Its output is usually a claim recorded somewhere durable (a comment, a doc, a table), so an error is published rather than merely held; and it arrives labelled *audited*, which is precisely the word that stops the next reader from checking. A wrong comment in a block written to prevent a specific future change invites that change while appearing to forbid it. Concretely: when the thing being audited is a function, grep for the function, not for a pattern in its file — a file with several functions will hand you the first match, which is often not the one you mean. Name the function in whatever you write down, so the claim stays checkable. (d-morrison/altdoc#78, 2026-07-27: a commit written to get ahead of a one-finding-per-round loop claimed mkdocs’ sidebar matched only `\.md`. It matches `\.md$|\.pdf$`; the grep had returned a different function 120 lines above the sidebar builder in the same file. Caught by the very next review round.)

Adding an explanation supersedes whatever the file already said about the same thing, so re-read the older passage — your own diff is the likeliest source of a contradiction nobody flags. The sync rules in [address-every-comment](#) all fire on an external trigger: a reviewer quotes a phrase, or behavior changes and a changelog goes stale. This one has no trigger at all. You add a paragraph explaining that something was misunderstood, and the note recording the original misunderstanding sits a few lines below, still stating it as fact. Nothing conflicts, no check fires, and both passages read plausibly on their own — but a reader who reaches the older one first comes away with exactly the belief the new text was written to remove.

The tell is a diff that adds an explanation, a correction, or a “what this actually means” paragraph near existing prose. Re-read the surrounding passage as a whole rather than diffing your addition in isolation, and treat a historical record (“we observed N of X”) as a claim your explanation may have just falsified. When the older passage recorded a *different* session’s observation, correct it with reasoning that stands on its own rather than restating it as though you had seen it — an inference presented as an observation is the same defect one level up. (ai-config#770, 2026-07-28: an added explanation established that seven reported orphans were misclassifications, while the note two lines below went on calling them “already deleted from the repo.” Caught by review, in the same hunk as the text that contradicted it.)

The same rule applies within a single diff, and there nothing prompts the check at all. The version above compares your addition against the *existing* file, so re-reading the surrounding passage catches it. The harder case is a diff that both adds an explanation arguing against some older wording **and** rewrites that wording, in the same changeset. Then the argument survives while its target does not, and every file still reads plausibly on its own: the new prose is coherent, the rewritten passage is coherent, and only the cross-reference between them is stale. There is no older text to go back and re-read, which is the cue the other version relies on.

So when a diff rewrites a passage, grep the rest of the diff for references to what that passage used to say, not just for references to the file. A rebuttal of the form “the section below already says X” is the shape to watch, since it pins the argument to wording the same commit may be deleting. Prefer stating the anti-pattern directly over citing another section as the thing being argued against; a self-contained sentence cannot go stale when its neighbour changes. (ai-config#801, 2026-07-28: a new UMS entry argued against the `/clear` section’s “disclose the owed pass in the flag” line while the same PR rewrote that line to say the opposite. Review caught it before merge; the fix was to drop the cross-reference and state the point inline.)

And when the explanation you add is a *mechanism* claim, test the class it distinguishes, not just the sample in front of you. The bullet above is about contradicting old text; this is about the new text being unfalsifiable on the evidence you gathered. A classifier validated on a population containing no positive instance of the class it is supposed to catch will report a clean result either way, so “it returned zero” is not evidence it works — it is the same missing-input-variety tell the regression-test bullet above describes, moved from a fixture to a diagnostic. Ask what a true positive would look like, confirm one exists in what

you tested, and if none does, say so instead of claiming the mechanism separates the cases. (ai-config#770, same day: a `git log -- skills/<name>` probe was said to separate “deleted from the repo” from “never ours” *exactly*, on the evidence that it reported zero false orphans. The repo contained no deleted-but-still-installed skill at all, so there was nothing for it to get wrong; and `git rev-parse --is-shallow-repository` returned `true`, meaning anything deleted before the shallow boundary would have been silently misread as harness-provided. The claim went into a PR reply before either check was run, and ai-config#765 had independently reached the correct conclusion.)

A symptom that stops reproducing is a fix having landed, until you have checked otherwise — reaching for nondeterminism is the attractive wrong answer. The bullet above governs a mechanism claim you write into a file. The same defect arrives in a status report, and there it is easier to publish, because “the check is just flaky” sounds like a complete explanation while resting on nothing. It is also unfalsifiable from a single observation and predicts nothing, which is exactly why it feels safe to say.

The shape to watch for: a known-failing check, a tracked false positive, or a reproducible bug goes quiet, and you explain the silence with a property of the *tool* rather than a change in the *world*.

Check for the merge first, because the instrument is a timestamp comparison and it costs one API call. Compare when a candidate fix merged against when each observation was made. That turns “it seems flaky” into a before/after table with a negative control — a far stronger claim than the one you were about to make, and one a reader can act on.

- **Do:** look for a merged fix, and date it, before attributing a vanished symptom to anything.
- **Do:** report the before/after with its timestamps, so the negative control is visible rather than asserted.
- **Don’t:** explain a symptom’s disappearance as nondeterminism on the strength of one clean run.
- **Don’t:** carry such a claim into an issue or a decision doc, where it argues against the very fix that produced the silence.

(ai-config#827, 2026-07-29: the Jules AI reviewer approved a diff carrying both of its known false-positive triggers, and the first explanation drafted was that the false positives are non-deterministic. ai-config#817 had in fact merged an `extra_instructions` fix at 21:30:51Z, between #820’s block at 19:43 and #827’s approve at 22:51. The nondeterminism claim was about to be posted to gha#366 as evidence, where it would have argued against porting the fix that actually worked.)

Verify a command, path, or flag *you* write into a doc, with the same rigor [address-every-comment](#) demands for one a reviewer suggests. That rule and the verify-a-blocker rule above both point outward, at a claim someone else made or at a limit you hit. This is the one you author from scratch, and it is easier to miss than either, because inventing a plausible command does not feel like making a claim at all — it feels like remembering one.

The shape is a CLI invocation, a file path, or a flag written into documentation, a comment, or a registry, where the surrounding prose is carefully sourced and the literal is not. A reader reaching for it gets **unknown command**, and they get it while following a document whose every other sentence checked out, so they are likelier to doubt their own setup than the doc.

Settle it against the tool's own source or `--help` rather than recollection. For a CLI, the subcommand registration list is definitive and usually one fetch away, and it beats a docs page because it cannot lag the release you are describing. Quote what you checked in the commit message, so the next reader inherits the evidence instead of the assertion.

- **Do:** confirm every literal you invent against the tool's own source or help output before it lands in a doc.
- **Do:** cite the file or command you checked, so the claim stays falsifiable.
- **Don't:** infer a subcommand from a family that has its siblings (`gh label list/create/edit` does not imply `gh label view`).
- **Don't:** treat a literal as exempt because the prose around it is well-sourced — the literal is the part a reader executes.

(Morrison-Lab/ai-config#834, 2026-07-29: a `GET_LABEL` registry row shipped `gh label view <name>`, which does not exist — `cli/cli's pkg/cmd/label/label.go` registers `list`, `create`, `clone`, `edit`, `delete`. Caught by review, in the same file where the previous round had declined to extend an untested claim from `gh issue create --label` to `gh issue edit`. The reviewer's own enumeration of the real subcommands also missed `clone`, which is why the fix cited the registration list rather than the finding.)

Run that check over your own fix, too — the remedy for an unverified literal is where the next unverified literal goes. The rule above fires when you notice you are writing a literal. Answering a finding does not feel like that: it feels like careful work, and the care is real, so the fix inherits an assumption of rigor from the diligence of writing it. The specific rule you are in the middle of applying is therefore the one least likely to be applied to its own application.

A correction also tends to *add* literals rather than merely repair one. Citing a source properly means naming the tool, the flag, the version, the file — each an assertion, each as guessable as the one under review, and none of them the thing the finding was about. So the fix can carry more unverified surface than the original did.

Nothing external catches this. The reviewer sees a fix that addresses the finding and confirms it, per the clean-verdict entry above, and the thread then records the item as settled twice over.

- **Do:** re-run the rule you are applying against the text of your own fix, before committing it.
- **Do:** say in the thread when a fix's own draft tripped the same rule, since that is the only place the near-miss is visible.

- **Don’t:** treat the effort of writing a correction as evidence the correction is verified.

(Morrison-Lab/ai-config#929, 2026-07-30: a review found `--failed` documented from recollection. The fix quoted `gh run rerun --help` correctly and anchored it to “`gh 2.83.0`” — a version invented in the same breath, where `gh --version` reported `2.96.0`. Caught before committing only by running this rule against the fix, which is the entire mechanism; the round-2 review confirmed the corrected text and would have confirmed the wrong version just as readily.)

When regenerating a generated tree makes it most of the diff, say so in the PR body — otherwise a reviewer reads it as pollution and blocks. Two failures share one root here, and both cost a round.

The first is editing the generated file rather than its source. A generated file usually declares itself in a header, and that header is the last thing you read when you have already found the line you wanted to change; the generator then reverts the edit and the sync check fails. Grep for **generated** by before editing any file you did not create.

The second is what happens once you regenerate correctly. A one-line source change can rewrite hundreds of files, and a reviewer working from a truncated diff sees only the generated bulk — so the finding comes back as “revert the unintended changes”, where complying would break the very check that demanded them. The PR body is the only place that can pre-empt this, since it is what a reviewer reads before the diff. Lead with the ratio and a per-path table marking each path generated or hand-written.

- **Do:** grep a file for a generated-by header before editing it, and change the source instead.
- **Do:** state in the PR body how many of the changed files are generated, and name the hand-written ones.
- **Don’t:** revert generated output because a reviewer calls it noise — check first whether the sync check requires it.
- **Don’t:** assume a reviewer sees the source files; on a large diff they frequently do not.

(Morrison-Lab/ai-config#834, same day: a fix was applied to the generated `tool-mappings.md`, which `sync-codex-skill-wrappers.py` then overwrote, failing `validate` with `stale tool-mappings.md`. Redoing it in `tool-mappings.yml` regenerated 175 `codex-skills/` wrappers, and Jules returned `VERDICT: block` twice for “bulk pollution”, its second verdict noting it had read only a truncated diff. `claude-review` called the same finding a false positive at the same head.)

Run the whole test suite before pushing, not the files you predict the change touches — and check that the ones you ran were not silently skipped. The pre-push self-review above assumes your local green means something. Two things quietly hollow it out, and they compound: you choose a subset, and the subset then reports success without having run.

The subset is chosen by predicting blast radius, which is exactly the judgement the change calls into question. The tests that break are frequently *not* in the files you edited: a test elsewhere asserts the behaviour you are changing, often with a comment stating the rationale your diff invalidates. Nothing about editing `R/foo.R` suggests opening `test-bar.R`, so the prediction feels complete while omitting the one file that matters.

The second is worse, because it looks like evidence. A conditional skip — `skip_on_cran()` without `NOT_CRAN=true`, `skip_if(!.venv_exists())`, `skip_if_offline()` — turns the test that would have caught the bug into a pass. `devtools::test()` sets `NOT_CRAN` for you — it applies `withr::local_envvar(r_env_vars())`, and `devtools` documents that set as “the standard environment variables set by devtools”, singling out `NOT_CRAN` as “of particular note for package tests” ([R/test.R](#), [R/check.R](#)). `testthat::test_file()` and `testthat::test_dir()` do not, so the harness you reach for by hand for a quick targeted run is exactly the one that skips. Setting it explicitly anyway (`NOT_CRAN=true Rscript -e '...'`) costs nothing and is what the rest of this corpus does. So read the **skip count**, not just the failure count: a run reporting 0 **failed**, 20 **skipped** has told you almost nothing, and is indistinguishable at a glance from one that verified everything.

Match the environment the change will be judged in, too. Running with an env var set that CI does not set (or vice versa) reproduces a *different* configuration, and its failures and passes both mislead.

- **Do:** run the full suite before pushing, and state the tests/failed/ skipped triple rather than “tests pass”.
- **Do:** set the flags that un-gate conditional skips, and re-run if the skip count is non-trivial.
- **Don’t:** scope a local run to the files you edited — the test asserting the old behaviour is usually somewhere else.
- **Don’t:** read a green subset as a green suite, or a skip as a pass.

(d-morrison/altdoc#95 and #96, 2026-07-29: twice in one session. On #95 a test asserting “aborts when no venv is configured” read that precondition from the ambient environment; the local run missed it because `NOT_CRAN` was unset and `skip_on_cran()` skipped that very test, and Windows R-CMD-check caught it. On #96 `test-11ms_txt.R` asserted non-recursive discovery for `docsify` — the exact behaviour the PR changed, with a comment stating the now-false rationale — and was not among the files run locally, even though `.11ms_txt_vignettes()` was one of the functions edited. Windows caught that one too. The subsequent full-suite run was itself misleading in the opposite direction: run *with* an env var `CI` does not set, it reported two failures belonging to the sibling PR.)

What “Fully Clean” Means

“Fully clean” is the terminal state the ARDI review loop drives toward. A PR/MR is **fully clean** when **both** of these hold:

1. **All CI workflows and check runs are green AND completed.** Every workflow and check run passes — not just the required checks and not just the review job. “Green” means finished with a passing outcome (success or skipped), not merely “currently reporting green while still running” — never treat a workflow or check run that’s still queued or in progress as clean, even if nothing has failed yet. **A reviewer’s posted verdict does not mean the review check has finished, so don’t let a clean verdict stand in for criterion 1 on its own job.** The bot posts its comment and then its run keeps going (bookkeeping steps, a cost tally, the gate job that consumes its result), so a full Ready for merge comment can sit on the PR for minutes while claude-review still reads in_progress and the require-review gate is still queued. Reading the verdict and moving straight to “clean” skips the very state this criterion exists to catch. The gap runs the other way from the stub-review case below: there the check is green and the verdict is missing, here the verdict is real and the check is unfinished. Re-read the check runs after the verdict lands, not just before. (ai-config#712, 2026-07-24: the round-2 verdict posted at 04:06, about two minutes before its own claude-review job completed at 04:06:56 and require-review at 04:07:03.) (The exact field names and casing for these states differ by API surface — REST’s check-runs endpoint returns lowercase status/conclusion strings like completed/success, while gh pr checks/GraphQL’s rollup returns uppercase state values like SUCCESS; don’t hard-code one casing when scripting a check.) **A raw Actions workflow run and a check run are not the same thing, and the usual lookups (gh pr checks, get_check_runs) only cover check runs (plus legacy commit statuses) — not every workflow run necessarily produces one.** A workflow run that’s blocked on action_required (e.g. pending manual approval) before any job starts can complete with zero jobs and consequently zero check runs, making it invisible to a check-runs-only poll. This normally doesn’t affect mergeability (GitHub’s branch-protection required-checks gate operates on checks, not raw workflow runs, so a check-run-less run can’t be wired as required), but if something about a PR’s CI state looks off despite gh pr checks reporting all-clear, cross-check the raw workflow runs before trusting the checks-only view. **gh run list --commit <head-sha> is not a reliable substitute for this cross-check on its own:** it returns every attempt for that SHA (including superseded/cancelled re-runs, so an old failed attempt can look like an outstanding blocker), and a run triggered by issue_comment or a workflow_dispatch invoked without an explicit ref can be recorded against the default branch’s SHA rather than the PR’s head SHA and be missed by a --commit filter entirely. Neither --commit nor --branch is fully reliable for this, because GitHub itself does not record a reliable PR linkage for these trigger types: an issue_comment-triggered run on this very PR (#635, run 29967418653) recorded head_branch: main and an empty pull_requests array via the raw REST API (GET /repos/{owner}/{repo}/actions/runs/{id}) — verified directly, not assumed — so no single filter (commit, branch, or the API’s own PR-linkage field) reliably narrows these runs to the ones for this PR. Treat this cross-check as best-effort: gh run list -R <repo> --workflow <name> (unfiltered, or windowed by approximate timestamp) and eyeball for anomalies near when the PR activity happened, rather than trusting

any one filtered command to be exhaustive. This includes non-gating checks like the Coverage / codecov job: don't merge around a red Coverage run just because it isn't a required check, unless there's a specific, stated reason for that merge (the project wants to maintain decent coverage, so a red Coverage job is a real signal to fix, not to ignore). **codecov/patch is a separate check from the repo's own Coverage workflow job, and both must be green.** The Coverage job runs the coverage-instrumented test suite; codecov/patch is the Codecov service's own status check, gating the PR's DIFF against a minimum patch-coverage percentage — a repo can have a fully green Coverage job while codecov/patch still fails (uncovered new lines in the diff). When delegating implement-a-PR work to a subagent, name this check explicitly in the brief (“ensure codecov/patch passes, not just the test suite”) — a subagent that only runs the local test suite and checks it's green has no way to know it also needs to check a service-side status check unless told. **The set of checks is not fixed while the run proceeds, and a check run's name is not unique — so “the ones I was waiting on went green” is not the same statement as “every check is green”.** A job can spawn further jobs when it completes, so the total grows *after* you started watching. Nothing announces that, and the natural mental model is a fixed list draining toward zero, which makes the growth invisible precisely when you are closest to declaring ready. Re-fetch the whole list each time and re-count it, rather than checking off the names you remember.

The name collision is the sharper half, because it turns a careless check into a confidently wrong one. Two check runs can carry the *same name* on the same head — an earlier one that already succeeded, and a later one still running — so matching on the name returns the stale green and reports the PR ready while the other is still going. They are usually not re-runs of each other: the common case is two separate workflow runs that each happen to define a job by that name, so neither replaces the other and both are legitimately present. Key on the check run's **id**, and read **status** before **conclusion**, since a run still **in_progress** has no **conclusion** to be misled by.

(ucdavis/bcs#458, 2026-07-29: a check-in found the three jobs it was watching all green and would have called the PR clean, except the count had gone from 17 to 20 — **update-snapshots** had finished and spawned a cross-platform R CMD check matrix. Two of those three were still running, and one was a *second* check run named **ubuntu-latest (release)**, alongside the original that had succeeded 14 minutes earlier.)

status itself can be stale, so never infer a job's *duration* from it. Reading **status** before **conclusion** is right, and it invites a second inference that is not: that a run still showing **in_progress** is still running, and therefore that the time since **started_at** is how long it has been going. The field lags. A job can read **in_progress** for minutes after it has actually finished, so “started at T, still **in_progress** now” measures the API's freshness rather than the job's runtime.

That is harmless while you are only waiting for a job to end, which is the usual reason to read the field — the lag costs a poll. It inverts the answer whenever **duration** is **itself the diagnostic**. A reviewer job that dies on a bad credential and one that

genuinely reviews a diff differ mainly in how long they take, so a stale `in_progress` is indistinguishable from exactly the recovery you are watching for, and it arrives as good news.

Take duration from the log’s own timestamps — first line to `Cleaning up orphan processes` — or from `completed_at` minus `started_at` once the run really is complete. Both are facts about the job; `status` at any given moment is a fact about the API.

- **Do:** read elapsed time from log timestamps whenever the length of a run is the thing being judged.
- **Don’t:** conclude a job is still running, or has passed some duration threshold, from `in_progress` plus the wall clock.

(d-morrison/altdoc#96, 2026-07-30: `claude-review` had failed six times in ~26 seconds each, the signature of the model call failing at auth. A re-run was polled twice, three minutes apart, and read `in_progress` both times — reported as “the reviewer has recovered”, and acted on by firing a second re-run on the sibling PR. The log showed that job starting at 04:05:25 and cleaning up at 04:05:51: 26 seconds, identical to the other six.)

2. **The latest review is totally clean:** no nits, and every item that wasn’t directly **Addressed** is either **Deferred** to a tracked follow-up issue, or **Rebutted with a rebuttal that actually convinced the reviewer** — i.e. the reviewer did *not* re-raise it on the next round. A rebuttal the reviewer still disputes does **not** count as clean. That review must be a genuine posted verdict at the current head commit, from an external reviewer if one is reachable — self-review is a fallback for when no working external reviewer is available, never a substitute once one is (see the `ardi` skill’s step 2 for the availability-recheck procedure). Re-check availability right before declaring clean, not just at whichever round self-review first started; an inferred “probably clean” from green CI and resolved threads does not satisfy this.

Criterion 2’s test is the absence of findings, not the presence of a verdict line saying so. A reviewer routinely asserts both at once: a `### Verdict` reading **Ready for merge**, and directly beneath it a findings section listing items nobody has addressed. Neither half is wrong, which is what separates this from the eight numbered cases below — those are all a reviewer producing an unreliable or absent signal, whereas here the comment is accurate throughout and the defect is in the reading. The verdict line answers a narrower question than the one criterion 2 asks, and it is the part that appears first and gets quoted into a status report.

So when the two disagree inside one comment, **the findings win**. Read to the end of the comment before calling anything clean, and count the items under every heading, whatever that heading is called — `address-every-comment` already establishes that “non-blocking”, “nit”, “minor”, and “optional” are prioritization labels rather than a pass, and a reviewer files findings under exactly those words in the section that contradicts its own verdict line.

The disagreement is measurable, and it is not a wording problem. Across 38 verdict-bearing `claude-review` comments sampled from 16 PRs, 8 (21%) carried a verdict line that disagreed with the findings in the same comment. Six read a pass over unaddressed nits, and two ran the other way, blocking over findings the reviewer itself called non-blocking, so the error is not a consistent bias that an offset could correct. The vocabulary is nearly closed by contrast: five outcome lexemes across five markup carriers, with 37 of the 38 naming “Verdict” somewhere. So neither detection nor parsing is the weak link.

That is the argument against gating on a machine-readable verdict field. Adding one would encode the reviewer’s own looser threshold, making roughly one review in five confidently wrong in exactly the form that invites automation. Structured review output should carry **finding counts**, which are checkable against the inline-comment and thread lists, rather than a pass/fail mood, which is checkable against nothing.

(Sampled 2026-07-31 across 12 `Morrison-Lab/ai-config` PRs and 4 in `Morrison-Lab/gha`. All eight are named, so the rate is reconstructable. The six pass-over-nits cases were `ai-config` #955, #941, #939, #935, #934, and #925 — #934’s verdict line *is* the hedge, reading ****One finding (nit), otherwise ready for merge.****. Both opposite-direction cases were on `gha#371`, which returned ****Needs minor changes**** over “two non-blocking, fact/scope findings”. Four other comments within the same 38 are counted out as unclear for a different reason: three of them passes — two with the verdict restated or hyperlinked from an earlier review, one whose findings never reached the PR — plus one where **Needs work** did double duty as both the verdict and an inline finding’s heading. Those are legibility problems rather than disagreements, and they sit inside the pass and non-pass groups rather than beside them. Of the 24 passes, 23 used **Ready for merge** and one used the hedged variant above; the four non-pass lexemes were **Needs more work**, **Needs minor changes**, **Needs work**, and **Needs one fix**.)

A reviewer’s own verification block can be wrong while its verdict is right. The verdict-versus-findings test above needs a disagreement to spot. This one offers none: the verdict is right, the findings section is empty and correctly so, and the defect sits in the arithmetic the reviewer posts to show its work.

A block labelled “verification” is the part of a review *least* likely to be re-checked, because it presents as the checking already having been done. That is what makes a wrong one worse than no block at all. It will usually sum, too, since a balancing partition is what the reviewer was aiming for, so only the composition is wrong.

Re-derive the groups rather than the total. Arriving at the right number says nothing about which groups the parts came from, and that is exactly the error a table that balances conceals.

Read it as the mirror of `ardi`’s “A systematic audit done by skimming is worse than the one-at-a-time version it replaces”. That entry governs an audit *you* produce; this one governs an audit arriving *as evidence*.

The secondary signal is worth acting on rather than merely noting. A reviewer’s reconstruction error usually traces to something genuinely ambiguous in the diff, so treat it as evidence about your own prose and not only about the reviewer.

- **Do:** re-derive a posted verification’s groups, not just its total.
- **Do:** fix the wording that invited a wrong reconstruction, even when nothing in the diff was false.
- **Don’t:** let the word “verification” stand in for having verified.
- **Don’t:** read a table that sums as one that partitions correctly.

(Morrison-Lab/ai-config#957 round 2, 2026-07-31: `claude-review` returned **Ready for merge** with no findings, above a table partitioning the same 38 comments as 24 passes + 10 blocking + 4 unclear. That sums, and the composition is wrong: the sample is 24 passes and 14 non-passes, with the four counted-out comments sitting inside those groups (three passes, one non-pass) rather than beside them. It balanced only because the four were subtracted from the wrong group. Nothing in the diff was false, but “Four further comments” read as a disjoint third bucket, which is how a careful reader reached the wrong partition.)

What “an approving review” means here is not a review state. Across the 25 most recent merged PRs, all 106 posted reviews are `COMMENTED` and none is `APPROVED` — `d-morrison`’s own included, so this is not a bot limitation:

```
gh api graphql -f query='{search(query:"repo:Morrison-Lab/ai-config is:pr is:merged", type:ISSUE, first:100)}' --jq '[.data.search.nodes[].reviews.nodes[].state] | group_by(.) | map({state: .[0], n: length})'
#=> [{"n":106,"state":"COMMENTED"}]
```

The key order there is not a typo: `gh api --jq` marshals through Go and sorts keys alphabetically, so `n` precedes `state` even though the expression builds `state` first. Plain `jq` would preserve the insertion order and print `{"state":..., "n":...}`.

A constant carries no information, so `.state` cannot confirm clean here, and waiting for a formal `APPROVED` would stall every PR indefinitely. Approval is established instead by the two reads criterion 2 and the **Threads** paragraph already name: zero findings in the latest review body, and zero unresolved inline threads. The one state that does still carry information is `CHANGES_REQUESTED`, which stays blocking however a later verdict line reads.

- **Do:** read the whole review comment and count findings under every heading before calling a PR clean.
- **Do:** establish approval from the findings and thread lists, since `.state` is `COMMENTED` on every review this repo receives.
- **Don’t:** quote a **Ready for merge** line as the clean signal while the same comment lists findings.
- **Don’t:** wait for a formal `APPROVED` review, or read `COMMENTED` as a defect in the reviewer.

(Morrison-Lab/ai-config#900, 2026-07-30: the verdict read “**Ready for merge**. No hallucinations, fabricated references, or factual errors found”, immediately above a “Findings (all nits, non-blocking)” section naming three inline comments and closing “None of these affect correctness or usability of the guidance”. All three threads were unresolved at that point, so the PR failed both halves of criterion 2 while carrying a verdict line that read like a pass.)

Findings hide on four different surfaces, and no single check sees all four — so read the verdict body, the inline comments, the thread list, and the verdict’s own conclusion every round. The entry above is about a reviewer contradicting itself inside one comment. This is about the *detection method* returning an answer that is technically true and substantively wrong, which is harder to notice because nothing looks inconsistent.

- **An out-of-diff finding never becomes a thread.** A finding about a line the diff did not touch cannot be attached as an inline comment, so it appears only in the body — reviewers say so explicitly (“inline comments were unavailable for out-of-diff lines”). A thread count therefore cannot see it. Zero unresolved threads is not evidence of zero findings.
- **An empty body hides the mirror case.** A review can post a completely empty top-level body and carry its entire finding in one inline comment, so a body-only read finds nothing to act on and concludes there is nothing.
- **“No verdict” is its own state, distinct from “a verdict with no findings”.** A review job can fail having posted *nothing* — not a stub, not an empty comment. Zero findings and zero review are indistinguishable by any count, and they call for opposite responses: one is done, the other needs a self-review and a re-run. Read the job’s step outcomes when a review is missing rather than inferring from the absence of comments.

The reason this defeats otherwise-good instruments is that each check answers a narrower question than the one being asked. “Are all threads resolved” is not “are there no findings”, and neither is “does the verdict say ready”. Per [algorithmatize-checks](#), prefer the instrument that decides the question exactly — and where none does, as here, say so rather than substituting the nearest available count.

- **Do:** read all four surfaces before calling a PR clean, every round.
- **Do:** distinguish “no findings” from “no verdict” explicitly, and treat the latter as unreviewed.
- **Don’t:** report clean on a zero thread count, however many checks are green.
- **Don’t:** treat an empty review body as an all-clear without checking the inline comments.
- **Don’t:** read a reviewer’s silence as a verdict — a job that posted nothing leaves the same zero counts as a job that found nothing.

A fourth surface, and the one that defeats the gate itself: the review check can pass on a blocking verdict. The three above are cases where a *reader* looks at the wrong place. This is the case where the repo’s own gate looks at the right place and still reports

green, because **require-review** tests whether a review **ran**, not what it **concluded**. So a “Needs more work” verdict and a “Ready for merge” verdict produce an identical check row.

It compounds with case 1 rather than sitting beside it. A review invoked without a **--comment** argument reports its findings in the run’s own comment and posts nothing as a thread — and the better reviewers say so in their last line, which is the tell worth grepping for. The result is a PR with every check green, zero inline comments, zero unresolved threads, and a blocking correctness finding sitting in plain text that no count reaches.

This is the third numbered case below – a check that cannot fail on its own content, so its green carries no signal – arriving on the one job whose whole purpose is to gate on review outcome. The difference is what makes it worse than the benchmark check recorded there. That one is *designed* never to block, and a reader who knows the design knows to read its comment. **require-review** is designed to block, is frequently a required check, and still reports green on a verdict that says the opposite. Read the verdict line itself, every round; a green **require-review** is evidence a reviewer spoke, and nothing more.

- **Do:** grep the verdict body for its own conclusion, and treat a **require-review** pass as orthogonal to whether the PR is clean.
- **Don’t:** let a green review-gate check stand in for reading what the review said.

(Morrison-Lab/ai-config#921, ucdavis/bcs#477, ucdavis/bcs#473, all 2026-07-30, within hours of each other. On #921 every mechanical check passed — all CI green, zero unresolved threads, verdict line reading “Ready for merge” — and the PR was reported clean twice while carrying an open out-of-diff finding. On #477 the review body was empty and the finding was inline-only. On #473 **claude-review** failed after its built-in retry, posting nothing at all, so there was no body to read past and zero threads because zero comments.)

(The fourth surface, ucdavis/bcs#468, same night. **require-review** passed, **claude-review** passed, all 18 checks were green, and there were zero inline comments and zero threads — while the verdict read “Needs more work” over a blocking finding that the new section’s own safety rule was false on one code path. The review’s own closing line said as much, noting that because no **--comment** argument was passed, it had not posted the findings to the PR. Every count-based check called that PR ready.)

A review comment’s header SHA can be stale, so take the reviewed commit from the run’s own head_sha. Criterion 2 requires the verdict to sit at the current head, and the obvious instrument for checking that is the unreliable one: the commit named in the comment’s own caption. A verdict captioned with a superseded commit can be a current-head review whose caption simply names a different commit than the run checked out.

The failure direction is the expensive one. It reads as a stale review, which invites a needless re-trigger, and under **concurrency: cancel-in-progress** that re-trigger cancels the run already in flight at the real head. So the caption costs you the verdict it was making you doubt.

The run's `head_sha` settles it, and the comment links the run it came from, so the check is one call. This is [algorithmatize-checks](#) applied to a verdict: prefer the API field over the prose caption.

- **Do:** follow the job link in the comment and read that run's `head_sha`.
- **Don't:** treat the SHA in a comment's heading as the commit reviewed.

(Morrison-Lab/ai-config#957, 2026-07-31: the `Ready for merge` comment is captioned “Review of `de72464`” while the run it links, 30614782680, records `head_sha: c8d5d8a` — the PR's head at the time, since a `main` merge had superseded `de72464` 64 seconds earlier. Both facts came from `get_workflow_run`; the caption was never rewritten, and the cancelled prior run 30614715159 is the one that actually ran at `de72464`.)

A clean CI run and a clean review verdict are a snapshot, not a standing guarantee of mergeability. `main` can advance after your last check — including gaining its own independent addition that collides with yours (see `sync-with-main.md`'s “two PRs append the same numbered subsection” case) — so re-verify the branch still merges cleanly against current `main` before reporting a PR ready, not just trust the last green run.

Re-check version parity in that same sweep, not only conflict-freedom. [sync-with-main](#) already covers comparing DESCRIPTION versions *after merging main in*. The case that rule misses is the one with no merge at all: `main` advances on its own after your last review round and lands on the branch's exact version, so an R package's `version-check` job (which requires the branch to *exceed main*) goes from green to red with nothing to point at. There is no conflict, no failing check yet, and no warning — the last run passed because `main` was still a version behind when it ran. So the declare-ready sweep needs both `git merge-tree` for conflicts and a direct version comparison; either one alone reports a PR ready that isn't. (UCD-SERG/serocalculator#392, 2026-07-25: the final pre-declaration check found `main` had reached 1.4.1.9016, exactly the branch's version, minutes after a clean `Ready for merge` verdict on an otherwise all-green head.)

Threads: at fully-clean, every **inline** review thread is resolved, and the only conversation left open is the final all-clear exchange — the reviewer's all-clear comment and your reply to it. (The all-clear is usually a top-level PR comment, not an inline thread.) Check this mechanically rather than from a memory of which threads you replied to. Which field name to look for depends on the surface: the GitHub MCP tool `pull_request_read` `get_review_comments` returns thread objects under a `review_threads` key with snake_case `is_resolved/is_outdated`, while a raw `gh api graphql` `reviewThreads` query — what [resolve-pr-threads](#), `pr-status`, and `ard` use — returns camelCase `isResolved/isOutdated`. Both are correct on their own surface; this is the same REST-vs-GraphQL casing split the check-state paragraph above already warns about, so read the response you actually get rather than assuming one spelling. Either way, sweeping for the unresolved ones is the entire check. An **outdated** thread (`is_outdated: true` — the code it anchored to has since changed) still counts as unresolved: addressing a finding and resolving its thread are separate actions, and only the second clears this criterion.

An addressed-but-unresolved thread reads as outstanding work to every later reviewer, which is exactly what this criterion exists to prevent.

One finding can own two threads, so sweep by thread id rather than by finding.

When a reviewer re-raises an item you already answered, the re-raise often opens a **new** thread instead of continuing the original — same file, same line, same finding, different `threadId`. Resolving the one you remember replying to therefore leaves a second thread behind, and it is easy to miss twice over: it is usually marked `is_outdated: true` (the line it anchored to has since changed), and your own memory of the exchange says the item was settled. Neither of those clears it. Re-read the thread list before declaring clean and resolve every entry whose `is_resolved` is false, whatever you recall about the finding it carries; reply on the second thread too, pointing at the first, so a reader landing on either one sees the resolution. (d-morrison/altdoc#61, 2026-07-25: the round-4 re-raise of an unused fixture parameter opened `PRRT_...TypeQ` alongside the original `PRRT_...TypeRc`; resolving the original left the re-raise outstanding, caught only by a mechanical sweep of all seven threads.)

Deadlock -> escalate to a human. If you and the reviewer(s) can't reach consensus on an item (a rebuttal was exchanged and neither side is budging), don't loop forever and don't unilaterally override the reviewer — request a **human reviewer**, @-mention them in a comment summarizing the impasse, and surface the open item.

An automated reviewer's verdict on a disputed factual/technical claim is not stable across independent runs, even with identical evidence available each time. Don't treat one round's "settled, no need to keep arguing" as durable: the very same review job, re-triggered later with no new code changes, can re-raise a claim it previously retracted — and then retract it again on a subsequent run — purely from re-deriving the question differently each time, not from anything changing in the PR. This means a rebuttal thread's outcome (however many rounds of citations and counter-citations) doesn't itself resolve a genuine deadlock the way a human's decision does; only escalating per the bullet above actually settles it. The one thing that DOES help going forward: fold the authoritative citation/evidence directly into the code or doc being reviewed (a comment, not just a PR conversation reply) — a fresh reviewer run re-deriving the claim from scratch is more likely to find the citation sitting right next to what it's evaluating than to dig through prior thread history for it, though even that is not a guarantee against a bot that ignores context already in front of it. (Sparta#852, 2026-07-14: the same @claude review job's independent runs on this PR gave three different verdicts on the identical `gitglossary(7)`-backed `pathspec` claim across three re-triggers with no intervening code change to the claim itself — "settled, accurate" -> "backwards, needs more work" -> "accurate after all, retracting my own prior finding" — resolved only once the human merged it directly rather than by winning the argument with the bot.)

A review job's pass/fail conclusion can diverge from whether a genuine clean verdict was actually posted — check both directions, not just the check's color. The familiar direction: a green review job that posted only a stub with no verdict (a stalled/crashed review run) is NOT a clean verdict — re-trigger and read the actual comment before trusting green. The inverse, easy to miss: a review job reporting FAILURE can still have posted a

complete, genuine “Ready for merge” verdict with real findings-review content — some guard scripts that gate the job’s own pass/fail on detecting a verdict string can misfire and report failure even though a full review ran and passed. Read the posted comment body, not just the check conclusion, before concluding a PR is or isn’t clean. If the check is a **required** check and you’ve independently confirmed the posted content is genuinely clean, that is still not authorization to merge past it yourself — a required check failing is exactly the “stop and ask” case even under a merge-when-confident grant (see `mwc`’s scope note); report the evidence and let the human decide whether to override, fix the guard script, or relax branch protection. (Learned on `sparta#590/#594/#598`, 2026-07-02: two independent PRs hit the inverse misfire in the same session, and an attempt to merge past the required check on verified-clean content was correctly blocked by the harness’s own permission system.)

That inverse has a second mechanism, and under this one the guard is not misfiring at all. The `sparta` case above is a guard that reads the transcript for a verdict and gets the answer wrong, so the red is a defect, and finding the defect explains everything. A guard can also fail a run **before** it ever asks about a verdict, and then nothing is malfunctioning.

`Morrison-Lab/gha`’s review guard is built that way, in both shipped versions. It reads the run’s result object and, on `is_error == "true"`, prints `Claude review ended in an error state` and exits 1, with a single carve-out for the quota case (`total_cost 0` at `num_turns 1`). At `@v1` the step is inline in `claude-code-review.yml` and contains no verdict test whatsoever. At `@v2` the logic moved into `check-review-execution.sh`, which does scan for a verdict, but every line of that scan sits below the `is_error` branch, on the `is_error: false` path. So an errored run is failed without either version asking whether a review was posted, and the version that ran is not the variable.

The claim the guard makes is therefore true, and narrower than it looks: the run ended in an error state. The false step is the reader’s. A review reaches the PR through tool calls **as it works**, so its comment can be complete minutes before the run’s own result object reports a failure. Read `is_error` as a fact about how a run **ended**, never about what it **accomplished**.

That is what lets this survive a long investigation rather than a careless one. Every reading of the run data is correct, and the natural check against the case above, whether the guard misfired, comes back **no**, which reads as confirmation that the red can be trusted. Nothing in the run’s conclusion, result object, or step list differs between “the reviewer produced nothing” and “the reviewer produced a full verdict and then errored”, because the verdict is an artifact on the PR rather than a field of the run. Be precise about how far that goes: at `@v2` the execution output does carry the posted text, so the fact is present in the run and merely unreachable — the guard exits above the scan that would read it — while at `@v1` there is no such scan to reach. Gated by control flow rather than absent, and on neither version does any path the guard takes evaluate it.

So read the guard’s own failure branch once, and let it tell you what its red is worth. A branch that exits before evaluating the artifact yields a red carrying no information about that artifact,

which is the mirror of the benchmark-check case this file records, whose green carries none about its content. Two timestamps then localize it exactly, per `algorithmatize-checks`: bracket the verdict comment's `updated_at` inside the `review` step's own `started_at` and `completed_at`. A comment written while that step was running is a comment that run produced. Comparing it against the guard step's `started_at` instead does not discriminate, since the guard always runs after the review step in the same job, so a stale comment from any earlier round clears that bar just as easily.

- **Do:** read the PR's own comments before accepting that a failed review run produced no verdict.
- **Do:** read a guard's failure branch to learn whether its red is evidence about the artifact at all, rather than keeping “go look” as a habit to remember.
- **Don't:** infer that a run produced nothing from a true report that it ended in an error.
- **Don't:** treat “the guard did not misfire” as establishing that its red is informative.

(Morrison-Lab/ai-config#984, 2026-07-31, job 91208954246: Run Claude Code Review concluded **success** over 16:17:16Z to 16:28:13Z; the review's comment carries a `### Verdict` heading reading **Ready for merge** above a substantive findings review, last updated 16:28:12Z; **Fail the check if the review did not complete** then failed at 16:28:13Z, and the job went red. A session spent several hours on that failure and its siblings, asserting throughout that the runs had produced no verdict, without reading either PR. The unsuffixed step name is the `@v1` tell, and `@v1` is the version with no verdict test at all, which is why the mechanism had to be read at both tags rather than at the one the extracted script lives in.)

A third case, distinct from either misfire above: some checks are designed to NEVER fail regardless of their own posted content, so their green color carries zero signal at all. A CI-runner-relative benchmark check that gates a soft threshold (e.g. “regressed beyond 20% vs. baseline”) may deliberately report success/pass at the GitHub-check level even when it posts a `:warning:` regression comment, precisely because the project has decided that threshold is “a human call, not an auto-block” rather than a hard gate. `gh pr checks` (or the equivalent status API) showing this check as PASS is consequently not evidence there is nothing to look at — it only means the check ran, not that its content was clean. Read the check's own posted comment body every time, the same discipline the review-job case above already demands, but don't expect the check's pass/fail conclusion to ever flip for this class of check even on a real, large regression. (Sparta#995/#998/#999, 2026-07-19: `gh pr checks` reported **benchmark** as PASS across three separate PRs while the actual posted comment showed regressions of 45%, 38.8%, and 36.9% respectively against the CI-runner baseline — two were real, fixable redundant-computation bugs; the third traced to a stale baseline that predated an earlier PR's own accepted cost increase and hadn't been refreshed yet, since the refresh workflow only runs on a weekly schedule, not on every main push.)

A fourth case: a review job can post a syntactically valid, confidently stated verdict that is nonetheless invalid because it rests on a hallucinated premise about the

PR’s own state — not a stub (no verdict) and not a misfire (guard-script/check-conclusion mismatch), but a fabricated fact baked into an otherwise well-formed review. A reviewer that infers PR state from a commit message rather than querying the PR’s actual `state/merged` API fields can mistake a routine `Merge remote-tracking branch 'origin/main'` into `<PR-branch>` commit — pushed to resolve a sync conflict on the still-open PR branch itself — for evidence the *PR* was merged into `main`, and confidently report “PR is closed, no action taken” while never actually reviewing the diff. This reads exactly like a legitimate all-clear (a `### Verdict` section is present, the job reports success), so the stub-detection guards described in CLAUDE.md’s “Do the review yourself when the (claude?) workflow doesn’t produce a verdict” section don’t catch it. Sanity-check any surprising verdict — especially “nothing to review” or “already merged/closed” — against the PR’s real API state before trusting it, and re-trigger for a genuine review rather than accepting a verdict-shaped comment built on a false premise. (gha#293/gha#295, 2026-07-24: after a merge-conflict-resolution push, the re-triggered `claude-code-review` run reported “The PR is closed — it was merged as commit `db11634`” even though the PR was still open and `db11634` was only the PR branch’s own merge-with-main commit; re-triggering once more produced a genuine review of the actual diff.)

A fifth case, and the one that decides what “reachable” means in criterion 2 above: an external reviewer can decline to review at all, posting a refusal in the shape of a review. Unlike the four cases above — all of which are a review that ran and produced something misleading — this is a reviewer that never ran, and says so in a `COMMENTED` review whose whole body is the refusal (e.g. Copilot’s “*unable to review this pull request because the user who requested the review has reached their quota limit*”). Three consequences for driving a PR to fully clean:

- **A refusing reviewer is not “reachable,”** so criterion 2’s external-verdict requirement falls to whichever external reviewer *is* working. Don’t stall a PR waiting for a reviewer that is refusing — but don’t quietly downgrade to self-review either while another external reviewer is answering normally.
- **Reviewers fail independently.** One can be quota-dead while another reviews the same head normally, so check each one rather than generalizing from the first refusal.
- **Keep re-requesting each round anyway.** A quota resets on its own schedule, so a reviewer that refused a few pushes ago can come back mid-session — which is exactly what criterion 2’s “re-check availability right before declaring clean” is for. Say so explicitly when reporting a PR ready: name which reviewer’s verdict the clean call rests on, and which one never weighed in at this head.

The mechanics of detecting a refusal (it arrives as a posted review, not an API error, so the request call’s success proves nothing) are in [memories/github-mcp-tools.md](#). (ucdavis/rampp#111, 2026-07-24/25: Copilot refused three times across two heads for quota while `claude-review` posted genuine verdicts at both; the PR was reported clean — and merged — on `claude-review`’s verdict, with Copilot’s absence stated in the ready-for-merge comment rather than papered over.)

A sixth case runs the other way from all five above: the review is genuine and complete, but the workflow posts the reviewer’s own tool invocation instead of the review body. The comment opens with a literal `gh pr comment <N> --repo <owner>/<repo> --body "$(cat <<'EOF' and closes with EOF\n", wrapping a real, correct verdict as un-rendered text — the model emitted a shell command as its final response and the workflow posted that string verbatim. Nothing is lost, and the same body usually also lands as a properly-rendered sibling comment, so the PR carries the review twice. Two reasons not to shrug at it: a comment opening with a raw gh invocation reads as a broken run, so a human is likely to discount a review that actually passed; and a verdict-detecting guard script (check-review-execution.sh) is now matching against a shell command rather than prose, which can misfire into a needless stub-retry and a second full review’s cost. Read the body and extract the verdict from inside the heredoc rather than re-triggering. (UCD-SERG/serocalculator#392, 2026-07-25; filed as d-morrison/gha#312, which proposes unwrapping the pattern before posting.)`

A seventh case: a reviewer can post a **BLOCKING** verdict on a false positive that will reproduce on every future round. The six cases above all turn on what a reviewer said about the *code* — or, in the fifth, on its declining to say anything. This one is a policy detector firing on the repo’s own conventions, and it behaves differently from every case above in the way that matters for the loop: **re-triggering cannot clear it**, because it keys on text that is still there and that you are declining to change. A timeout or a quota refusal resolves itself on a re-run; this does not.

The shape is an injection detector reading imperative prose as instructions aimed at the reviewer. That misfires badly on an agent-instruction corpus, where imperative mood is the medium rather than a signal of compromise — the distinction that matters for injection is **provenance**, not grammar. Repo-authored guidance in a PR against that repo is not untrusted input, and a detector that cannot tell the difference will flag most of the corpus.

That reading is not an inference from one misfire. The detector went on to block **this very entry**, citing its “Do not count the re-raise” line, and in the same verdict flagged the PR *description* — text that is not in the repository at all and cannot be a convention, a file, or anything a later reader would see. So the trigger is mood alone, on whatever text is in front of it. Treat a third data point arriving on the write-up of the first two as confirmation rather than as coincidence: it is the cheapest possible demonstration that re-running and rewording both miss the point, since the only rewrite that would satisfy it is one that stops giving instructions — which is the entire function of a `shared/` fragment.

Three consequences:

- **Answer with corpus evidence, not argument.** One command usually settles whether the flagged form is a convention: `grep -l "^## In review" shared/coding/*.md | wc -l` against the directory total. Eight of eighteen is a convention; one of eighteen would be a real finding.

- **Do not count the re-raise against the rebuttal test in criterion 2.** That test assumes a reviewer that can be convinced. Reply once naming the evidence, then hold, per [address-every-comment](#)'s per-item noise rule — and keep processing that reviewer's *other* findings normally.
- **Escalate rather than comply, and say why in the thread.** Complying means either a one-file exception to a convention already merged many times, or a corpus-wide change; both are the human's call. State plainly that the check is red **by decision, not oversight**, so a later reader does not treat it as an unaddressed finding and silently “fix” it.

(Morrison-Lab/ai-config#818, 2026-07-29: Jules returned VERDICT: block for “prompt injection attempt in diff” on a new `shared/coding/` fragment's `## In review` section, then repeated it verbatim at the next head without engaging the rebuttal. Eight of the eighteen existing fragments carry an identically-worded section. `claude-review` returned Ready for merge at the same head. The maintainer's call was to hold; the PR merged with `jules/review` red.)

An eighth case: the reviewer's workflow can fail outright on an upstream failure, so there is no verdict of any kind — and its error message may blame the wrong thing. All seven cases above concern a reviewer that produced *something*: a stub, a misfiled conclusion, a pass that cannot fail, a fabricated premise, a refusal, a wrapped verdict, a false positive. This one produces nothing. The job goes red, no review comment appears, and the check is simply absent as evidence either way.

It matters for the loop because the right response is neither of the two obvious ones. It is not a finding to address, so do not self-review as though the reviewer had spoken. And it is not the fifth case's unreachable reviewer either, so do not write the reviewer off yet: an infra failure is frequently transient, where a quota refusal is not. Retry the failed job once, per this repo's standing flaky-infra rule, and only treat the reviewer as unreachable if it fails again.

Read the log rather than the error message, because the message can name a cause the log rules out. A failure of this shape often surfaces as a credential hint (“check `<SERVICE>_API_KEY` is valid”), which is the most expensive possible wrong diagnosis — it sends you to repo secrets for something that will clear on its own. The log usually settles it: a request that *authenticated*, did work, and then failed on a follow-up call was never an auth failure, whatever the summary says.

Two pieces of evidence beat arguing about it, and both are cheap:

- **Prior successes on the same credential in the same session.** A reviewer that posted verdicts minutes earlier is not using an invalid key.
- **A retry with no code change.** If it passes, the failure was transient by construction. This is the mirror of [ardi](#)'s “a symptom that stops reproducing is a fix having landed” — there, silence after a merge needs the merge ruled out before you may call it flaky; here the retry is a genuine negative control, because nothing changed between the two runs.

Say which of the two you have when reporting it, so a later reader can tell a diagnosed transient from a hopeful one. And state plainly that the posted error text was wrong, since the next person to hit it will read that text first.

- **Do:** retry the failed job once, then read the log for where the request actually broke.
- **Do:** cite prior successes or a no-op retry as the evidence for calling it transient.
- **Don't:** treat a crashed reviewer as either a finding or a refusal.
- **Don't:** act on a credential hint that the same log contradicts.

(Morrison-Lab/ai-config#835, 2026-07-30: `jules/review` failed with a 404 on `GET /v1alpha/sessions/<id>/activities`, reported as “Check JULES_API_KEY is valid”. The log showed the key creating that session and confirming it *ready* 0.2s earlier, so the 404 was a propagation race on the sub-resource, not auth — an invalid key fails at creation with 401/403. Jules had already approved twice on the same key that session, and `rerun_failed_jobs` with no code change returned `approve`.)

A second shape of that failure is cheaper to diagnose, because the reviewer names its own session in the failure comment. `jules/review` can fail with Jules did not return a review within 15 minutes. Session: `<id>`, which is not an API error at all — the request authenticated, created that session, and then never delivered a verdict. The session id is itself the auth-succeeded proof, so this shape needs no log fetch: a credential that cannot authenticate never gets a session to name. Prefer that field to the log whenever it is present, per [algorithmatize-checks](#) — one value in the comment decides the question the log was going to answer.

And when a fix is already queued for the same round, **the push is the retry**, so a separate `rerun_failed_jobs` call is wasted: the push re-triggers every reviewer on the new head anyway. Say which of the two you did, because they are not equally good evidence — a push changes the code, so it demonstrates only that the reviewer works now, rather than being the no-op negative control the bullets above prize.

- **Do:** read the failure comment for a session id or similar work-happened marker before fetching a log.
- **Do:** let a pending push serve as the retry, and label that evidence as weaker than a no-op re-run.
- **Don't:** spend a `rerun_failed_jobs` call on a head you are about to replace.
- **Don't:** report a push-triggered pass as proof the failure was transient by construction.

(Morrison-Lab/gha#374, 2026-07-30: `jules/review` reported “Jules did not return a review”, with the 15-minute timeout and session 4236561570323034536 in its own comment. A review fix was already staged, so the push carried the re-trigger, and Jules returned `VERDICT: approve` on the new head about four minutes later.)

A third shape is the one the retry rule hands off to and then stops short of: the failure that reproduces identically on every attempt. The eighth case tells you to retry

once and call the reviewer unreachable if it fails again, which is right, and it is where that case ends. But “unreachable” covers two situations with different owners and opposite next actions. A service-wide outage clears on its own, so waiting is correct. A credential scoped to this repository or organization never clears by itself, so every further retry is wasted and the real deliverable is an issue naming a human with admin access. Retrying cannot separate them, because both keep failing.

The discriminator is a repository you are not asking about: run the same reviewer on a **different** repo in the same session. A success there proves the service is up, which leaves the failing repo’s own credential as the only remaining explanation. This is an [algorithmatize-checks](#) case rather than a judgment call – two check runs decide it – and a multi-repo session usually has the second one for free.

The inversion is what makes this worth writing down, because it reuses the eighth case’s own evidence and points the opposite way. That case offers “prior successes on the same credential in the same session” as grounds for calling a failure transient. Read it carefully: it holds only when the successes are on the **same repo**. A cross-repo success is a different credential, so treating it as evidence of transience argues for waiting out precisely the failure that will never clear. Same evidence type, opposite conclusion, and only the scope tells them apart.

The duration signature is the corroborating half. A reviewer that authenticates and then works takes minutes; one whose credential is rejected dies in seconds, with `is_error: true`, zero cost, and zero permission denials, because no work ever started. Take those seconds from the completed run’s own `started_at/completed_at` rather than from `status`, per criterion 1 above.

- **Do:** run the same reviewer against another repo in the session before concluding a service is down.
- **Do:** stop retrying and file an issue naming the credential once a cross-repo success has localized the failure.
- **Don’t:** read a success on a different repo as evidence that this repo’s failure is transient.
- **Don’t:** keep spending retries on a failure whose every attempt dies at the same short duration.

(d-morrison/altdoc#95 / altdoc#96, 2026-07-30: `claude-review` failed seven times across those two PRs – six on #96, one on #95 – each run finishing in the 26-to-35-second band, with `is_error: true`, `total_cost_usd: 0`, and no permission denials. The nearest pair is 38 seconds apart: the run on altdoc#95 failed 04:07:37Z -> 04:08:12Z, and the same reviewer returned a full `Ready for merge` verdict on Morrison-Lab/ai-config#858 over 04:08:50Z -> 04:11:41Z. So the service was fine and the `d-morrison` credential was not, which no number of re-runs would have shown. Tracked in d-morrison/altdoc#99.)

A check-run reading failure is a fact about one *attempt*, not about the whole `run_id` – a later attempt of that same run can still resolve on its own, with nobody having

triggered it. The section above is right that repeated *identical* failures at the same short duration point to a durable cause, and that retrying blindly is wasted motion once that pattern is established. It does not say the reverse: that a run which has failed once, or even twice, is done. GitHub records each `rerun_failed_jobs/rerun_workflow_run` as a new **attempt** of the same `run_id`, and `actions_get`'s `get_workflow_run` exposes that directly via `run_attempt`, `run_started_at`, and `previous_attempt_url`. A check-run's `failure` conclusion describes the attempt it belongs to; it says nothing about whether attempt 3 of the same `run_id` might still post a genuine verdict, from anyone – a scheduled retry, a maintainer's manual re-run, or a mechanism this session never identified.

So don't infer "this `run_id` is exhausted" from a failed attempt, however many rounds have already failed. Read `run_attempt` before writing that off, and treat a later successful attempt as the real, final verdict – not as an anomaly to explain away.

- **Do:** check `run_attempt/run_started_at/previous_attempt_url` on the actual `run_id` before declaring a review permanently stuck, even after more than one failed attempt.
- **Do:** accept a later attempt's genuine verdict as authoritative, without needing to know who or what triggered it.
- **Don't:** assume a `run_id` is done because its most recent check-run you read was `failure` – fetch the run fresh rather than trusting a cached conclusion.
- **Don't:** claim a specific cause (a scheduled retry, an org-level rerun) for an attempt you did not trigger yourself, without evidence naming it.
- **Don't:** trust a contemporaneous explanation for why a prior attempt failed – your own included – without checking it against that attempt's actual job logs.

(Morrison-Lab/gha#390, 2026-07-31: run 30646364412 failed twice – attempts 1 and 2 both stubs (no verdict, low denial count, on both the initial call and its own built-in gha#185 in-job retry), confirmed against attempt 2's own job logs rather than recalled – and was treated as reproducibly stuck, with self-review relied on instead of a further retry. Attempt 3, `run_started_at`: 2026-07-31T23:34:41Z, `previous_attempt_url` pointing at attempt 2, resolved with `conclusion`: `success` and posted a genuine, itemized "Needs more work" verdict – without this session triggering it, and with nothing on the PR explaining who or what did. A same-thread comment offered a different, already-documented explanation for the earlier failures (a downstream guard misreporting failure after a real verdict had posted) – checked against attempt 2's actual job logs and found not to match: both prior attempts genuinely produced no verdict at all, so that explanation was itself an unverified guess, not a checked one.)

That duration signature does not run backwards, and reading it in reverse is how several unrelated bugs get filed as one. The paragraph above offers a short run as **corroboration**, once a credential is already the hypothesis on other grounds. It is not a test that *produces* the hypothesis, and the difference is easy to lose because the sentence reads the same in both directions.

The reason it cannot run backwards is that every failure occurring before the model call takes about the same time. A job that dies at checkout, at the App token exchange, or at authentication has spent its whole life on setup, so 13 seconds and 28 seconds are the same observation. The duration tells you the run stopped early. It says nothing about **which** early step stopped it, and a credential problem is only one of several candidates.

The failure this produces is worse than an ordinary wrong guess, because it **merges** distinct bugs. Reading a cluster of short failures as one credential fault yields a single tidy story covering all of them, and every separate root cause underneath it goes unfixed. Grouping by symptom feels like pattern recognition, which is why nothing about it prompts a second look.

So read each job's own terminal error before naming any cause, and expect short failures sharing a repo and an afternoon to have nothing to do with each other.

- **Do:** open the log and quote the line the job actually died on.
- **Do:** treat a cluster of short failures as several candidate bugs until each one's error says otherwise.
- **Don't:** infer a credential or quota problem from a short duration alone.
- **Don't:** let one explanation absorb every failure that resembles it.

(2026-07-30, auditing which repos held a `CLAUDE_CODE_OAUTH_TOKEN`: three failures in the 13-to-28-second band were reported as one App-permissions problem. They had three unrelated causes. `UCD-SERG/ucd-serg.github.io` run 30529959398 (25s) failed `App token exchange failed: 401 Unauthorized - User does not have write access on this repository`, because the triggering actor was the Copilot coding agent, which is not a collaborator – filed as `ucd-serg.github.io#84`. Run 30509709695 (13s) on the same repo logged `Actor has write access: write` and then failed `Command failed: git fetch origin --depth=20 pull/77/head:main`. `d-morrison/qwt` run 30391041128 (28s) reached the model and returned `is_error:true` after a workflow-modification denial. Only the first was about permissions at all.)

A group established on real discriminating evidence can still admit a case that was never held to it, and widening scope is when that happens. The section above concerns a weak signal, duration, being read as though it produced a hypothesis rather than corroborating one. This one fires later and is narrower. The signature is strong, it was established correctly, and the defect is in what gets added to the group afterward.

The shape is a first pass done properly, followed by an admission done on less. You read two or three failures' own output, find a genuine shared signature, and group them. Then a further case turns up sharing only the **symptom** that made you look at it, which is usually no more than “this check failed today”, plus a plausible shared cause story. It goes into the group without anyone rereading a log.

So the check is one question, asked of every case after the first. Does this match on the **discriminating evidence** that defined the group, or only on the symptom that made me look?

Note that the remedy above does not catch this on its own. It says to open the log and quote the line the job died on, and that is exactly what was done for the cases that formed the group. Applying a standard to the first N cases is what makes the $N+1$ th feel already covered by it.

A cross-repo or cross-project case is the likeliest to be admitted this way, and the one that most needs the bar raised. It arrives feeling like independent corroboration rather than like another instance, so it reads as strengthening the finding rather than extending it. A scope claim is also the most quotable thing you will write about a bug. “This affects two repositories” is what other people act on, and it usually gets published in a tracking issue, where it outlives the session that produced it. Widening scope is therefore the moment to demand the same evidence again, rather than the moment to accept a weaker kind.

Three disconfirming signals are cheap and general enough to look for by name.

- **How far the pipeline got.** A run that produced two attempt artifacts reached a retry path, and a run whose guard rejected it before any retry cannot have produced a second attempt. A structural difference in progress is evidence about which failure this is, independent of any log line.
- **The error text itself.** Two failures printing different messages came from different code paths, and both strings are usually already in front of you.
- **Whether the run produced the artifact the check exists to gate.** For a review job that means asking whether a verdict is on the PR, which is a different question from whether the job went red.

That third one is the last resort and the sharpest, because it is the only one that survives the two above agreeing. Two runs can print the identical error, from the identical code path, at the identical stage, and still be opposite phenomena – one where the reviewer failed, and one where the reviewer succeeded and the guard failed it anyway. Nothing in the run data distinguishes those, because the distinguishing fact is not in the run: it is on the PR.

So when a check’s own output is the only evidence, remember that a check is a claim about an artifact, and go read the artifact.

The cost is not only a mislabelled case. Dropping the second repository also removed the support for a real inference that had been drawn from it, that two repositories sharing one action implies the bug lives in the action. That support was never real, so the false claim cost an inference on top of a case record. The retraction only revealed the loss rather than causing it.

- **Do:** re-read the new case’s own terminal error before adding it to an existing group, however well established that group is.
- **Do:** compare the attempt or artifact count for a structural difference in how far each run got, before treating two failures as the same one.
- **Don’t:** admit a case on a shared symptom plus a shared cause story when every earlier member was admitted on quoted evidence.

- **Don't:** publish a widened scope claim without holding the added case to the standard the original ones met.

(2026-07-31, `claude-review` failures on Morrison-Lab/ai-config #984, #985, and #986: two run results were read directly and shared a genuine signature, `is_error: true` alongside `subtype: "success"` after real work (\$4.10 over 13 turns, \$0.97 over 2 turns). Both of those reads were #986's own two runs, though. PRs #984 and #985 were admitted on nothing but a `claude-review` failure the same day, with no result object read for either – so that grouping was already the pattern this section condemns, one step before the one it was written about. Reading them later made it worse rather than merely unverified: both had posted complete **Ready for merge** verdicts, minutes before their guards failed the check. They were the *opposite* phenomenon – the reviewer succeeded and the check was wrong – filed as instances of the reviewer failing. Neither signal above would have caught it, since their error text and their stage are identical to #986's; only the third one is, and it was added to this list because of them. The duration rule above was invoked explicitly to confirm that #986's 9-minute and 53-second runs were the same bug. Morrison-Lab/gha#390 was then added to the group because its own `claude-review` had failed the same day, and a scope correction widening the finding to two repositories was posted to the tracking issue, Morrison-Lab/gha#391. It was a different bug. That PR's log reads `Attempt 1 produced a stub review (gha#185) and the retry ALSO ended without a verdict with a low denial count`, a path reachable only when `is_error` is false, so the grouped signature is rejected by the guard before any retry can happen. The two attempt artifacts and the differing guard wording were both visible at the time. The claim was retracted on the same issue.)

Address Every In-Scope Review Comment

When iterating on a PR with a reviewer, **address every in-scope flagged item**, regardless of severity label. The reviewer's "Not a blocker", "minor", "nit", "optional", "consider", or "if you want" labels are for prioritization, not a free pass for the implementer.

For each flagged item, do exactly one of:

1. **Fix it in this PR.** The default path — most nits are 1–3 line changes.
2. **Defer.** Only when the fix expands the PR's scope (new feature, broader refactor, separate concern), the requester has explicitly said this PR shouldn't grow, or the flagged content isn't actually yours to fix here (see the `main-sync` case below). File a follow-up issue and reference it in a PR comment so the item isn't lost — except in the `main-sync` case, where the "follow-up" is fixing it on `main` directly, not a new issue.

Then trigger another review and repeat until the PR is **fully clean** — zero flagged items under any heading, no "non-blocking", "harmless", "minor observation", or "could improve" sections. "Looks good" / "no findings" / "approved" with no follow-on bullets is the bar. Resolve every inline review thread along the way, leaving only the final all-clear exchange.

Always resolve an inline thread the moment its comment is successfully addressed — the fix pushed and a reply posted naming it — in the same pass, whatever workflow you’re in: a formal **ard**/**ardi** round, a CI-monitor nudge, or a one-off fix outside any loop. Addressing without resolving leaves a thread that reads as outstanding work to every later reviewer, blocks **fully-clean**’s every-inline-thread-resolved criterion, and drags stale noise into the next review round. The per-disposition settlement rules in **ard** step 4b still govern the exceptions: a **Rebut** stays open until the reviewer drops it, and an **Address** you’re not confident fully settles the concern gets a reply asking for confirmation instead of a resolve. The **resolve-pr-threads** skill sweeps any stragglers, but it’s a backstop — resolve-on-address is the default, not a cleanup step.

Do **not** report “ready to merge with one minor nit noted” / “harmless as-is” / “can address if you want” — that hedging just pushes triage back to the requester. If after 3–4 rounds the reviewer keeps generating new nits each cycle (asymptotic noise), surface that and ask whether to keep going or accept the current state.

Noise is per-item, not per-round — don’t stop the whole loop over one recurring flag. A long-running PR can have both real findings (worth fixing every round) and one specific item the reviewer re-raises verbatim round after round even though it’s already deferred/tracked (e.g. a file-length guideline already split into a follow-up issue). Keep fixing every *new* finding as it appears — don’t let the recurring item make you stop processing genuinely new ones. But stop re-litigating *that one item* every round: reply once pointing at the tracked issue, and hold on it specifically rather than re-deferring it on each pass. Surface the pattern to the user (which item, how many rounds, where it’s tracked) and let them decide whether to resolve it now (e.g. do the split) or leave it as accepted recurring noise — don’t decide unilaterally to either keep re-processing it or silently drop it. (rme#706 ran 100+ review rounds: each round’s *new* findings — a missing derivation step, a missing i.i.d. hypothesis, an unverified citation locator — got fixed every time; the one recurring file-length flag got a single reply-and-hold each round until the user weighed in.)

When a finding is a pattern (a formatting/style rule broken in one spot), apply it everywhere it recurs in the same file, not just the flagged line. A reviewer that flags one inconsistent list-item format is telling you about the rule, not just that one item — fix every occurrence in the same file that breaks it in the same pass, rather than waiting for the reviewer to flag each occurrence in a separate round. Re-scan the whole changed file for the same pattern before pushing the fix.

When a prose fix changes wording that’s also paraphrased elsewhere in the same PR (a CHANGELOG entry, a PR description, a cross-reference), sync that copy too. A CHANGELOG entry written before the review lands often quotes or paraphrases the exact phrase a reviewer later flags; fixing the source prose but leaving the paraphrase stale reintroduces the same wording issue one file over. Grep the diff for the flagged phrase before considering the finding closed. (ai-config#373: fixed “routing/dispatch site” in the skill per review, but the CHANGELOG entry still said it until a follow-up commit.)

The PR description is on that list and is the one copy grepping the diff cannot find, so check it separately. A PR body is not a file, so it appears in no diff and no reviewer reads it as part of the change under review. That makes it the copy most likely to survive a fix, and the copy most likely to be *read* by someone deciding whether to merge — so a stale one teaches the reader exactly the thing the diff was corrected to remove.

The tell is a fix to something the PR body summarizes: a behaviour change, a mechanism, a rationale. Re-read the description against the corrected diff before declaring the round done, and say in the update that it was corrected, so a reader who saw the original knows it was revised rather than always having said this. Where the correction has history worth keeping — a claim that was wrong and is now right — state it as history in the body rather than silently overwriting, since the wrong version is what earlier comments respond to.

- **Do:** re-read the PR description after any Address that changes what the PR does or why, alongside the changelog check above.
- **Don't:** treat a clean **grep** over the diff as evidence every paraphrase is synced — the description was never in it.

(ai-config#829, 2026-07-29: a review nit led to correcting a gha#350 attribution in `memories/github-actions.md`. Both reviewers then approved, and the PR body still carried the original wrong claim verbatim — “**continue-on-error** there, the dropped implicit `success()` here” — because it had been written before the correction and was not part of the diff either reviewer read. Caught only while assembling the ready-for-merge summary.)

Following that “state it as history” advice is what produces the next block, because an automated reviewer reads the body as a flat statement of intent. The paragraph above is right that a correction with history worth keeping should be recorded rather than silently overwritten, since earlier comments respond to the old version. It has a failure mode it does not warn about, and the failure lands precisely on the authors who follow it.

A past-tense paragraph saying a thing *was* excluded is, to a bot, not distinguishable from a claim that it *is* excluded. Tense is doing all the work, and nothing in the reviewer’s reading of the document preserves it. So the more faithfully the reversal is recorded, the more confidently the reviewer reports the diff as contradicting its own description — and the remedy it proposes is to revert the change, which means undoing whatever the reversal was.

Distinguish this from an ordinary stale snapshot before answering. A reviewer that started before your edit never saw the correction and needs only a pointer to it. The timestamp check further down is written about a missed *rebuttal*, but the same **started_at** comparison decides a missed *body edit*: a body corrected after the run began is invisible to it for exactly the same reason, since the whole PR is snapshotted once at run start. This one re-raises *at the corrected text*, so the timestamps clear and the finding still stands. Compare the run’s start time against the edit, then read which passage the new verdict quotes — if it is quoting your history section, this is the case, not that one.

- **Do:** state the current content first, marked as current, before any history.
- **Do:** put the reversal in its own section that opens by saying it is history.
- **Do:** make sure the “what is excluded” section does not name the reversed item at all, in any tense.
- **Don’t:** rely on past tense alone to carry the distinction.
- **Don’t:** revert a maintainer-requested change because a reviewer read the history as current — rebut, and escalate rather than comply.

Be honest about the residual: all of that can be applied and a further run can still block, at which point the only remaining move is deleting the history outright, which costs the earlier comments their referent. That trade belongs to the human, not to the agent driving the PR.

(Morrison-Lab/ai-config#843, 2026-07-30: the maintainer asked for a fourth tool on an allowlist that had shipped with three. Jules blocked twice. The first was an ordinary stale snapshot — its run started at 02:56:28Z, two seconds before the body was corrected. The second re-ran at 03:00:31Z against the corrected body, with a new session id and a different cited line number, read the reversal-history section, and praised the description for “explicitly documenting which permissions should be intentionally excluded” while demanding the requested tool be removed. `claude-review` saw the same inconsistency at the same stale head and graded it a minor, explicitly non-blocking prose note. Escalated; the human merged past it.)

The same sync is needed when the review fix is to CODE BEHAVIOR rather than to wording — and that case is easier to miss, because nothing about fixing a bug points at the changelog. The rule above fires on a recognizable trigger: a reviewer quotes a phrase, so you go looking for that phrase. A behavior finding gives you no phrase to grep. You change the code, update the PR body’s description of what it now does, and the `NEWS.md/CHANGELOG.md` entry — written before the review, in prose that described the *old* behavior correctly — goes on asserting it. Every later round then reviews a diff whose changelog contradicts its own code, and no reviewer flags it, since each file reads plausibly on its own. The shipped result is worse than a stale paraphrase: a user reading the release notes is told the opposite of what the release does. So after any Address that changes behavior, re-read the PR’s changelog entry against the new behavior — not just the code and the PR body. Fold it into the same pre-push self-review pass `ardi` already requires; a changelog entry is a claim about the diff, so `fact-check-prose` applies to it exactly as it applies to any other prose in the PR. (d-morrison/altdoc#78, 2026-07-27: review round 2 established that `mkdocs` serves `/man/foo/`, not `/man/foo.html`; the code and the PR body were corrected that round, while `NEWS.md` kept saying links point at `.html` under “`mkdocs` and `quarto_website`” through two further clean review rounds. Caught by a `main-sync` merge conflict that happened to land in that entry — not by any review, and not by any check.)

Tighter still: a changelog entry can contradict its own commit message, in the same commit, with no review in the loop at all. Both cases above need a review round to set them up — a reviewer quotes a phrase, or a finding changes behaviour — so the trigger

to go looking is external. Here there is none. The commit message and the changelog entry are written minutes apart, by you, in the same commit, and disagree.

The reason it survives is that the two are drafted in different registers. A commit message argues for the change and reaches for the sharpest true statement of the mechanism; a changelog entry describes the change for a release note and reaches for the tidiest one. Nobody reads them side by side afterwards. A diff review sees one, a `git log` sees the other, and no check compares them — so the contradiction ships, and the release notes are the half a user actually reads.

The check is mechanical and belongs in the pre-push self-review pass [ardi](#) already requires: after writing a rationale into a commit message, grep that same commit’s prose changes for a claim about the same mechanism, and read the two together before pushing. Where they differ, the commit message is usually the correct one, because it was written while the mechanism was in front of you.

(ucdavis/bcs#463, 2026-07-30: a0f4113d’s commit body said `update_trigger` “can set `enabled: false` and rewrite a routine’s prompt and cron outright”. The `NEWS.md` entry edited by that same commit justified excluding a different tool on the grounds that the allowed set only changes *when* routines run — which the commit body directly refutes, and which is wrong about `create_trigger` too, since it authors a whole new routine. Caught by `claude-review` as an inline finding, not by any check, and the correct framing turned out to be deferred versus immediate effect.)

One step further back: a figure inherited from the tracking issue is both the copy git keeps and the copy nobody verified. The entry above explains a mismatch by *register* — a commit message argues for the change while a changelog describes it, so they get drafted differently and never read together. Here both claims sit in the same register and describe the same fact. Only one of them was checked. What separates them is **provenance**: a number produced by running something, versus a number carried over from the issue you wrote before you had anything to run.

Two properties make it worse than an ordinary wrong number.

One of the two copies becomes permanent, and you cannot tell which from inside the PR. A PR body stays editable forever, while a commit message does not survive a merge in editable form — but which text a squash merge actually keeps is a repository setting, and it can be either. Configured one way the commit messages land on `main` and the PR body is discarded; configured the other the PR body becomes the commit body and the commit messages are dropped.

That is why the rule is *both must be right* rather than *check the important one*. The copy that survives is chosen by a setting most authors have never looked at, so treating either as the draft is a coin flip. And the odds are not even: the commit message is the one written earliest, from the least evidence, so the configuration that keeps it is the one that makes the weaker copy permanent.

Read a recent squash commit on `main` if you want to know which way a given repo is set — `git log -1 --format=%B <a squash merge>` shows it directly, and beats reasoning about settings pages. (Checked this way on this repo, 2026-07-30: `5670f9f`, the squash of [#855](#), carries that PR’s commit message rather than its body. So here the weaker copy is the one that persists.)

And verifying once feels like verifying. Running the check for the PR body produces a real sense of having established the fact, which is what stops you checking the other place it appears. The verification is genuine; the coverage is not.

Note the shape is the same as [ardi](#)’s “an instruction’s own suggested code is not exempt”, one artifact over: content inherited from a planning document does not feel authored, so the checks you apply to your own claims do not fire on it.

- **Do:** re-run the check when a figure moves from an issue into a commit message, even having verified it once for the PR body.
- **Do:** read `git log -1 --format=%B` before pushing, against the same source the body’s claims came from — a commit message is not greppable from the working tree once written.
- **Don’t:** copy a count, version, or path out of the tracking issue on the strength of having written that issue.
- **Don’t:** treat “permanent in history” as settled while the PR is unmerged — `git commit --amend` still works, and is usually worth a fresh CI round against a wrong figure reaching `main`.

(`ucdavis/bcs#465`, 2026-07-30: a submodule-pin bump whose PR body said `CLAUDE.md` resolves `33 @.ai-config/... imports` and whose commit message, written minutes apart, said `25. 33` came from a script; `25` came from the tracking issue, written from recollection. Review named the mechanism exactly — inherited from the issue rather than re-checked against the file — and graded it non-blocking on the grounds that it was already permanent, which was the one part that was not yet true.)

A corollary for checking any of this in a semantic-line-break corpus: a single-line `grep` returns false negatives on your own prose. The instruments above and elsewhere in this file assume you can search for a phrase you wrote. In a corpus that mandates one clause per line, a phrase of any length routinely spans a newline, so `grep 'flat statement of intent'` reports zero against a file that plainly contains it. The failure direction is the dangerous one: a missing-content check that answers “absent” when the content is present reads as a merge having dropped your work, which invites re-doing something already done. Normalize whitespace before matching — read the file, collapse `\s+` to a single space, then search — rather than trusting a line-oriented tool against deliberately broken lines. (Same day, verifying that [#855](#) had landed: two of three greps reported present and the third reported absent, purely because that phrase happened to straddle a line break.)

Inline markup breaks the same search, and that variant aims the false negative at someone else’s work rather than your own. Whitespace is the obvious thing a line-oriented tool gets wrong, so the fix above normalizes it. Markup is the one nobody normalizes, because the two strings *look* identical: a rule titled `Run a local session in an isolated `git worktree` by DEFAULT` is quoted in a citation as “Run a local session in an isolated git worktree by DEFAULT”, since prose quoting a title drops its code spans. Grep the quoted form and the definition does not match; only the citation does.

The consequence is worse than the line-break case, and in a specific way. There a false negative says your own merged work is missing, so you re-do something already done. Here it says the **cited** thing is missing, which reads as a dangling citation — and the prescribed response to a dangling citation is to file an issue. So the wrong search does not merely waste effort, it puts a false claim about the corpus into the tracker, against a citation that resolves. A result of exactly one hit, in the citing file, is the tell: a genuinely dangling citation and a formatting mismatch produce the same count, and only reading the hit distinguishes them.

Normalize backticks along with whitespace, or search a distinctive unformatted fragment rather than the whole title.

- **Do:** account for inline markup as well as whitespace before concluding a quoted phrase is absent — see the next block for which side to normalize.
- **Do:** read the single hit when a search for a citation’s target returns only the citation itself.
- **Don’t:** file a dangling-citation issue while the only evidence is a literal grep that found nothing but the citation — that is the search failing, until a normalized one agrees.

(Morrison-Lab/ai-config, 2026-07-30: `skills/ums/SKILL.md:109` cites `memories/preferences.md`’s `worktree-by-default` rule, and a literal grep for the quoted title returned only the citation, which was reported as a dangling reference. The rule is at `memories/preferences.md:264`, differing from the quotation by two backticks; a backtick-normalized search found both files.)

Apply whatever normalization you choose to the search term as well as to the text, or the fix produces a third false negative of its own. Both cases above are answered by transforming the haystack — collapse whitespace, strip backticks — and that framing invites transforming only the haystack, since the needle is the string you already know. But a normalizer is a function, and testing `f(text)` against a raw needle compares two different alphabets. Strip `_` to catch **emphasis** and a snake_case identifier stops matching itself: `SH_WORD_SPLIT` becomes `SH WORD SPLIT` in the file while your pattern still carries the underscores, so a term that is present reports absent.

This failure gets *more* likely as the normalizer gets better, which is the part worth naming. Every character class added to catch another markup form is another class that occurs inside real identifiers, so the enumerate-what-to-strip approach converges on breaking the searches it was extended to fix. Enumerating is the wrong shape, not merely an incomplete list.

Running the same function over both sides dissolves the question, whatever the function is:

```
norm = lambda s: re.sub(r"[*_\s]+", " ", s)
norm(needle) in norm(haystack)
```

- **Do:** normalize the needle with the identical function applied to the text, so the comparison is between two transformed strings.
- **Do:** re-test any earlier absent verdict after extending a normalizer, since the extension can break a term the previous version matched.
- **Don't:** enumerate which markup to strip and treat that list as the fix.
- **Don't:** test a raw search term against normalized text, however plain the term looks.

(Morrison-Lab/ai-config, 2026-07-30, verifying #919 on `main`: a probe collapsing backticks, asterisks, underscores, and whitespace in the file alone reported `SH_WORD_SPLIT` ABSENT, while `git grep -c` found it. The same needle normalized reported present. That was the third normalization-caused false negative of the session, and the first produced by the remedy rather than by the raw search.)

A flagged item that came in via a main-sync merge, not your own diff, is still a Defer — just one where the follow-up is fixing it on main directly, not filing a per-PR issue. This is not the ARD skill's "Acknowledge" disposition: `skills/ard/SKILL.md` reserves Acknowledge for praise or a no-ask observation, and explicitly warns against stretching it to dodge a real finding — a redundant config line a reviewer flags is a real finding with an implied fix request, so it needs a real disposition, not a label that means "no change requested." When a reviewer flags something (a redundant config line, a stale pattern) inside a file your branch only touches because you merged `main` in to resolve a conflict, check provenance before fixing it: `git log/git blame` the flagged line, or just compare against `origin/main`'s current content. If it's identical to `main`, "fixing" it on your branch alone doesn't fix anything — it just makes your branch disagree with `main` on unrelated content the next person to touch that file will have to reconcile again. Reply agreeing the finding is correct but out of scope for this PR, and leave it for whoever owns that file's actual content to fix on `main` directly — no follow-up issue needed, since the fix target is `main` itself, not this PR's own change. (UCD-SERG/serocalculator#503: a review flagged `.Rbuildignore`'s `^\.posit/assistant$` as redundant with the existing `^\.posit$` pattern above it — both lines had landed together in an already-merged `main` commit (#579), picked up via a routine `main-sync` merge, not introduced by #503's own diff. Deferred to `main` instead of fixed on the branch.)

This generalizes to a skill's own inline restatement of a fragment it links to. A `SKILL.md` that links a backing `shared/` fragment for the full detail often *also* restates the fragment's approach or word list inline (in its `description` field, or a short procedure-step summary) so a reader doesn't have to open the linked file. Fixing a bug in the fragment doesn't automatically fix these inline restatements — they're a second, independent copy of the same claim, and a review round after the fragment fix can catch them going stale exactly like a `CHANGELOG` paraphrase does. Grep the whole PR diff for the fixed phrase/word-list, not just

the fragment file, before considering a fragment fix complete. (ai-config#507: fixing `forward-references.md`'s regex left `fix-forward-references/SKILL.md`'s own `description` field and Step 2 summary describing the old, already-fixed approach — caught in a second review round.)

A bot that re-raises an item as “not addressed” may simply not have seen your reply — check the timestamps before treating it as an impasse. An automated reviewer gathers the PR's comments once, when its run starts. A rebuttal posted after that snapshot is invisible to it, so the next round reports the item as still open and unaddressed even though a substantive reply is sitting in the thread. The tell is a re-raise that repeats the original finding verbatim and speaks only to whether the *code* changed, without engaging any argument you made. Before escalating, compare your reply's timestamp against the review run's `started_at` (`gh run view <id> --json startedAt`, or the `started_at` field each run carries in `get_check_runs` when `gh` is absent): if the reply landed after the run began, it is a stale re-raise, not a genuine disagreement. Reply once pointing at the earlier rebuttal (link it directly — the next run will see it), and don't count that round toward the rebuttal-didn't-convince-them test in `fully-clean.md`.

The ordering fix is cheap: when a round is Rebut-only, post the rebuttal **before** anything that triggers the next review (a push, an @claude mention), so it is in the snapshot the next run reads. When a round mixes Address and Rebut, post the rebuttals first and push the code second, for the same reason. (d-morrison/altdoc#34: a `\pkg{}` rendering rebuttal carrying a `pandoc` run that disproved the finding's implied hazard was posted about a minute before the follow-up review job started; that review reported the item “wasn't addressed in 9398d5d” and re-posted the identical suggestion.)

Reply-first collides with citing the fix's SHA, and the way out is to commit between them rather than to pick one. The rule above is easy to agree with and still lose, because on a mixed round the reply you want to write says “Addressed in `<sha>`” — and that SHA does not exist until you have committed. So the two instructions read as mutually exclusive: reply first and you have no SHA to cite, push first and the reply misses the next review's snapshot. Pushing first wins that standoff by default, since it is the half that *unblocks* the sentence you were trying to write.

The conflict is only apparent, because committing and pushing are separate steps and only the push triggers review:

1. **Commit** the round's fixes. The SHA now exists and is stable.
2. **Reply** on each thread, citing that SHA.
3. **Push**. The next review's snapshot already contains the replies.

A commit that is never pushed is invisible to CI and to the reviewer, so step 2 is citing something real but not yet reachable — which is fine for a few seconds, and is exactly the window step 3 closes. Note the one thing this does *not* license: the SHA you cite must come

from `git rev-parse HEAD` or `git log`, never from recollection, per the PR-body bullet in [ardi](#).

- **Do:** commit, reply citing the committed SHA, then push — in that order.
- **Don't:** treat “I need the SHA for the reply” as a reason to push before replying; that is the ordering the bullet above exists to prevent.

(ai-config#871, 2026-07-30: a four-finding round with three Addresses and one Rebut was pushed first and replied to about a minute later, so the round-2 review run started before the rebuttal was visible to it. It happened to engage the rebuttal anyway — the evidence was in the diff as well as the thread — but that was luck, not the ordering working.)

A finding can be right while its suggestion block is wrong — verify the suggested literal before applying it. A GitHub ``suggestion block is one-click-appliable, which is exactly what makes an unverified one dangerous: the surrounding prose argues for a change you agree with, so the concrete replacement rides in on that agreement without being checked itself. Treat any file path, version, flag, or command inside a suggestion as a claim to verify, not as text to accept — the same standard [fact-check-prose](#) applies to the diff. Accepting a bad literal is worse than ignoring the finding, because it publishes a specific wrong value under the reviewer's apparent authority. When the suggestion is wrong but its point stands, fix the underlying issue your own way and say in the reply why the suggested form was set aside — silently deviating reads as having missed it. (ai-config#726: a review correctly flagged that a `<path>` placeholder didn't say where a script came from, but suggested `<path-to-gha-checkout>/check-new-line-breaks.py` — one directory level too high, since the composite action's directory and the script inside it share a name. `git ls-files` in the gha checkout settled it in one command. Applying the suggestion verbatim would have documented a nonexistent path in the entry whose whole purpose is getting someone to run that script.)

The same check applies to a fix a reviewer describes in prose rather than in a suggestion block, and the sharpest test is the reviewer's own example. A finding that ships a concrete repro case has handed you a test fixture: run the proposed fix against that very case before adopting it. A reviewer reasoning about a fix in the abstract can propose one that is directionally right and still insufficient — it closes the failure mode they named while leaving the case they cited broken — and adopting it verbatim converts their partial diagnosis into your shipped bug, with the review thread reading as though the item were settled. When the proposed fix falls short, prefer eliminating the failure mode outright over layering another patch onto it, and post the evidence (the fix applied to their example, and what it still produces) rather than just asserting it was insufficient. (gha#318, 2026-07-26: a review correctly found that a heredoc-terminator regex lacked an end-of-line anchor, and suggested adding one. Tested against the reviewer's own indented-EOF example, the suggested anchor still truncated the body, because the terminator's leading `[ \t]*` accepted a space-indented closing line real bash rejects. Matching whole lines against the tag — how bash itself ends a heredoc — removed the whole lazy-quantifier/anchor failure mode instead of narrowing it; the reply carried the failing output of the suggested form.)

A reviewer’s corrected citation is another factual claim, so verify the replacement before adopting it. The finding can be right: the citation in the PR can name the wrong source. That does not make the reviewer’s proposed source right. A replacement issue or PR number is a fresh provenance claim, and it needs the same check as the original citation. For text provenance, prefer history over word association: `git log -S "<exact line>" -- <file>` asks which commit introduced the line, while matching a word in another PR plus a nearby merge time only builds a story. Keep the review’s conclusion when it is right, but set aside the replacement when the evidence points elsewhere, and say which query decided it.

- **Do:** verify a proposed replacement citation with the source’s own history before editing the PR to use it.
- **Do:** use `git log -S "<exact line>" -- <file>` or an equivalent provenance query when the question is which PR introduced text.
- **Don’t:** adopt a reviewer’s corrected issue or PR number because the original was wrong.
- **Don’t:** use word overlap and same-day timing as a substitute for source history.

(Morrison-Lab/ai-config#971 round 2, 2026-08-01: a review correctly found that PR #955 did not cover a “default nobody chose” case record. It then proposed #951 as the source because #951’s `memories/tools.md` entry used the word “default” and merged the same day. That was the wrong default and the wrong file: #951 did not touch `shared/workflow/metacognitive-monitoring.md`. `git log -S "An unexamined default" -- shared/workflow/metacognitive-monitoring.md` identified #947 as the source for the default half, while #955 supplied the handed-premise half in the same fragment.)

The highest-yield version of that check: when a comment names an edge case in its own prose and also supplies a fix, run the fix against that edge case. The bullets above test a suggestion against the code, or against a repro the reviewer provided. This tests it against the reviewer’s *other paragraph*, and it is the cheapest of them, because the hazard has already been identified for you — the work left is only to check whether the proposed code handles it.

Nothing forces the two halves to agree. A comment’s prose and its suggestion are drafted separately, and a reviewer who spots an edge case while reasoning about the problem does not necessarily carry it into the snippet. So a comment can read as unusually thorough — it anticipated a failure mode you had not — while shipping a fix that falls into exactly it. That thoroughness is what makes the suggestion persuasive, which is the trap.

Applying it is worse than ignoring the whole finding. The prose half was right, so the reviewer’s authority is real; the snippet then lands under that authority carrying a defect the same comment already described, and the thread reads as settled. Worse still when the defect is one your own corpus documents, since the review has now talked you out of a standing rule.

Keep the finding and reject the snippet. Fix it your own way, quote the edge case back, and say plainly why the suggested form was set aside — silently deviating from a **suggestion** block

reads as having missed it.

- **Do:** check a suggested fix against every failure mode the same comment names, before checking anything else about it.
- **Do:** name the reviewer’s own caveat in the reply, so the rebuttal rests on their evidence rather than on your say-so.
- **Don’t:** let a comment’s demonstrated thoroughness transfer to its snippet — they are separate claims.
- **Don’t:** discard a finding because its fix is wrong; the half that named the hazard usually still stands.

(Morrison-Lab/ai-config#868, 2026-07-30: a review correctly found that `git merge-base --is-ancestor` prints nothing and answers by exit status, and its second paragraph noted the command exits 2 or higher when the ref has been pruned away. Its suggested ... `&& echo "ancestor" || echo "not ancestor"` maps that exit onto the `not ancestor` branch, since `&&` fails on any non-zero status — so the fix printed a confident verdict for precisely the broken-check case the comment itself had raised, which is the shape `fail-fast` names. A three-arm case `$?` was used instead, reporting 0, 1, and 2+ distinctly.)

A quieter variant: the suggestion introduces no defect at all, it restates the line above it — so applying it deletes coverage while reading as hardening. Every bullet above is about a snippet that would break something, so the reviewer’s authority is the trap and skepticism is the defence. Here nothing breaks. The tests still pass, the diff looks like a robustness improvement, and the comment is *correct about the problem*. What is lost is the only assertion covering a different property, replaced by a second copy of one already present a line earlier — so a two-assertion test becomes a one-assertion test that still looks like two.

The reason it survives review is that the surviving copy passes, which is indistinguishable from the fix working. So the usual after-the-fact check — run the tests — cannot detect it, and neither can CI. It is a `challenge-redundant-content` finding arriving from the reviewer, which is exactly the direction that makes deferring feel appropriate.

Watch for the comment citing the neighbour as *support*: “the check on the line above is already load-bearing for this claim” is the argument against the replacement, and it reads as an argument for it. Same structure as the edge-case bullet above — prose and snippet drafted separately, disagreeing — one artifact over.

Compare a suggested predicate against its **neighbours**, not only against the code it replaces, and evaluate both on real input rather than reasoning about them; one command decides it, per `algorithmatize-checks`. Then prefer removing whatever made the original fragile over swapping one fragile sentinel for another.

- **Do:** evaluate the suggested predicate and its neighbours on real input, and keep the finding while rejecting the snippet when they coincide.

- **Do:** fix the underlying coupling instead, and say in the reply why the suggested form was set aside.
- **Don't:** accept a **suggestion** block that restates an adjacent check — passing tests afterward prove nothing, since the survivor passes for both.
- **Don't:** read a reviewer's own "the line above already covers this" as support for their replacement.

(Morrison-Lab/ai-config#896, 2026-07-30: a review correctly called a test's "user-invocable" `not in body` sentinel fragile, and suggested `not in body.lstrip()[:3]`. Evaluated against the real body, that is the same predicate as the `not body.lstrip().startswith("---")` assertion directly above it — both test the first three characters, both returned `True` — so adopting it would have left one property checked twice and the other not at all. A synthetic fixture replaced the corpus coupling instead, plus a third assertion that body prose *survives* stripping, which neither the original nor the suggestion covered. Tracked as #905.)

A finding can be right, and its fix adequate, while the *reason* it supplies is too weak to ship — and in a corpus of rules, the reason is the deliverable. The bullets above all test whether the suggested fix *works*: against the code, against the reviewer's repro, against an edge case their own prose named. This one assumes it works, and asks whether the justification handed to you still holds when someone leans on it.

That distinction is invisible in code and decisive in a rule. A patch is judged by its behaviour, so a correct patch with a shaky rationale is merely under-commented. A **shared/** fragment is judged entirely by whether its reason forecloses the workarounds, so adopting a weaker reason ships a rule the next reader can talk themselves around — while the thread records the item as settled.

The tell is a suggestion that explains *why* something is forbidden in a single phrase, where the primary source carries a stronger provision. So ask what the strongest *available* reason is, rather than whether the offered one is defensible, and name the workaround the weaker reason would have licensed — that is what makes the choice checkable rather than a matter of taste.

- **Do:** read the primary source for the strongest reason before adopting a suggested rationale, even when the suggestion's conclusion is right.
- **Do:** say in the reply which reason you took and why the offered one was set aside, since deviating from a **suggestion** block silently reads as having missed it.
- **Don't:** accept a defensible-sounding mechanism because the conclusion it supports is correct.
- **Don't:** treat this as grounds to reject the finding — the conclusion usually stands, and only its reason needs strengthening.

(Morrison-Lab/ai-config#873, 2026-07-30: a review correctly found a CC-BY-ND table row that called verbatim copying allowed and then concluded idea-only with no bridge. Its suggested reason, "MIT grants modification rights; ND does not", frames the conflict as two grants differing in scope — which licenses the workaround of keeping the file under its own notice

inside the MIT repo, since on that framing no conflict arises. SPDX `license-list-data's CC-BY-ND-4.0.txt` §2(a)(1) grants a **non-sublicensable** license, so the material cannot be re-offered under MIT at all, which is exactly what vendoring does. The conclusion was right, and its stated reason stopped short of the provision that actually forecloses the workaround.)

And the mirror case: a finding can be wrong on its stated grounds while still pointing at something real. The bullets above check the reviewer's *fix*; this one checks their *premise*. A confidently reasoned factual claim – this pattern is valid, that value is in range, this call is safe – invites one of two lazy responses: accept it because it sounds authoritative, or dismiss the whole item once you notice the claim is false. Both lose information, because a reviewer usually arrives at a wrong premise while looking at something that genuinely bothered them.

So reproduce the claim before answering it, and answer the concern separately from the premise. When the premise turns out to be false, say so with the command and its output rather than by assertion, and then address what prompted it anyway – a reader who tested your example and got a different result has a real problem even if their explanation of it was wrong. Expect the corrected mechanism to be more useful than the original text: a premise worth disputing usually sits on something you had not fully explained. (ai-config#756, 2026-07-28: a review held that `[\x{2014}]` is valid PCRE and so could not produce the “code point value too large” error the fragment described, and proposed an out-of-range `[\x{110000}]` instead. Running it showed the original failing exactly as written – the cause is the locale, since PCRE in non-UTF mode rejects any `\x{}` above `0xFF`, and the same command succeeds under `LC_ALL=C.UTF-8`. The proposed replacement would have been worse, failing unconditionally and hiding that environment-dependence, which is the whole reason the swallowed error is dangerous. The reviewer's actual worry – that a reader might not reproduce it – was right, and sharper than stated.)

When a finding cites a source, read the cited source before reproducing anything – it is the cheaper instrument, and it is the one that can show the finding backwards rather than merely unsupported. The bullet above says to reproduce the claim. That is right, and it is the second thing to do when a citation is on the table, because reproduction tests the *behavior* while the citation tests the *reasoning*, and only the second can catch a finding whose own evidence contradicts it. A citation is also the most persuasive part of a review and the least likely to be checked: a linked changelog entry reads as settled fact, so the finding inherits authority it never earned, and a one-click **suggestion** block turns that borrowed authority into an applied edit.

Grep the cited document for the mechanism the finding names. One command usually decides it, which makes this an **algorithmatize-checks** case rather than a judgment call, and a fabricated mechanism produces a clean zero-hit result that is hard to argue with. Then quote the entry in the reply rather than paraphrasing the disagreement, and reproduce the behavior as the independent second leg.

Do not stop at winning the point. A finding that misread a source usually did so because the claim it questioned had nothing checkable next to it, so fold the citation into the file itself, per [fully-clean](#)'s note that a fresh review run re-derives from scratch and will not read the thread. (ai-config#762, 2026-07-28: a review held that `htmlwidgets::saveWidget(selfcontained = TRUE)` no longer needs `pandoc`, citing `htmlwidgets 1.6.0` as having “switched to `base64enc::dataURI()`”, and supplied a suggestion block deleting the `rmarkdown::pandoc_available()` gate. `grep -inE 'pandoc|base64'` over that NEWS file returned six `pandoc` hits and zero `base64` hits, and the 1.6.0 entry says the path “now uses the `{rmarkdown}` package to discover and call `pandoc`” – so the citation established the opposite of the finding, and incidentally made the gate the *same lookup* `htmlwidgets` performs rather than a proxy for it. Applying the suggestion would have removed the only warning before a hard error, in the one step that exists for running headless. The reviewer accepted the rebuttal on the next round and called its own prior claim a hallucination.)

When a reviewer hedges a finding because it depends on code it cannot see, check whether *you* can see it — the hedge is an invitation, not a verdict. Automated reviewers work from the diff, so a finding that turns on a reusable workflow, a dependency's internals, or another repo's behavior arrives with language like “moderate rather than high confidence”, “depends on behavior not visible in this diff”, or “worth the author confirming intent”. That hedge is a fact about the *reviewer's* visibility, not about how likely the finding is. You frequently have access it lacks: the repo cloned locally, a pinned dependency vendored in, or permission to fetch the source.

Reading it converts a maybe into a settled yes or no, and that changes the disposition. Confirmed, it earns a fix or a precisely-scoped follow-up issue with the mechanism recorded; disproved, it earns a Rebut with evidence instead of a vague “I think this is fine”. Either way the next reader is spared re-deriving it. Quote the specific lines you checked, since a follow-up issue that merely repeats the reviewer's hedge is barely more useful than the review comment it came from.

(UCD-SERG/serodynamics#274, 2026-07-28: a review flagged possible duplicate review dispatch at moderate confidence, explicitly because the reusable workflow in `d-morrison/gha` was not visible to it. That repo was cloned locally. Reading both matchers showed the reusable fires on `@claude[[:space:]]+review` and the local job on a punctuation-tolerant superset, so the plainest phrasing — `@claude review` — matches both and dispatches twice. The follow-up issue could then record the exact overlap table and note that the upstream gap motivating the local job had since been closed, making “broaden upstream, delete the local job” a real option.)

Timestamp the evidence before rebutting a finding with it — during a live incident, a log from twenty minutes ago describes a different system. The bullets above all say to verify a finding rather than accept it, and they assume verification is a fixed target: read the source, run the command, reproduce the case. That assumption quietly fails while something is actively breaking, because the evidence you gather is a *measurement*, and measurements

expire. Re-reading an existing CI log feels like verification — it is concrete, it is specific, it is right there — but it only tells you what was true when that job ran.

The tell is a rebuttal whose evidence you did not generate yourself in this turn. A log you fetched, a check-run conclusion you read, a status you were told about: each carries a timestamp, and the question is whether anything could have changed since. When the finding is *about* an outage, a migration, a permission change, or anything else in flight, the answer is almost always yes.

So prefer evidence you can regenerate now over evidence you can only cite. Re-running the failing thing is usually cheap and settles it outright — and in the best case it produces the cleanest possible proof, two attempts of the same run on the same commit disagreeing, which no amount of reading could have given you. When regenerating is genuinely not possible, say how old the evidence is in the rebuttal itself, so the reader can weigh it.

This matters more than an ordinary wrong rebuttal because of who it lands on. Telling an author their diagnosis is contradicted by the logs is a strong claim that invites them to stop investigating. Getting it wrong can stall a correct fix for the exact bug still breaking everything. (gha#351, 2026-07-28: a PR correctly diagnosed that Actions had stopped resolving `uses:` after a repo transfer. Its premise was disputed on the strength of two run logs showing the workflow resolving fine — logs from 45 and 30 minutes before the PR was opened, spanning the cutover. Re-running one of those very workflows reproduced `startup_failure` immediately, and the retraction had to be published in the same thread.)

A finding built on a *negative* result – “I searched and it isn’t there” – is only as strong as the paths that were searched, and the search scope is the part reviewers state loosest. The bullets above all check a reviewer’s positive evidence: the suggested literal, the proposed fix, the cited source. A negative result invites none of that scrutiny, because there is nothing to look up: the claim is that looking up would fail. It also arrives sounding the most settled of any finding – “no file or heading with that title anywhere” reads as exhaustive, and the reader’s natural move is to accept it and edit.

So read the search itself rather than the conclusion. Ask which paths were actually covered, and whether the obvious location is among them. One command usually settles it, which makes this an `algorithmatize-checks` case rather than a judgment call – and note the reviewer’s own tooling may have failed the way `fail-fast`’s hand-check section describes, matching too narrowly against text that was wrapped or reformatted.

When it turns out the thing does exist, name the gap rather than only the correction: which paths were searched, where it actually lives, and why the two did not overlap. That is what stops the same search being re-run the same way. And check whether the finding still points at something real, per the mirror case above – an unresolvable-looking citation is often a genuinely under-specified one.

- **Do:** ask which paths a negative finding actually searched, and check the obvious location yourself before editing anything.

- **Do:** name the gap when the thing does exist – paths searched versus where it lives – so the same search is not re-run the same way.
- **Don’t:** accept “it isn’t there anywhere” as settled because it is stated more confidently than a positive finding would be.
- **Don’t:** discard the finding once its negative result is disproved – the thing it tripped over is often a real ambiguity.

(Morrison-Lab/gha#338, 2026-07-28: a review reported a cited section as nonexistent, having “checked ai-config’s full tree (`shared/workflow/*.md`, `skills/`, `codex-skills/`)”. The heading was an H2 in that repo’s **root** `CLAUDE.md`, the one directory those three paths skip. The reviewer had even found the phrase in `shared/workflow/fully-clean.md` and read it as pointing at a *consuming* repo’s `CLAUDE.md`. The rebuttal carried the one-line grep; the underlying point was real anyway, since citing a section title without naming its file is what sent the search to the wrong directories, so the citation was fixed to name and link the file.)

References